

# Memory Fragmentation Analyzer - Mallice

Mohammed Al-Bow, Jeffrey Edgington, Yan Mayster,  
Daniel Pittman, Christian Grothoff

`{malbow,jedgingt,ymayster,dpittman,grothoff}@cs.du.edu`



# Memory Consumption Still Matters

- Cache Effects (e.g., small cache size)
- Power Consumption (e.g., embedded devices, battery life)
- Hardware Cost



# Causes of Inefficient Memory Use

- Inefficient Data Structures
- Memory Leaks
- Memory Fragmentation
- Memory Management Overheads
  - Garbage Collection — often a factor of 3!
  - Free Lists

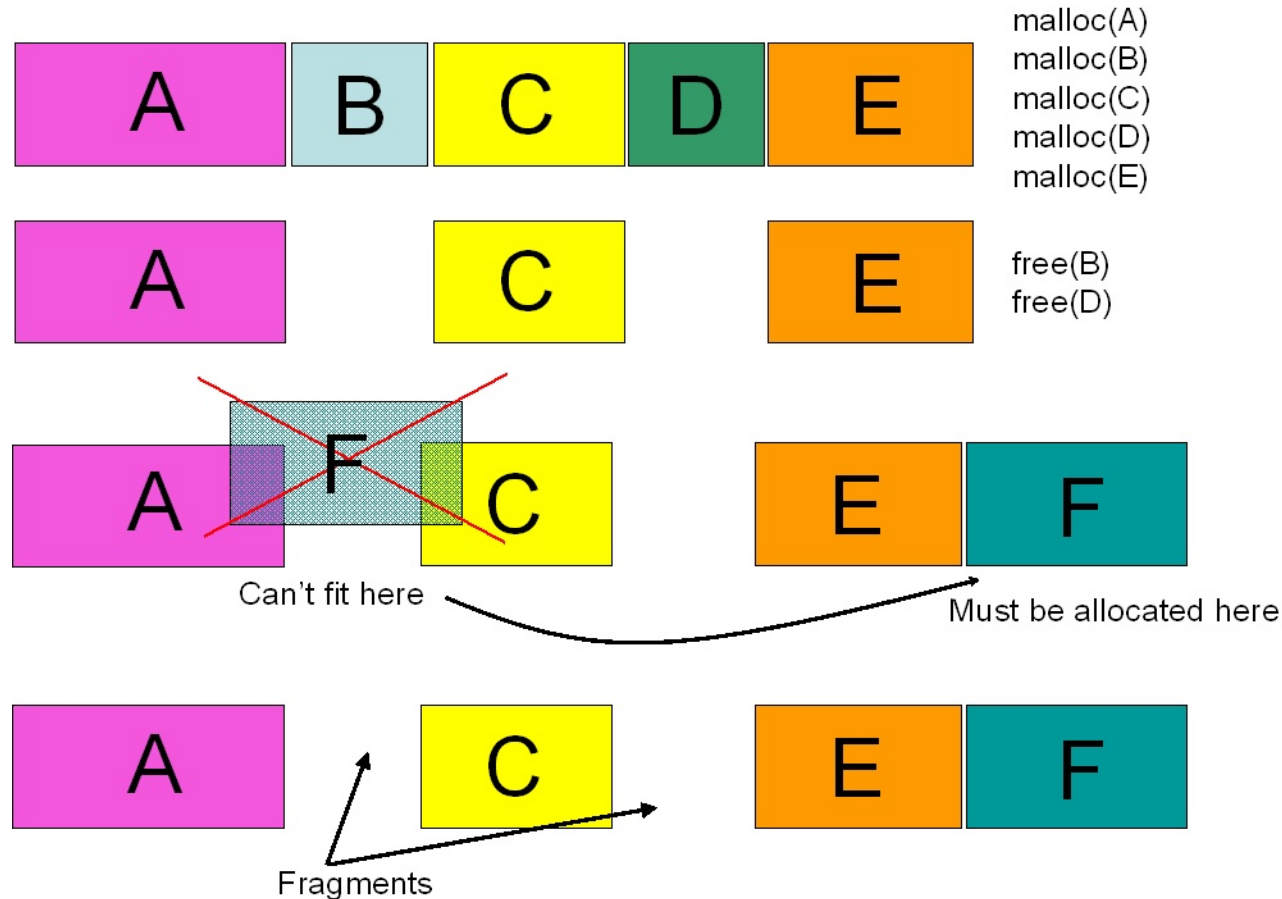
⇒ Today: Focus on manual memory management.



# Existing Solutions

Problem	Solution
Inefficient Data Structures	<i>Virgil, Valgrind/Massif</i>
Memory Leaks	<i>Valgrind/Memcheck, Purify, Electric Fence</i>
Memory Management Overheads	<i>GcSpy (for garbage collection)</i>

# Memory Fragmentation Illustrated



# MM Overheads Illustrated



Typical malloc memory layout.



# Other MM Overheads

- 8K OS page sizes
  - System call reduction
- ⇒ Overallocation with sbrk



# Code with Memory Fragmentation

**Challenge:** Give advice to the programmer where to improve memory allocation/freeing code.

## Example code:

```
a = malloc(4);  
b = malloc(400);  
c = malloc(4);  
free(b);  
\\ Compute  
free(c);  
free(a);
```

## Optimized code:

```
a = malloc(4);  
c = malloc(4);  
b = malloc(400);  
free(b);  
\\ Compute  
free(c);  
free(a);
```





# Performance Metrics

1. Total allocation (in Bytes  $\times$  Cycles):

$$\sum_{i=0}^{T_{max}} m_i \quad (1)$$

2. Peak allocation (in Bytes):

$$\max_{i=0 \dots T_{max}} m_i \quad (2)$$




# Problems with Memory Analysis

- Inefficient memory use is hard to diagnose
- Memory operations are frequent
  - ⇒ need lightweight instrumentation
- Instrumentation impacts application behavior



# Architecture



## Architecture (cont.)

Application  $\leftrightarrow$  libc  $\leftrightarrow$  libmallice  $\rightarrow$  Mallice Daemon

### Achieved Design Goals:

- No dynamic allocation by libmallice  
 $\Rightarrow$  accurate picture of the heap!
- Zero-overhead for programs that do not use malloc




# Techniques Used

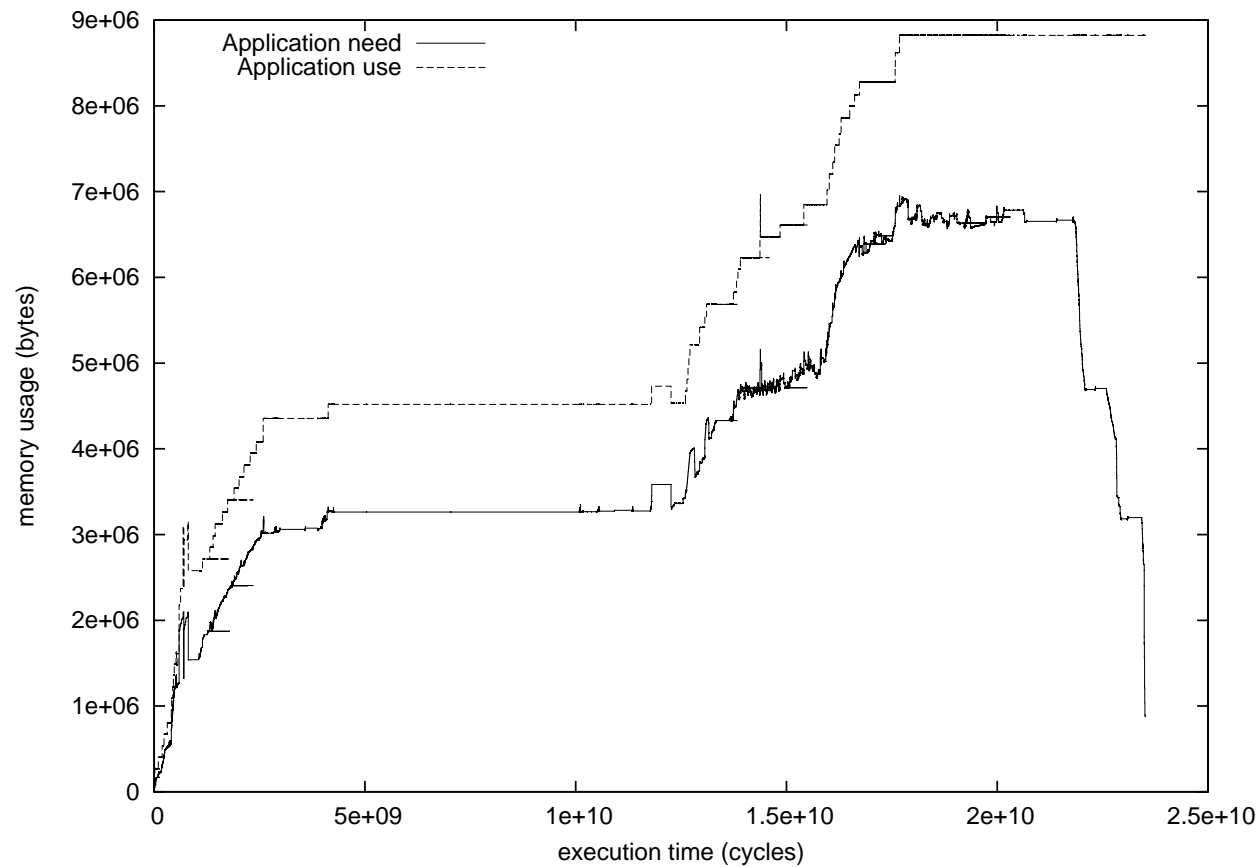
- Use malloc wrappers hooks from glibc
- /proc/self/maps: reverse engineers mmmaps by malloc
- LD\_PRELOAD is used to add instrumentation dynamically



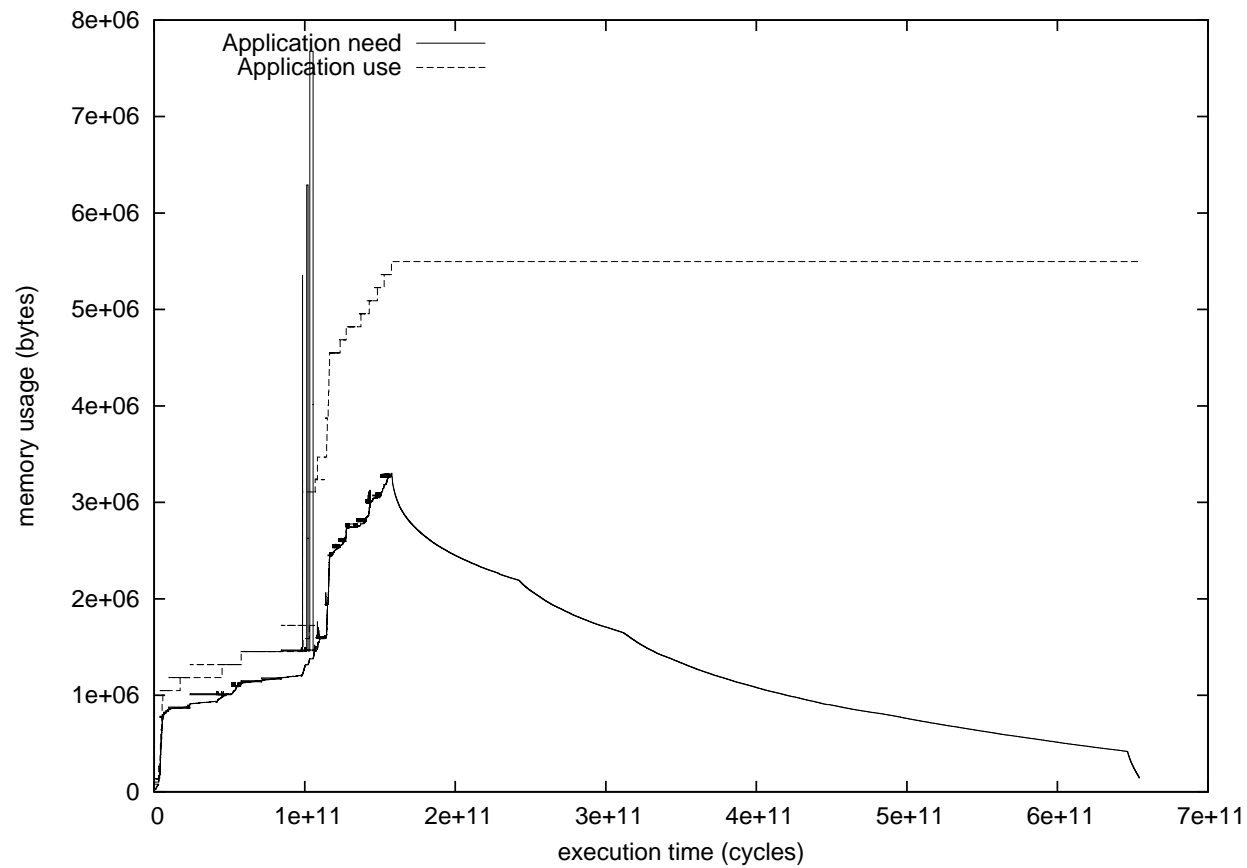
# dpkg --list (package management)



# Konqueror (browser)




# doodle (suffix-tree)







# xpdf (viewer)



# extract (file parsers)



# Pattern Detection



# Performance Measures

Instrumentation overhead (execution time):

Benchmark	Normal Run	with libmallice
gcc	0.170s	0.294s
xpdf	0.123s	2.851s
latex	0.272s	5.235s
xmms	2.594	6.391s



# Benchmarks under Consideration

## Embedded:

- firefox
- konqueror
- xpdf
- gzip
- display

## Background:

- postfix
- mysql
- X11
- xmms
- kicker
- sendmail

## Algorithms:

- latex
- doodle
- extract
- convert
- gcc
- gimp

## Common:

- bash
- konsole

Additional suggestions welcome!



# Questions

?