

# Parallelizing Irregular Gauss-Seidel Using Full Sparse Tiling

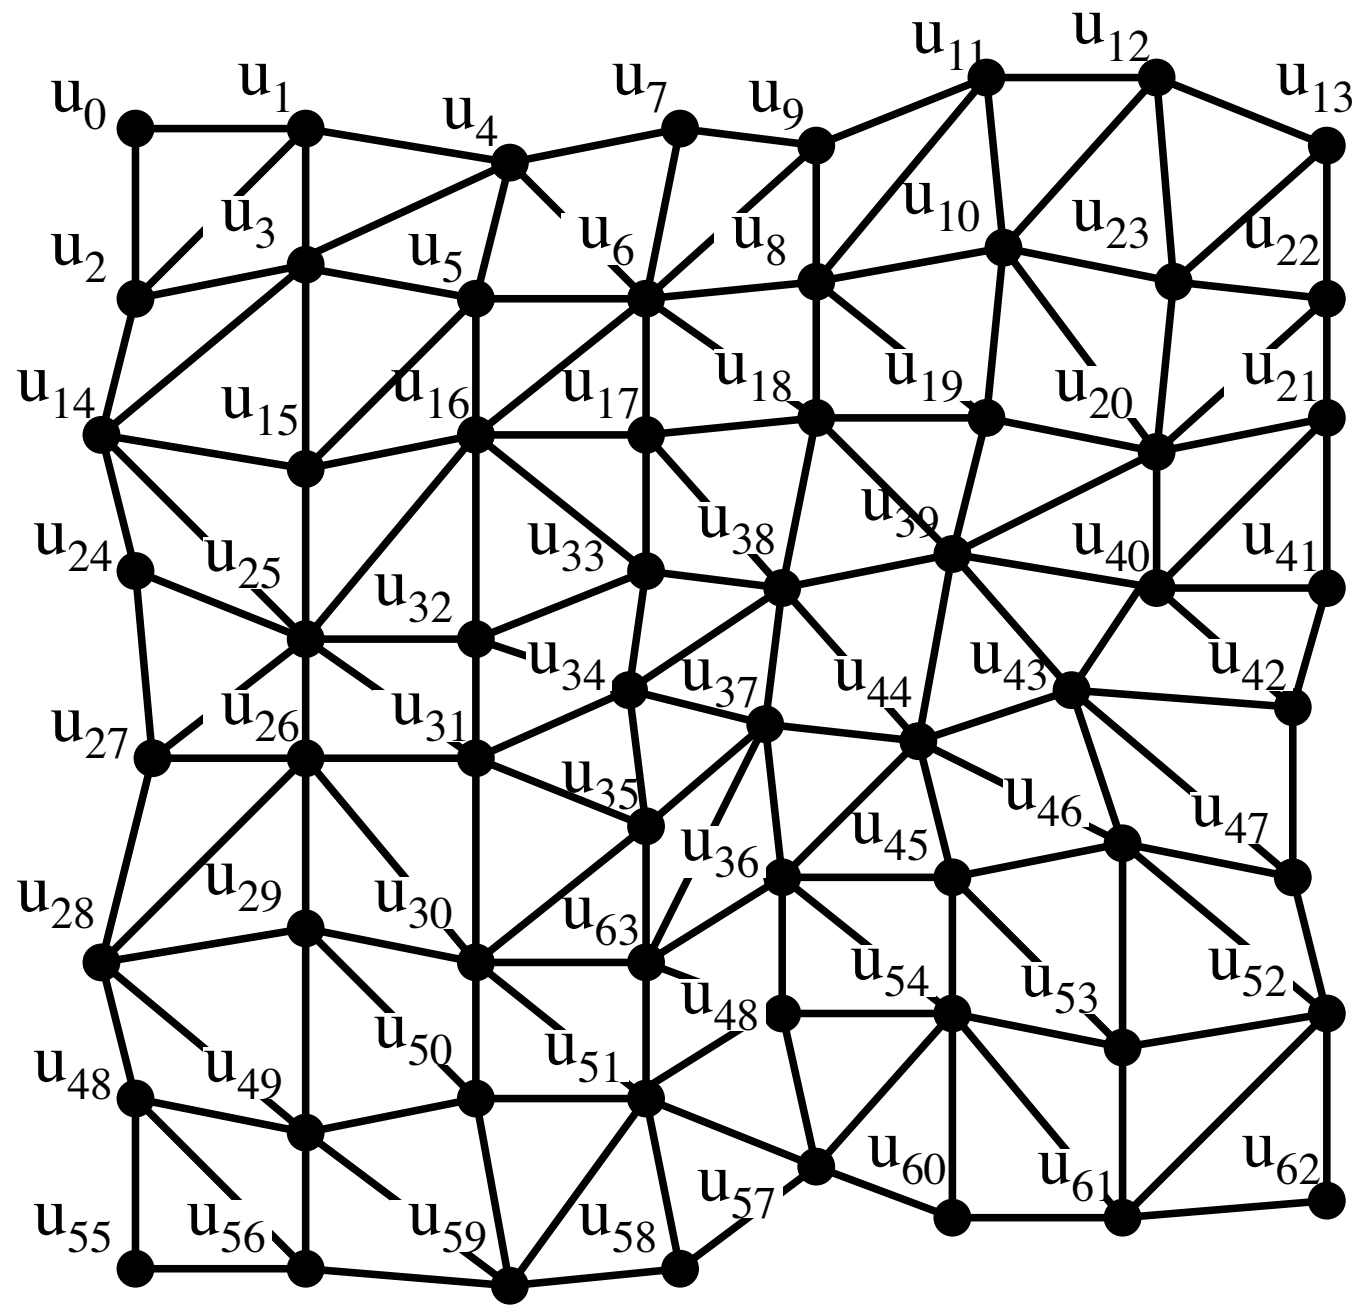
Michelle Mills Strout

in collaboration with Dave Rostron, Jeanne Ferrante, and Larry Carter  
partially supported by CSCAPES ([www.cscapes.org](http://www.cscapes.org)), an institute  
funded by DOE's Office of Science

FRACTAL

April 26, 2008

# Goal: Make computations over meshes fast



- Iterative smoothers such as Gauss-Seidel dominate the execution time of finite element applications
- Need to effectively utilize the memory hierarchy
- Performance should scale to multiple processors

# Full Sparse Tiling to the Rescue!

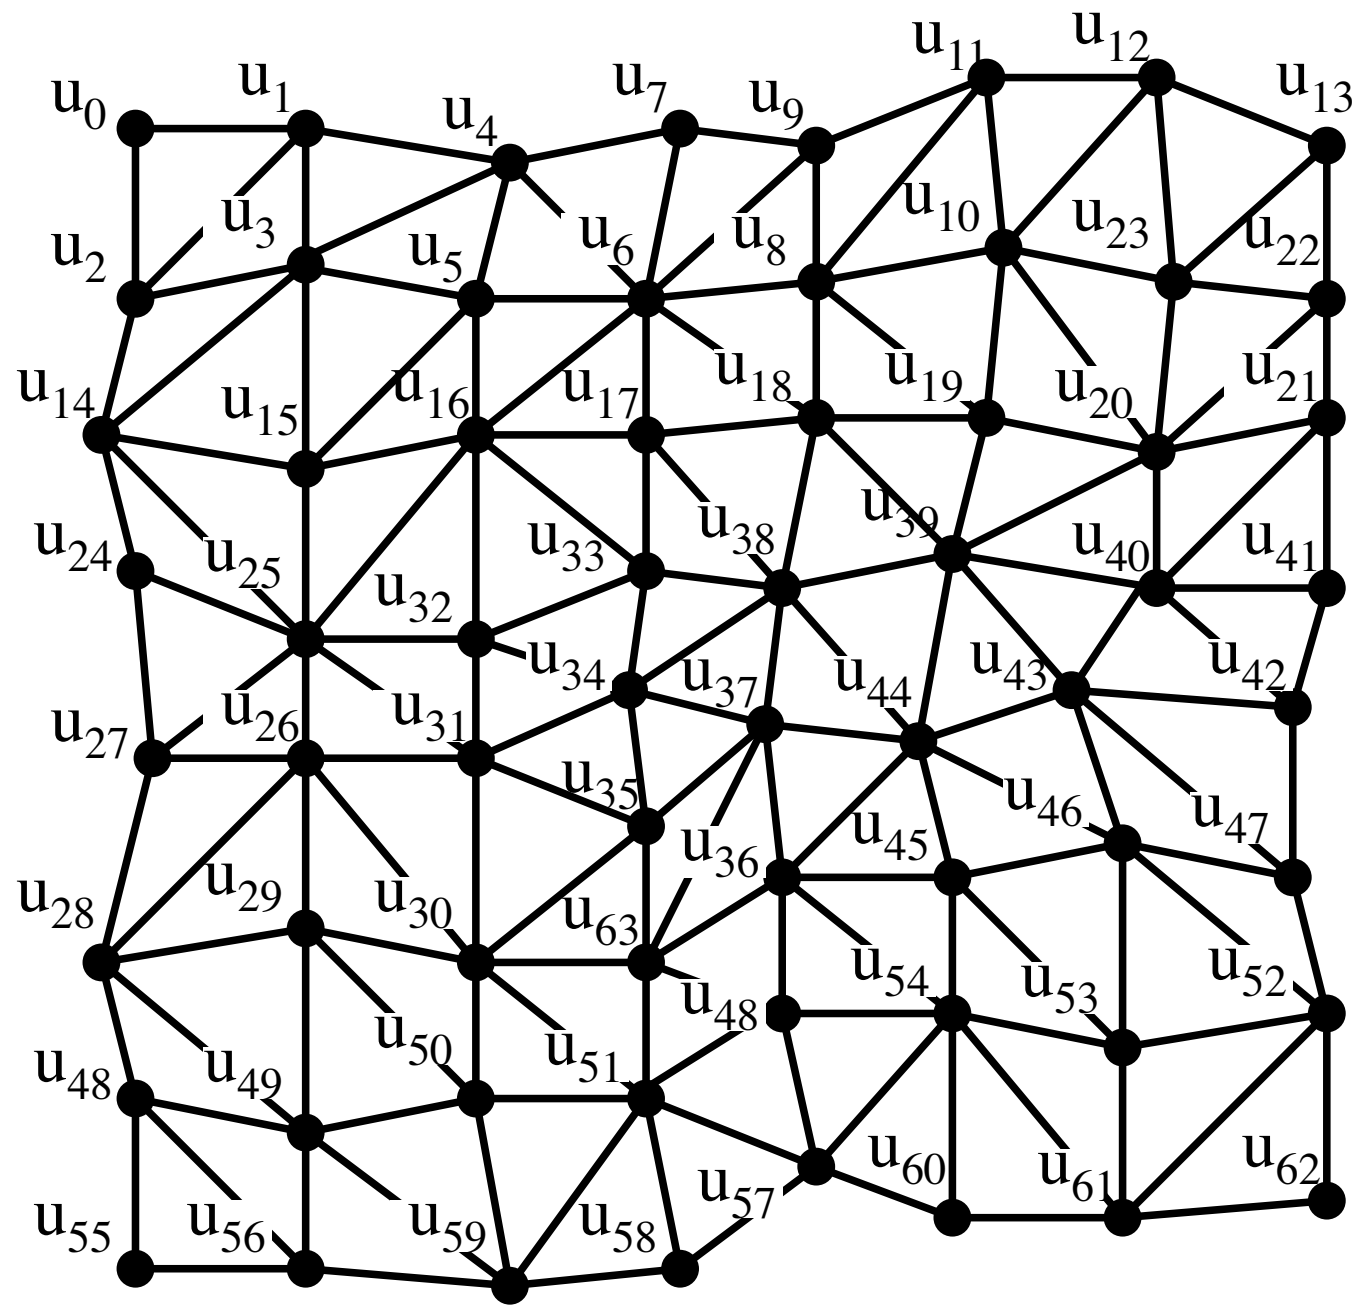
- **Problem**

- Compile-time data reordering and computation scheduling to improve data locality and parallelism is not possible due to irregular memory references  $A[B[i]]$

- **Solution**

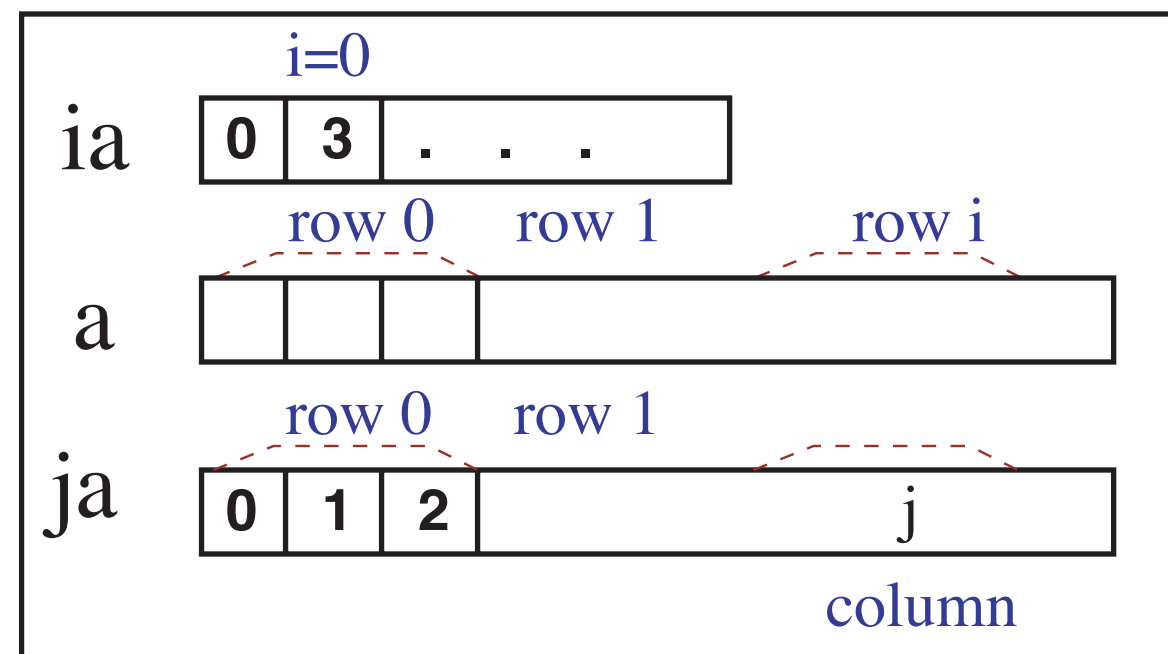
- Inspector/executor strategies perform data and computation reordering at runtime
- Full sparse tiling is one such strategy that improves performance by exploiting parallelism, intra-iteration data reuse, and inter-iteration data reuse

# Gauss-Seidel Iteratively Solves $Au = f$



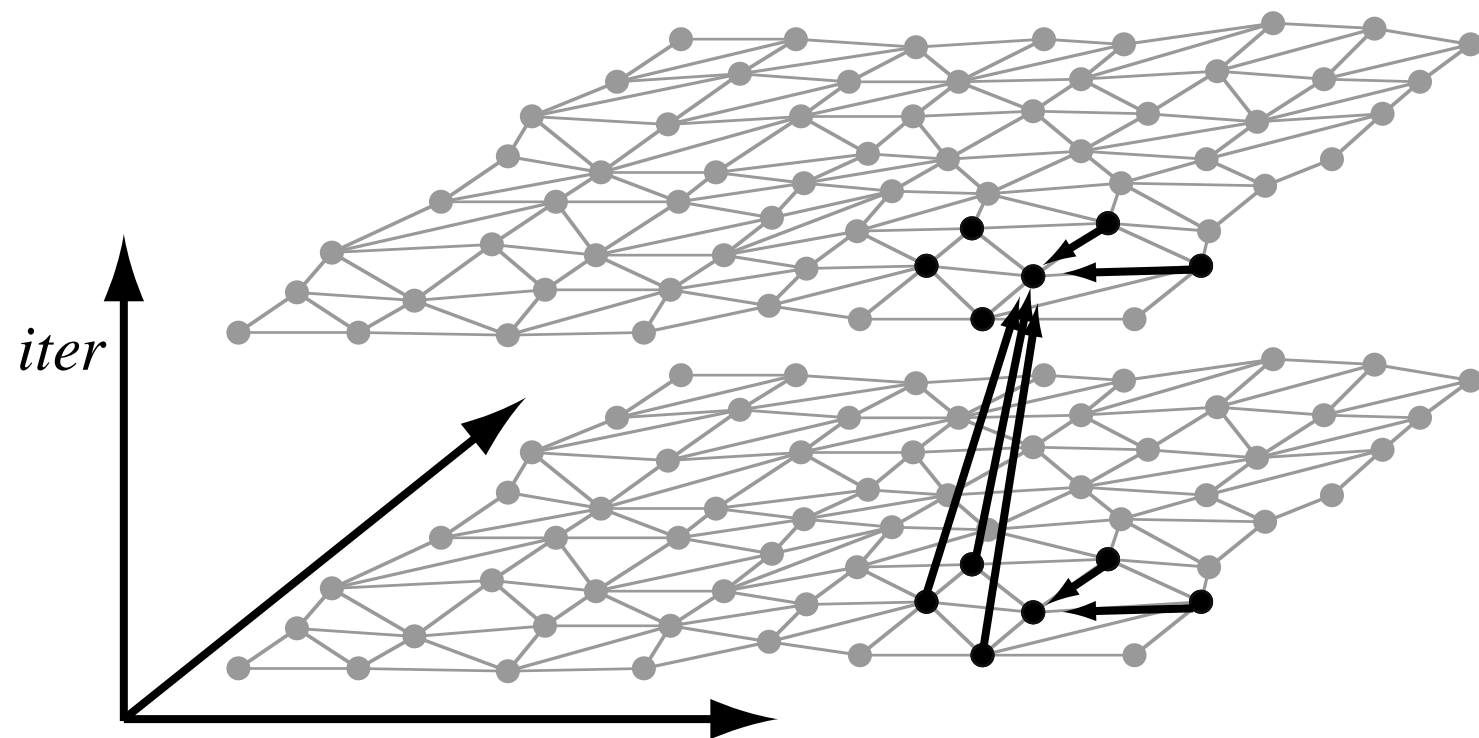
Matrix Graph

- $u$  is a vector of unknowns
- $A$  is a sparse matrix stored in the compressed sparse row format (CSR)

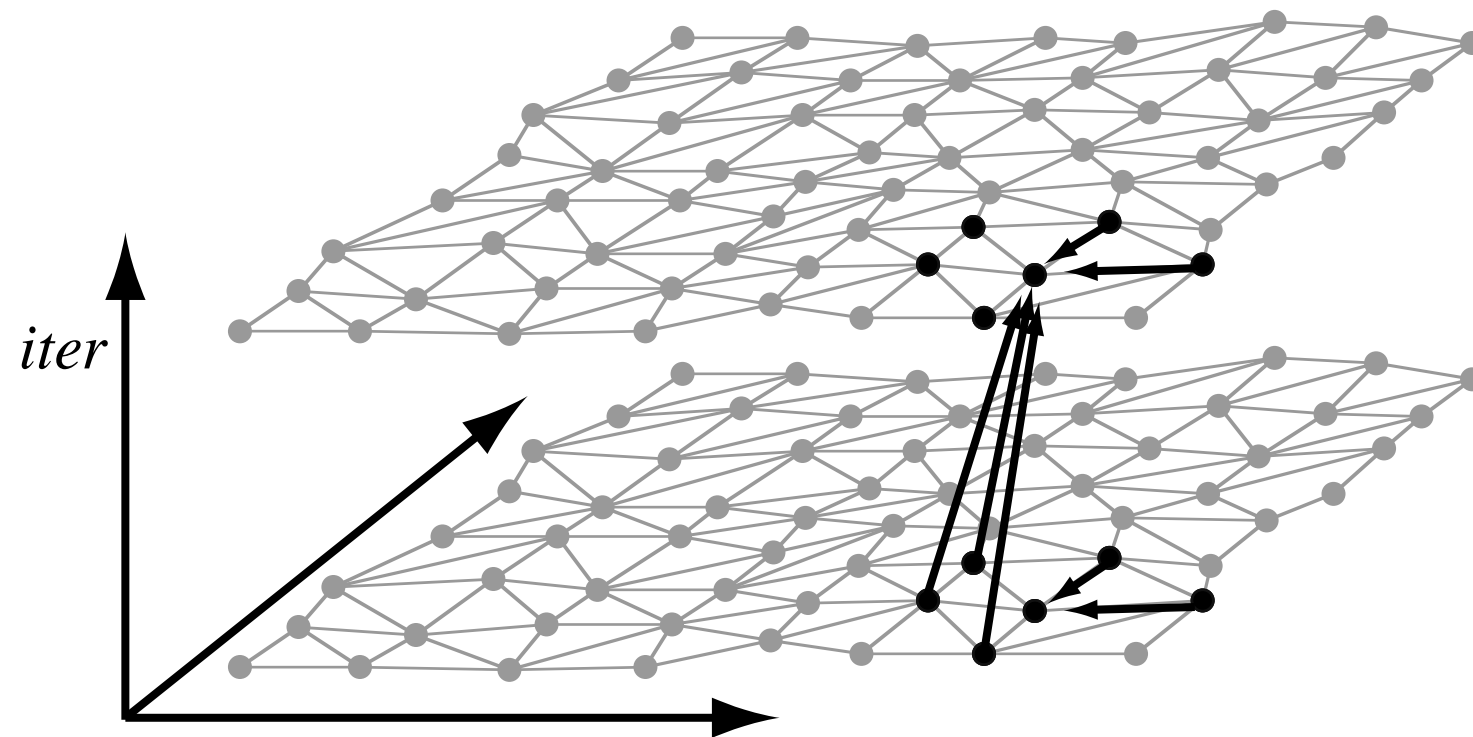


# Gauss-Seidel Iteration Space

```
do iter = 1, T
do i = 1, R
  u(i) = f(i)
  do p = ia(i), ia(i+1) - 1
    j = ja(p)
    if (j != i)
      → u(i) -= a(p) * u (ja(p))
    else
      diag = a(p)
    endif
  enddo
  u(i) = u(i) / diag
enddo
enddo
```



# Performance Improvement Opportunities

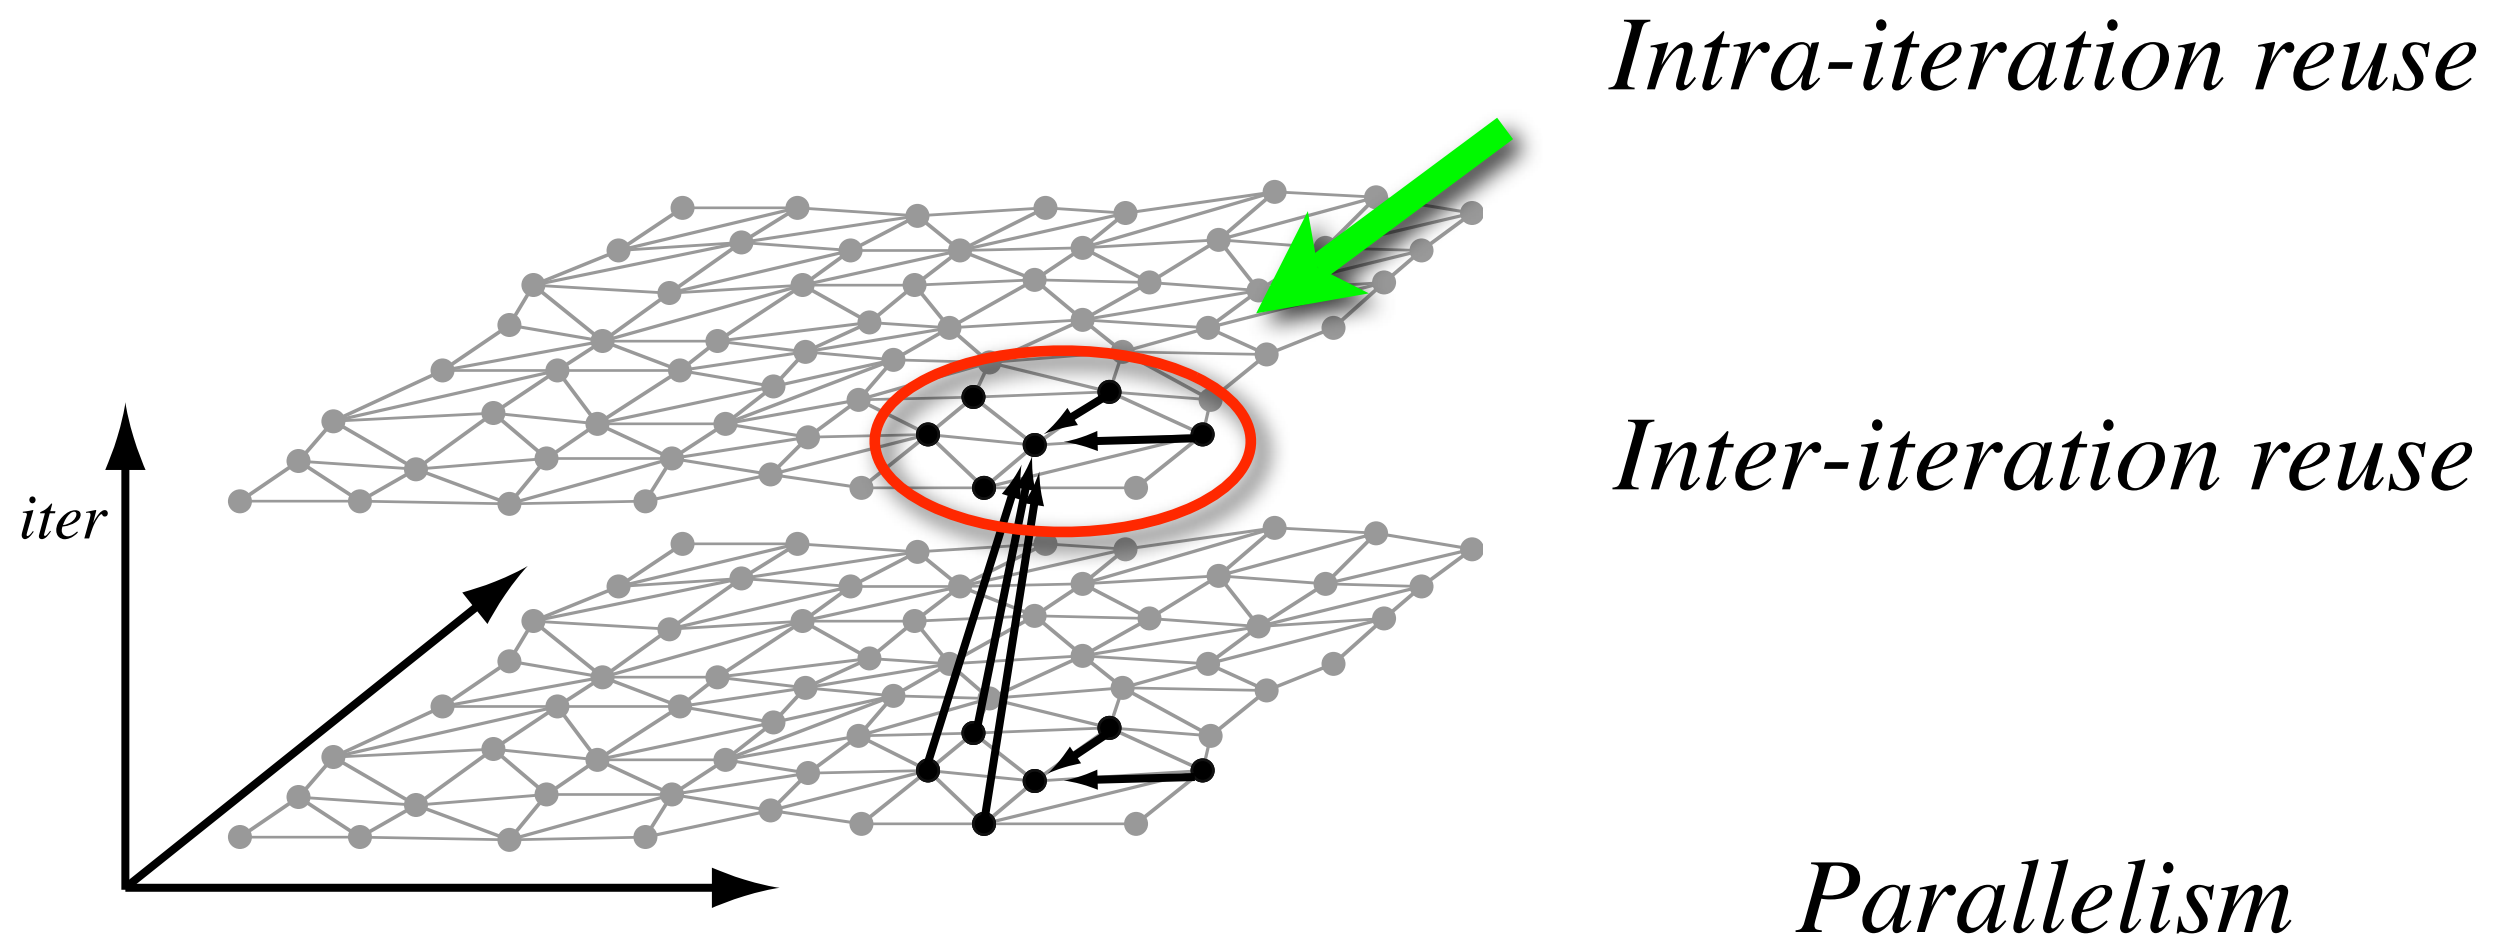


*Intra-iteration reuse*

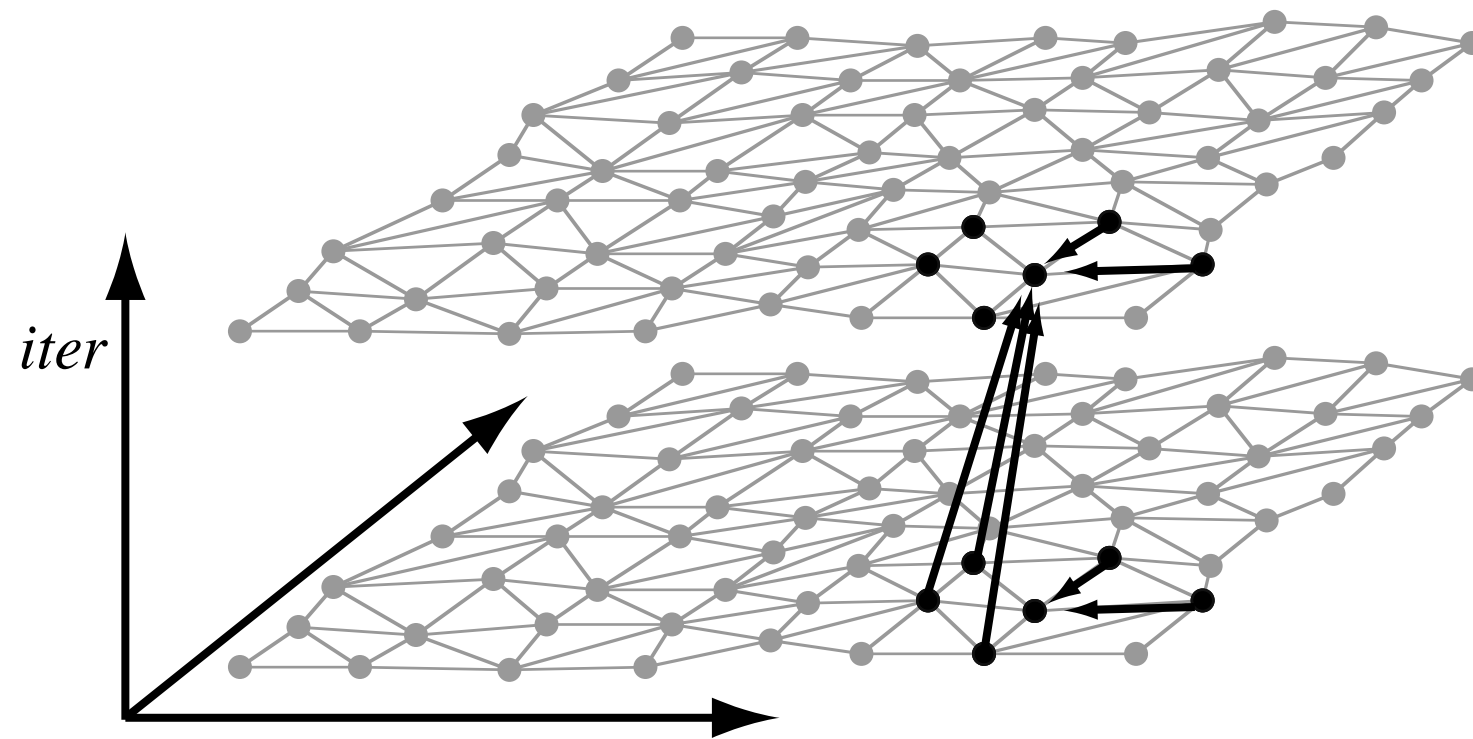
*Inter-iteration reuse*

*Parallelism*

# Performance Improvement Opportunities



# Performance Improvement Opportunities



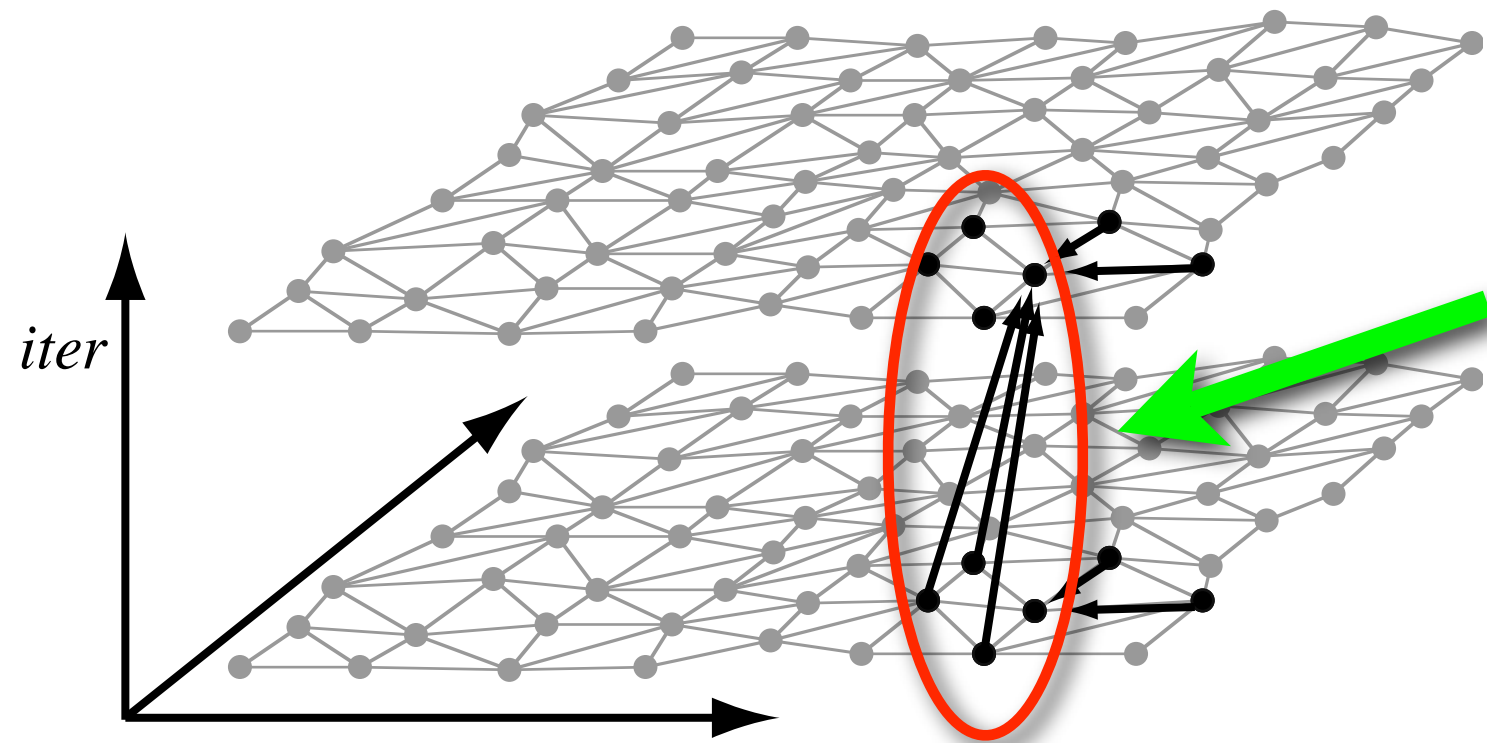
*Intra-iteration reuse*

*Inter-iteration reuse*

*Parallelism*



# Performance Improvement Opportunities

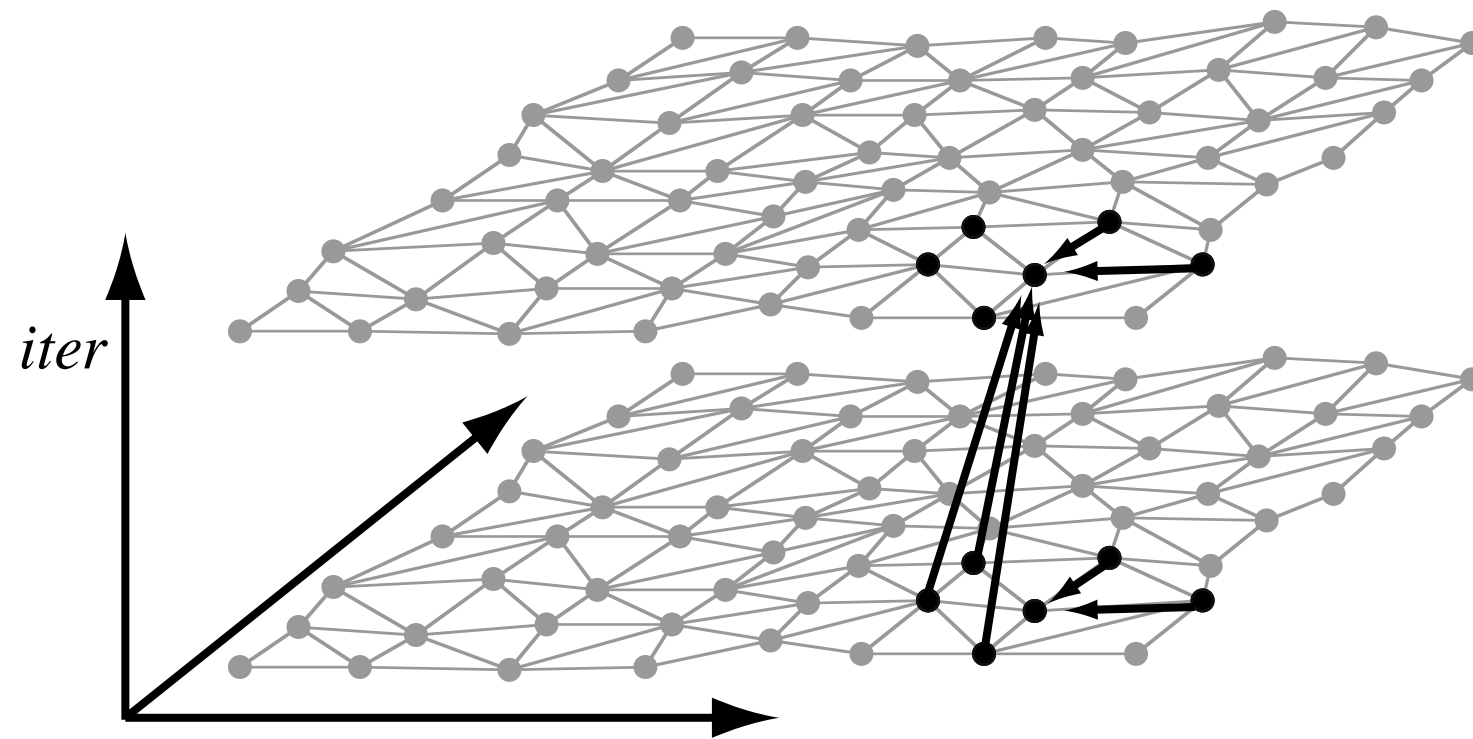


*Intra-iteration reuse*

*Inter-iteration reuse*

*Parallelism*

# Performance Improvement Opportunities

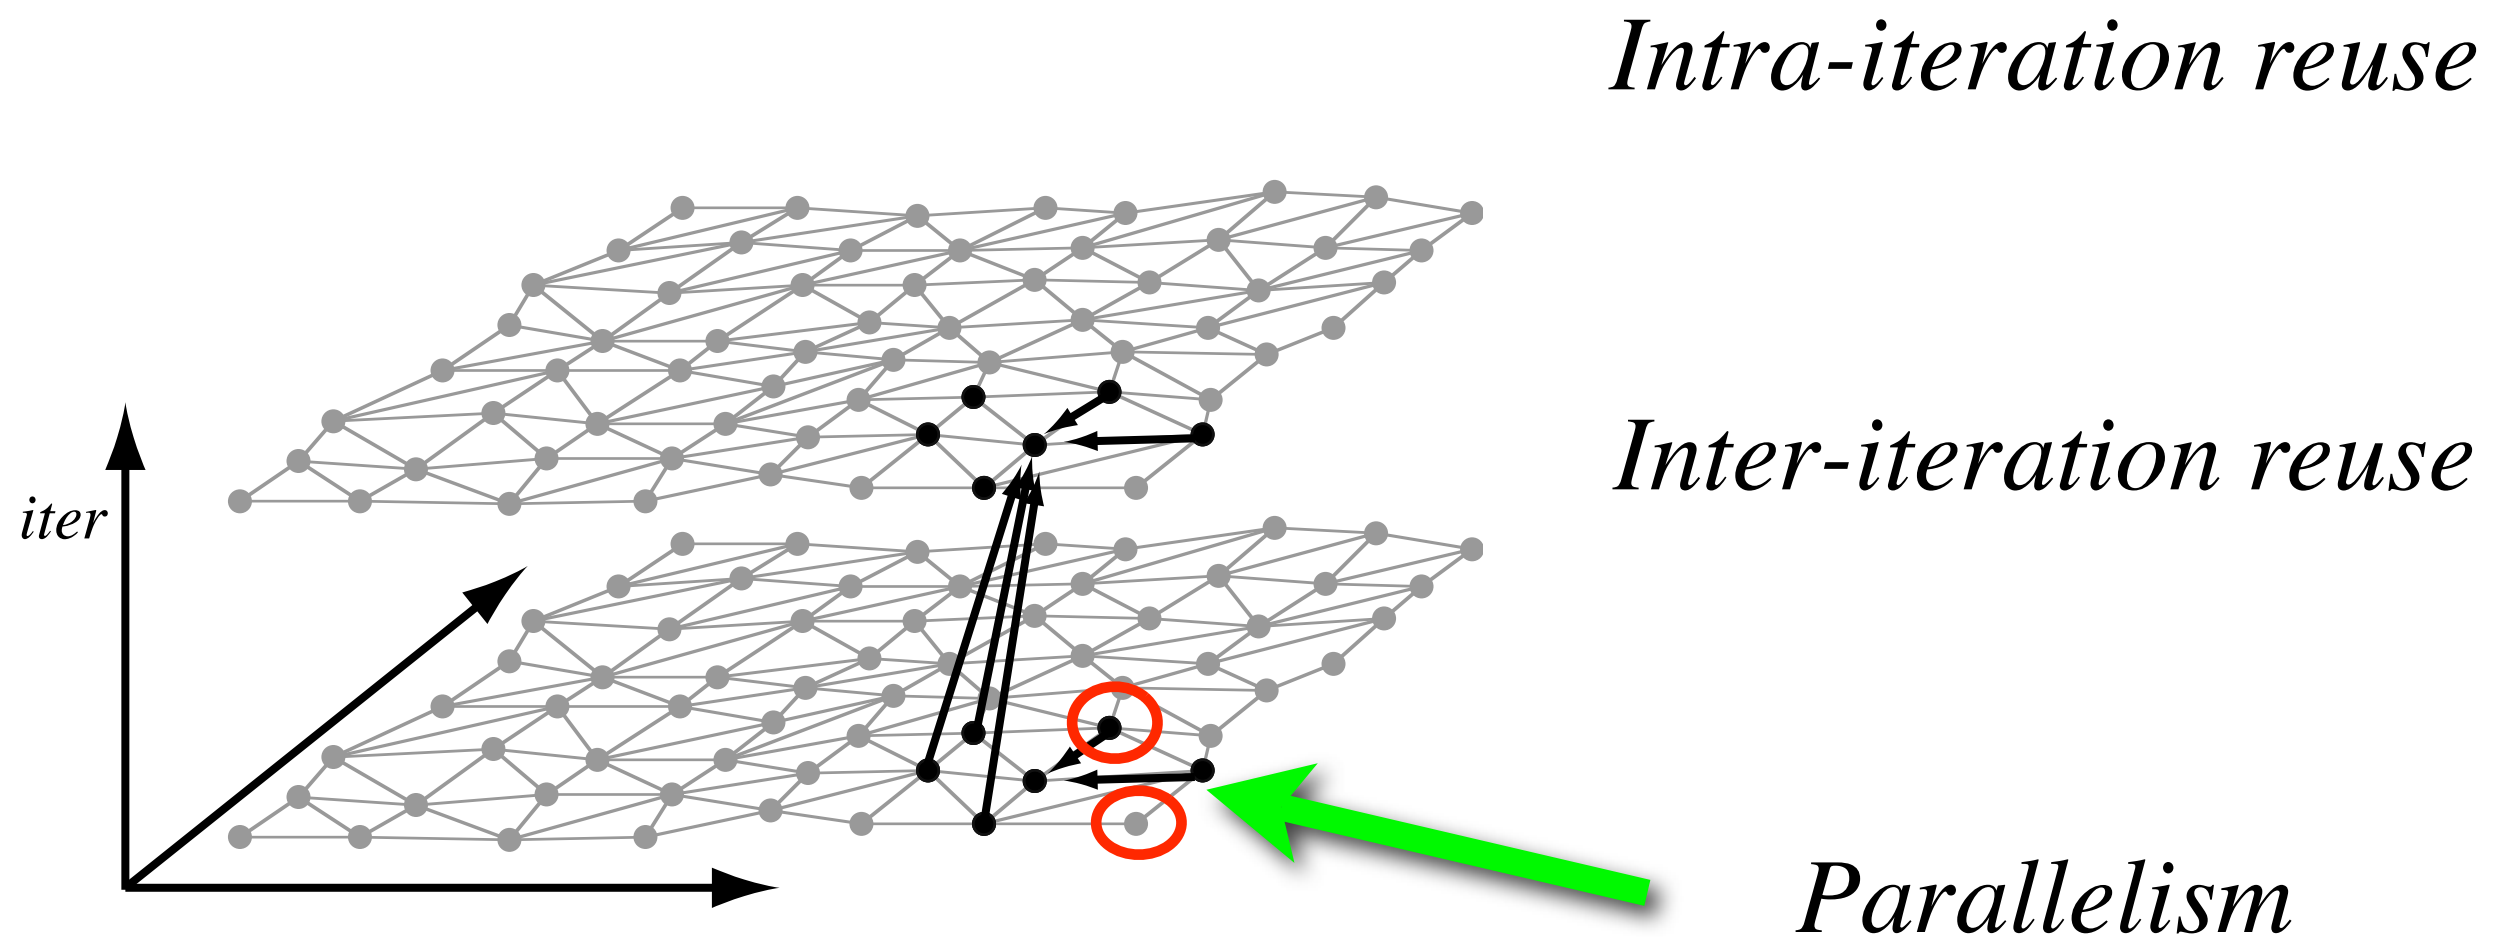


*Intra-iteration reuse*

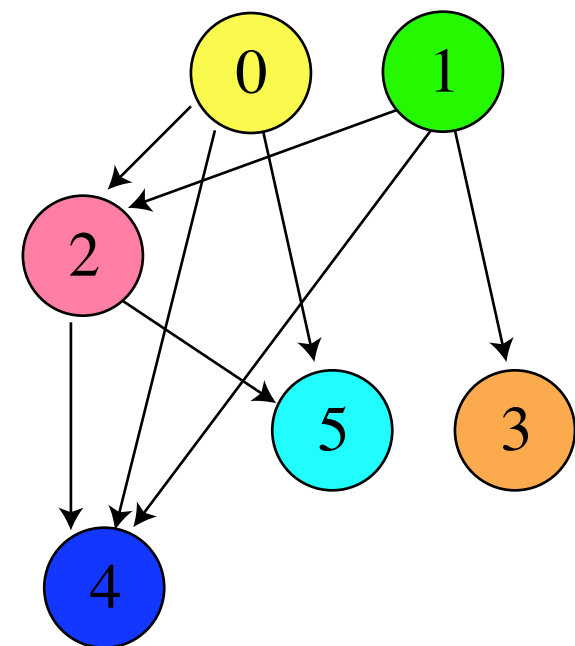
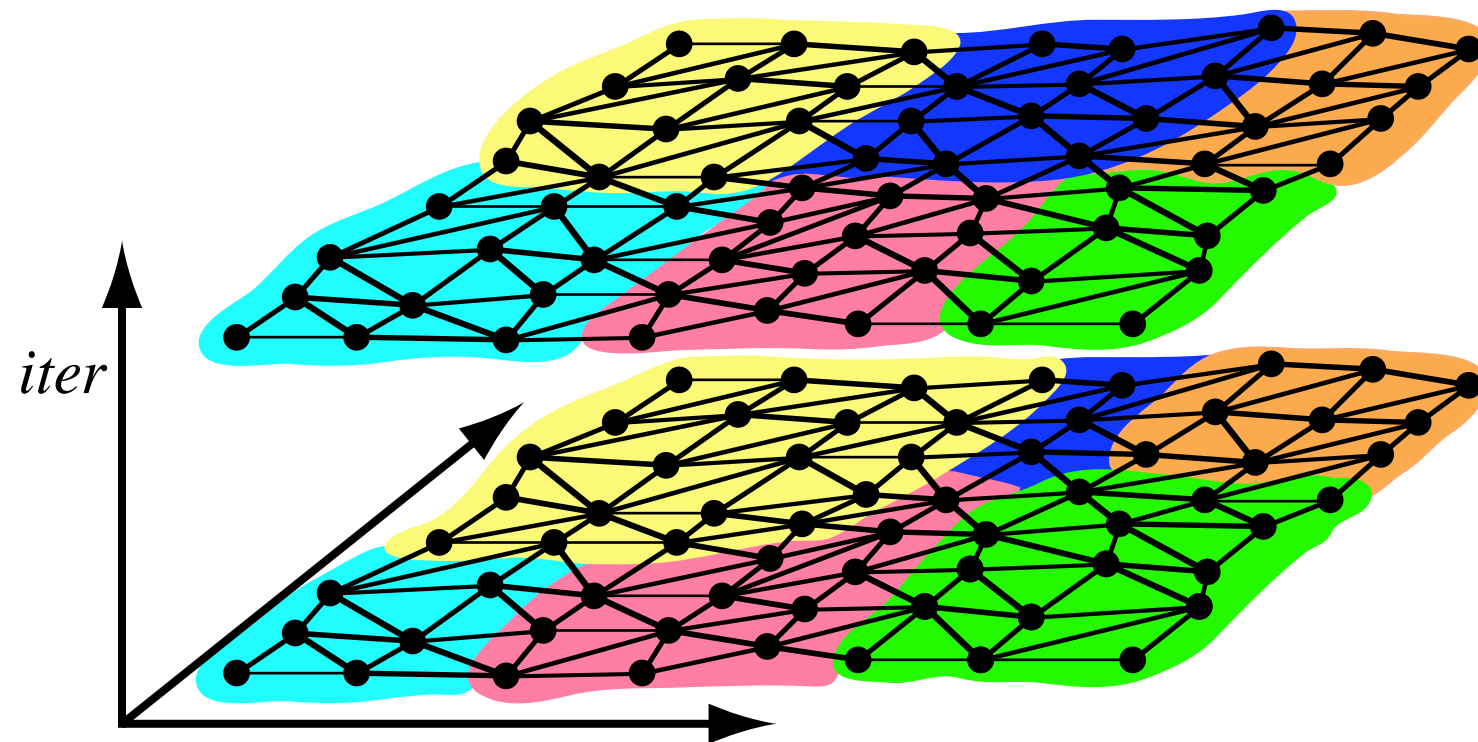
*Inter-iteration reuse*

*Parallelism*

# Performance Improvement Opportunities



# Punchline: Full Sparse Tiling



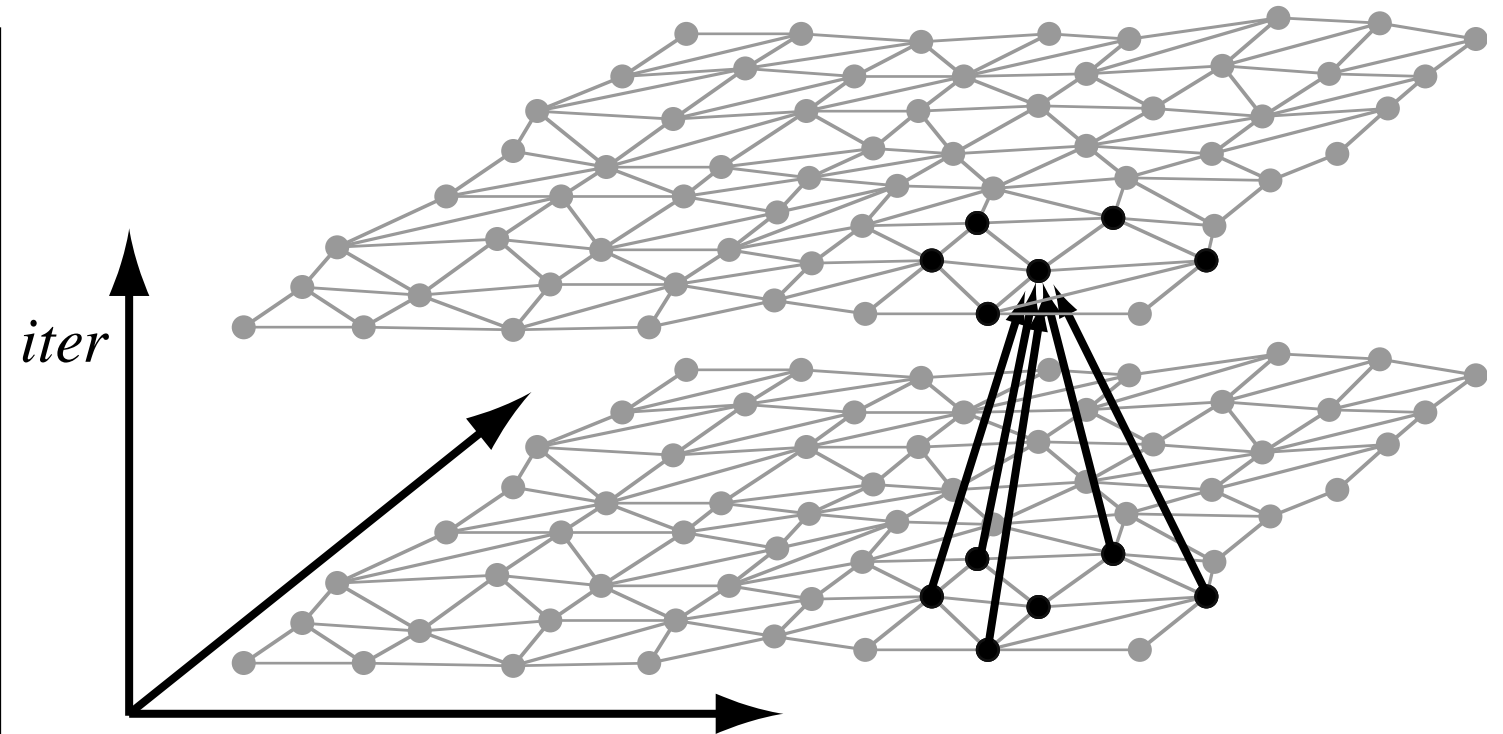
- Breaks up computation into pieces
- Each piece has intra and inter-iteration data locality
- Some pieces can be executed in parallel

# Talk Outline

- Overview
- Parallelizing Jacobi with the owner computes method xor full sparse tiling
- Increasing the parallelism in full sparse tile schedules
- Extra work needed to handle Gauss-Seidel
- Experimental Results

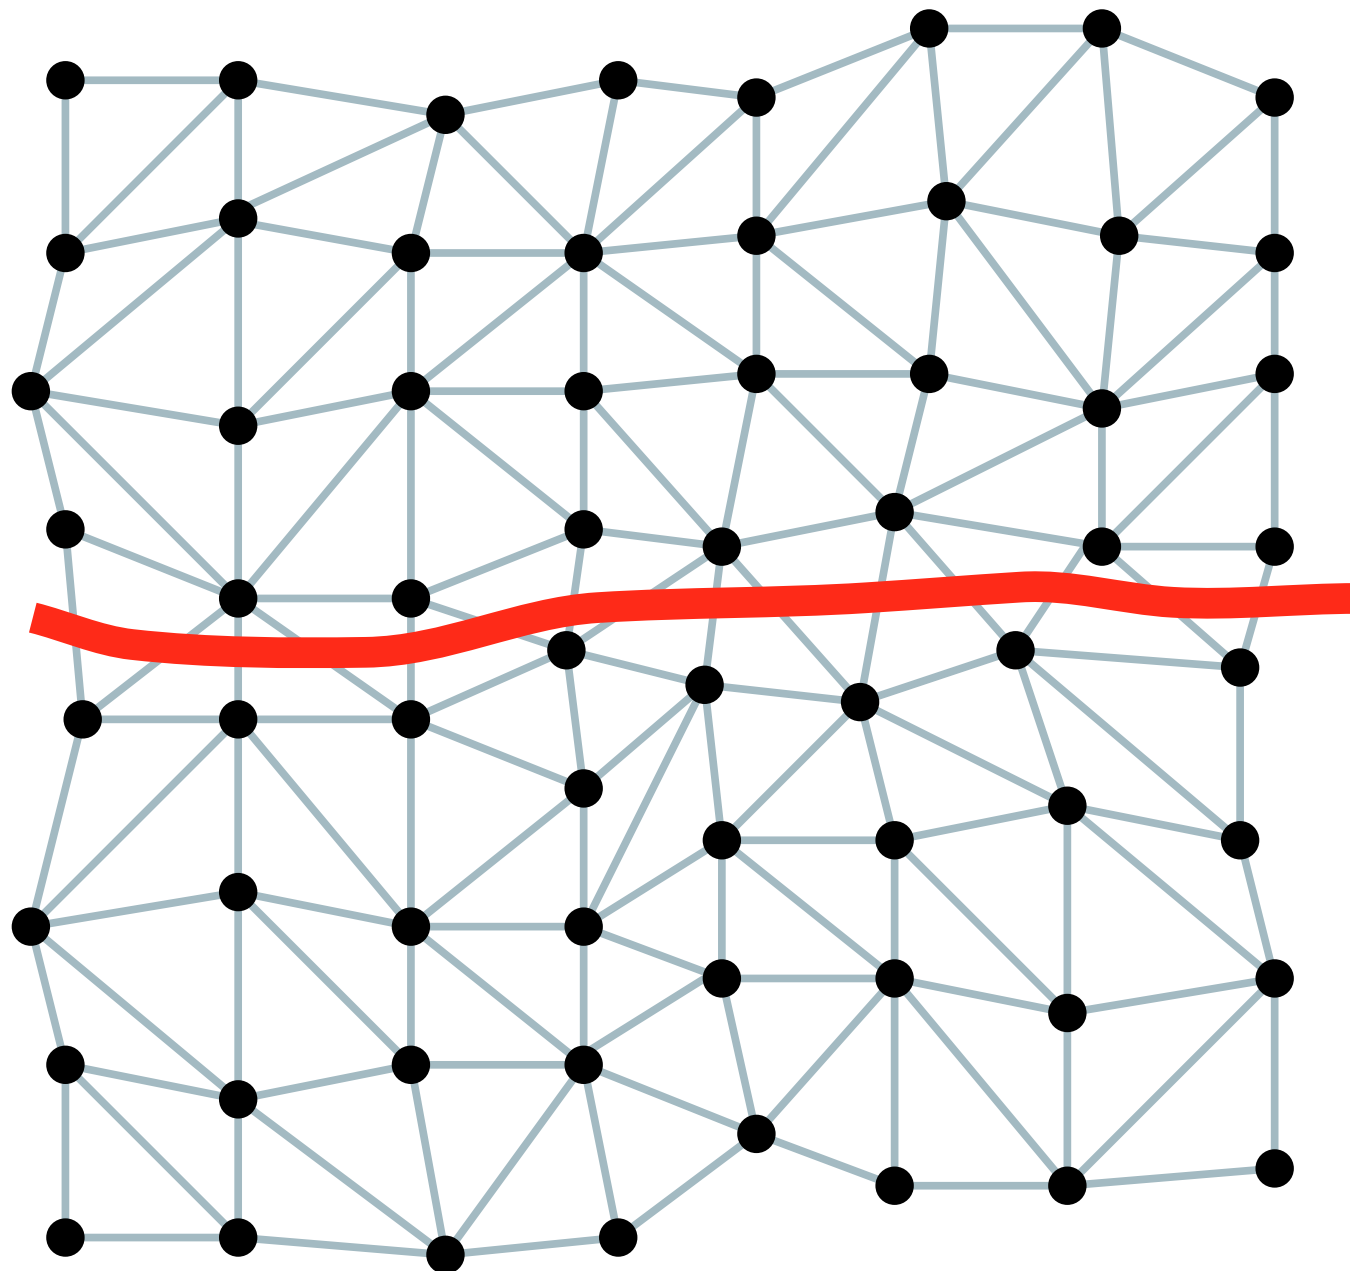
# Jacobi: similar yet simpler

```
do iter = 1, T
  do i = 1, R
    u(i) = f(i)
    do p = ia(i), ia(i+1) - 1
      j = ja(p)
      if (j != i) then
        → u(i) -= a(p)*tmp(j)
      else
        diag = a(p)
      endif
    enddo
    u(i) = u(i)/diag
  enddo
  do i = 1, R
    tmp(i) = u(i)
  enddo
enddo
```



- data dependences not known until runtime
- no intra-iteration dependences
- i loop is a reduction, therefore parallelizable

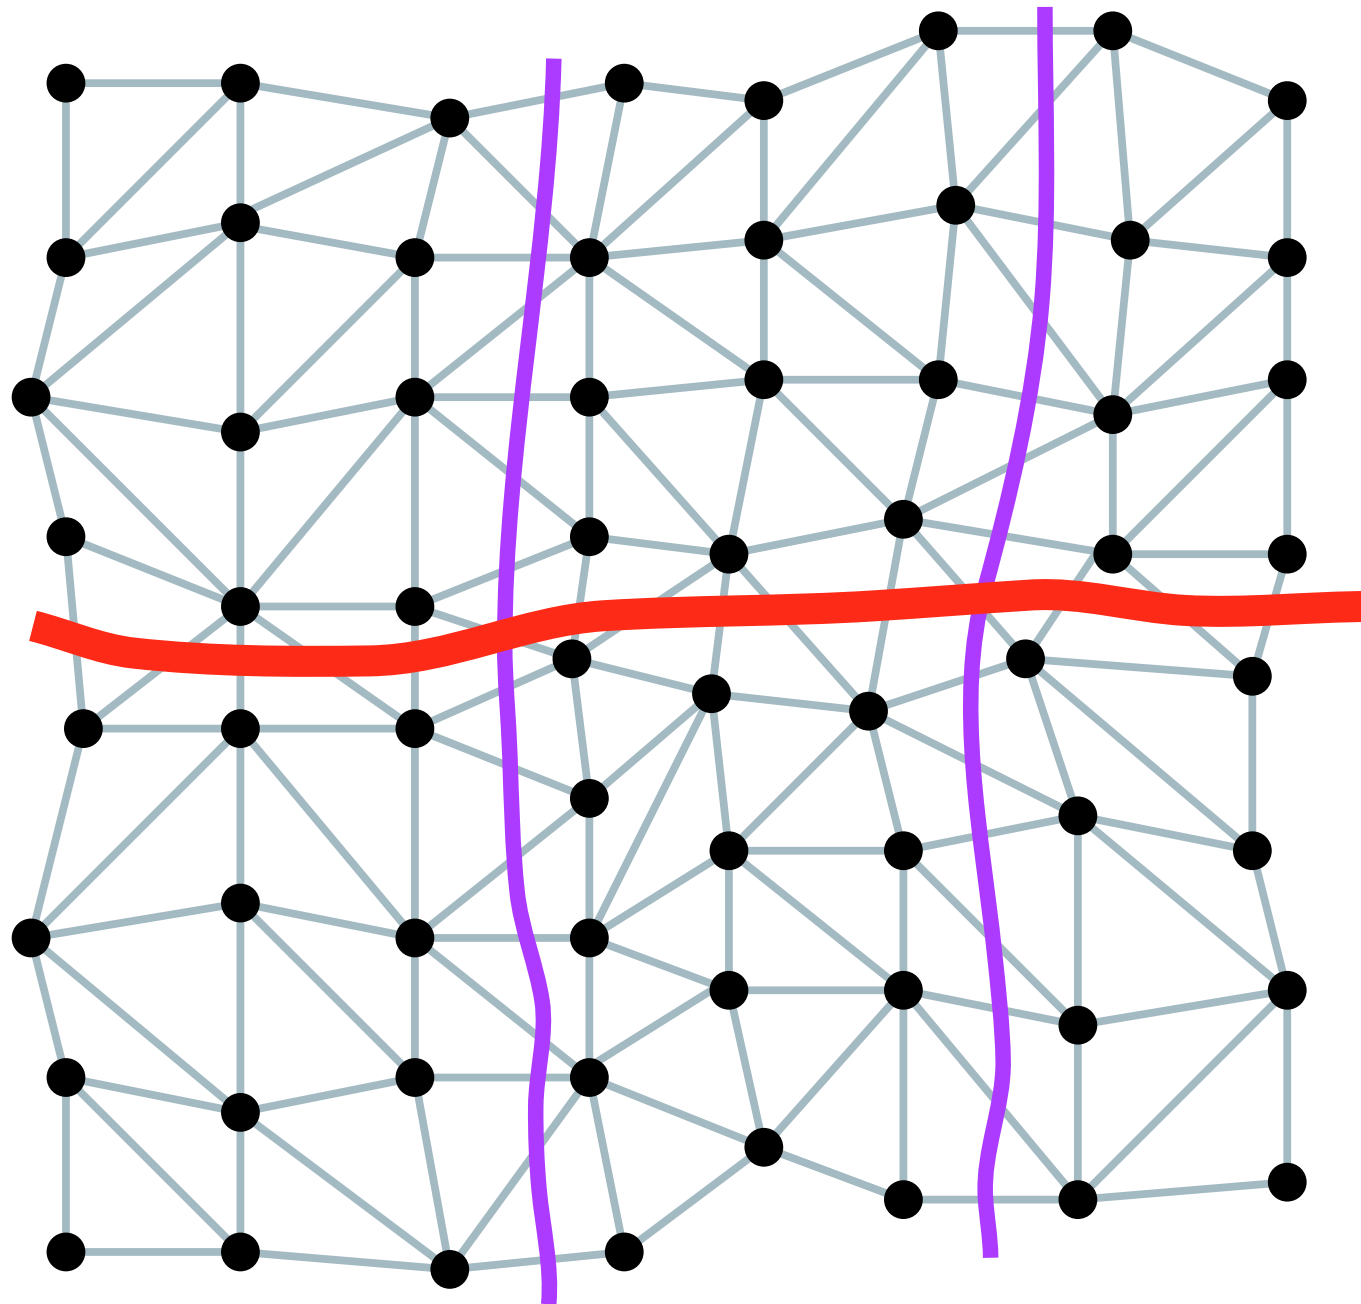
# Parallelize with Owner Computes Method



At each convergence iteration ...

- each partition receives data from previous convergence iteration
- each partition executes in parallel
- each partition sends data

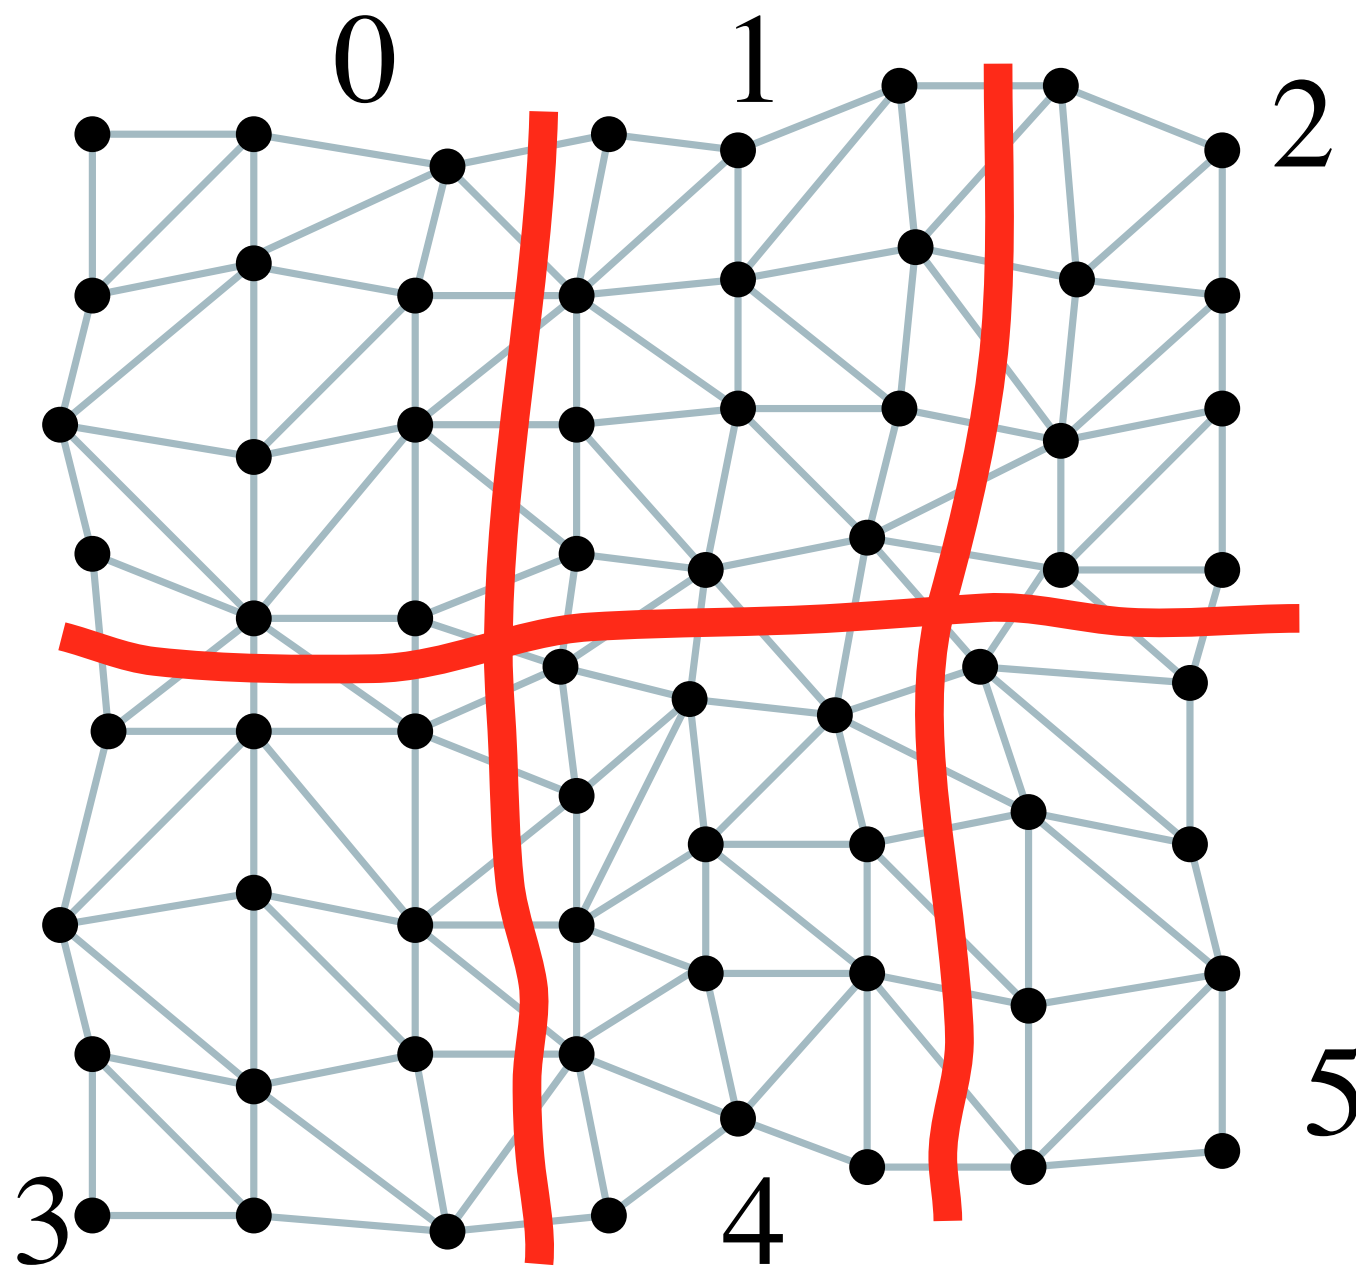
# Locality in Owner Computes Method



- Intra-iteration locality
  - schedule by sub-part
  - reorder data for consecutive order in sub-parts
- NO Inter-iteration locality, main partitions don't fit in cache

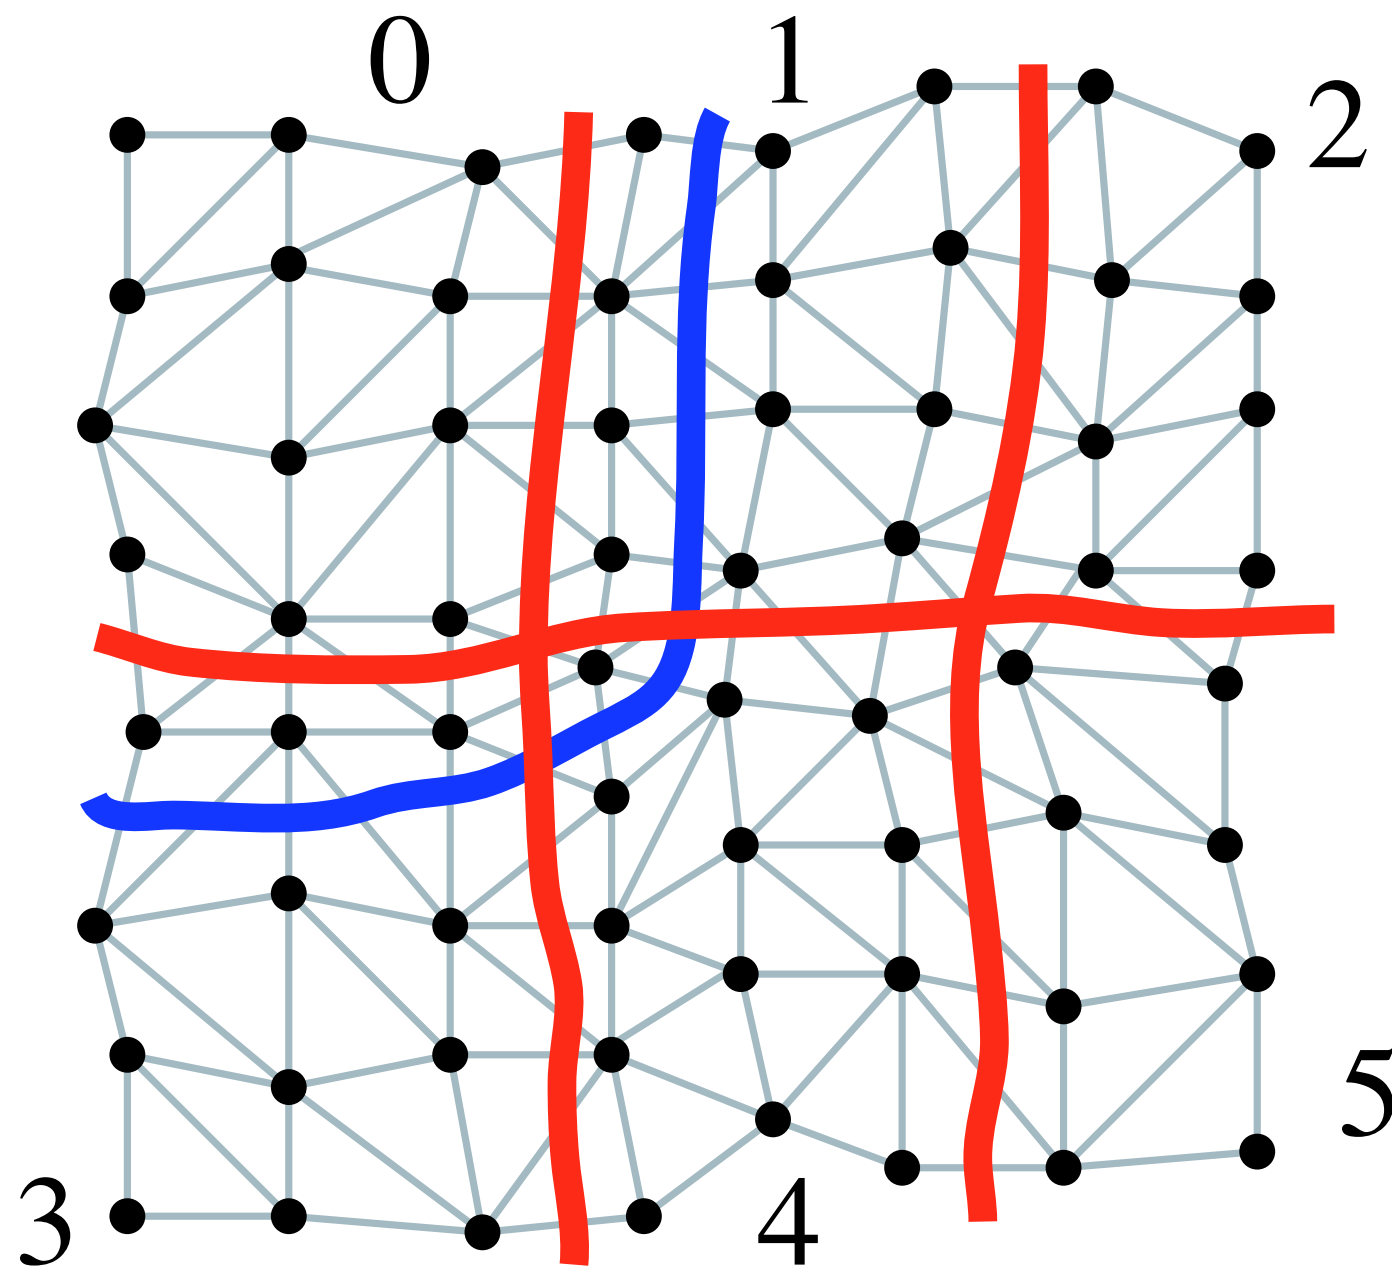


# Parallelize with Full Sparse Tiling



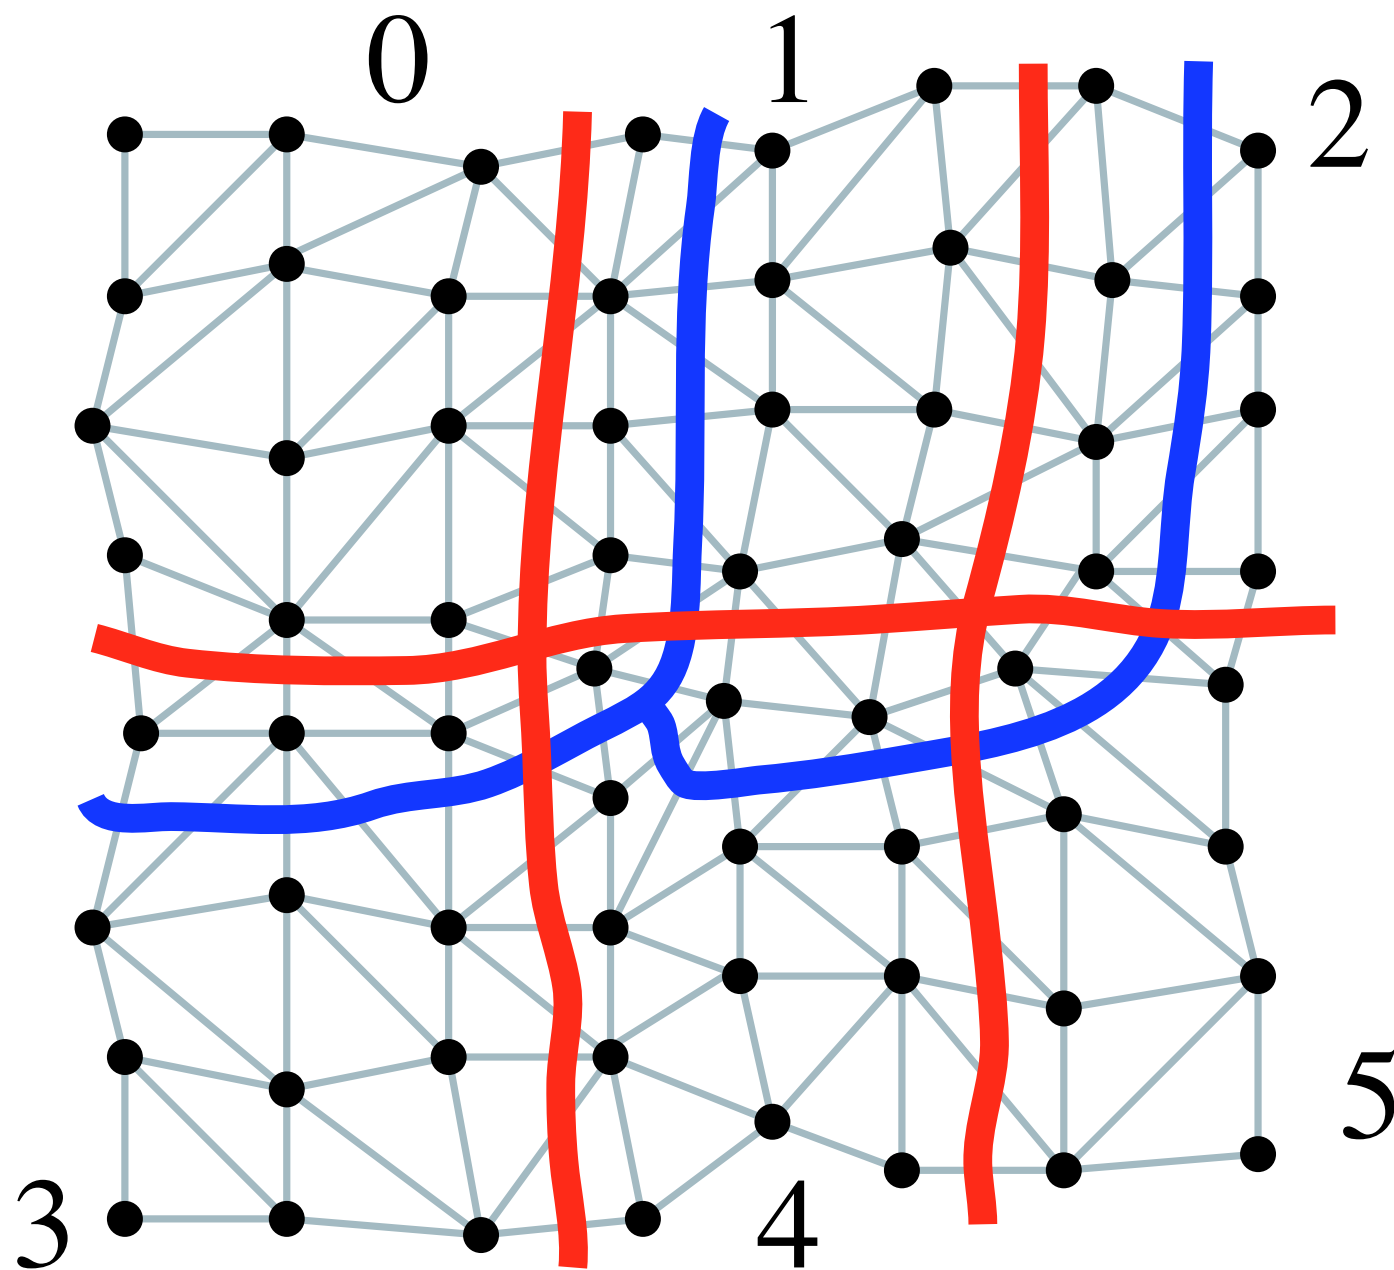
- Create seed partitioning at  $iter = 2$

# Parallelize with Full Sparse Tiling



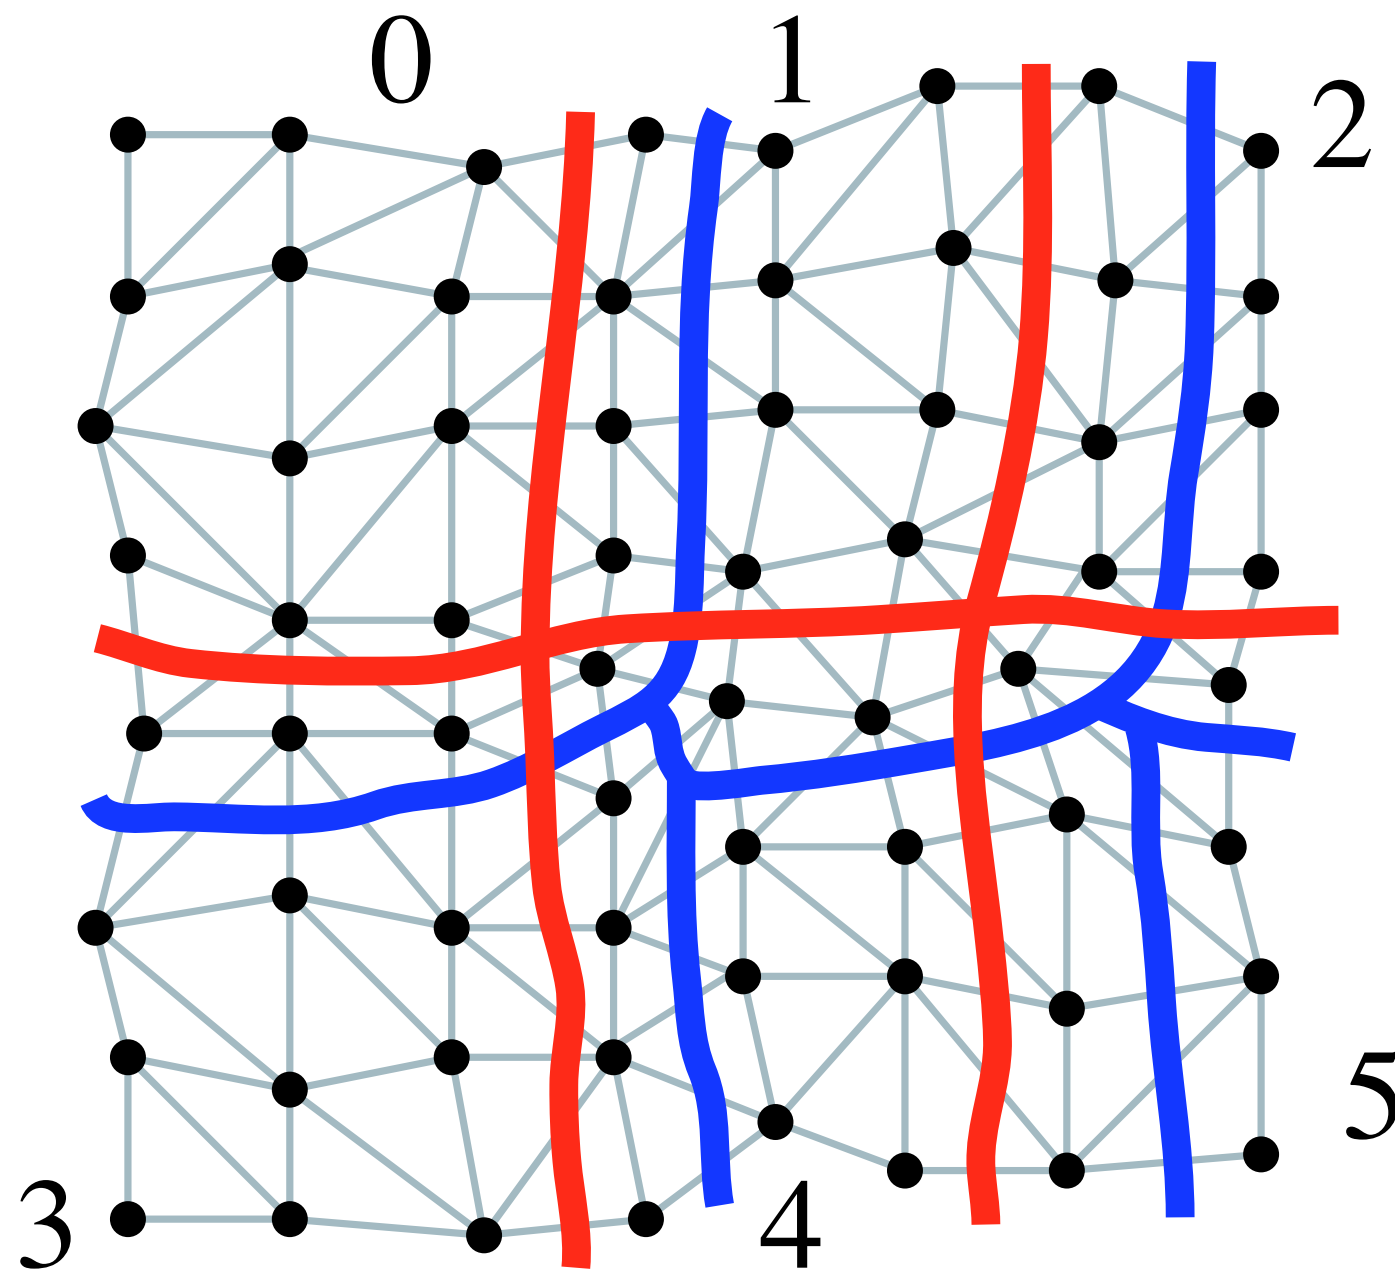
- Create seed partitioning at *iter* = 2
- Grow tiles to *iter* = 1 based on ordering of partitions

# Parallelize with Full Sparse Tiling



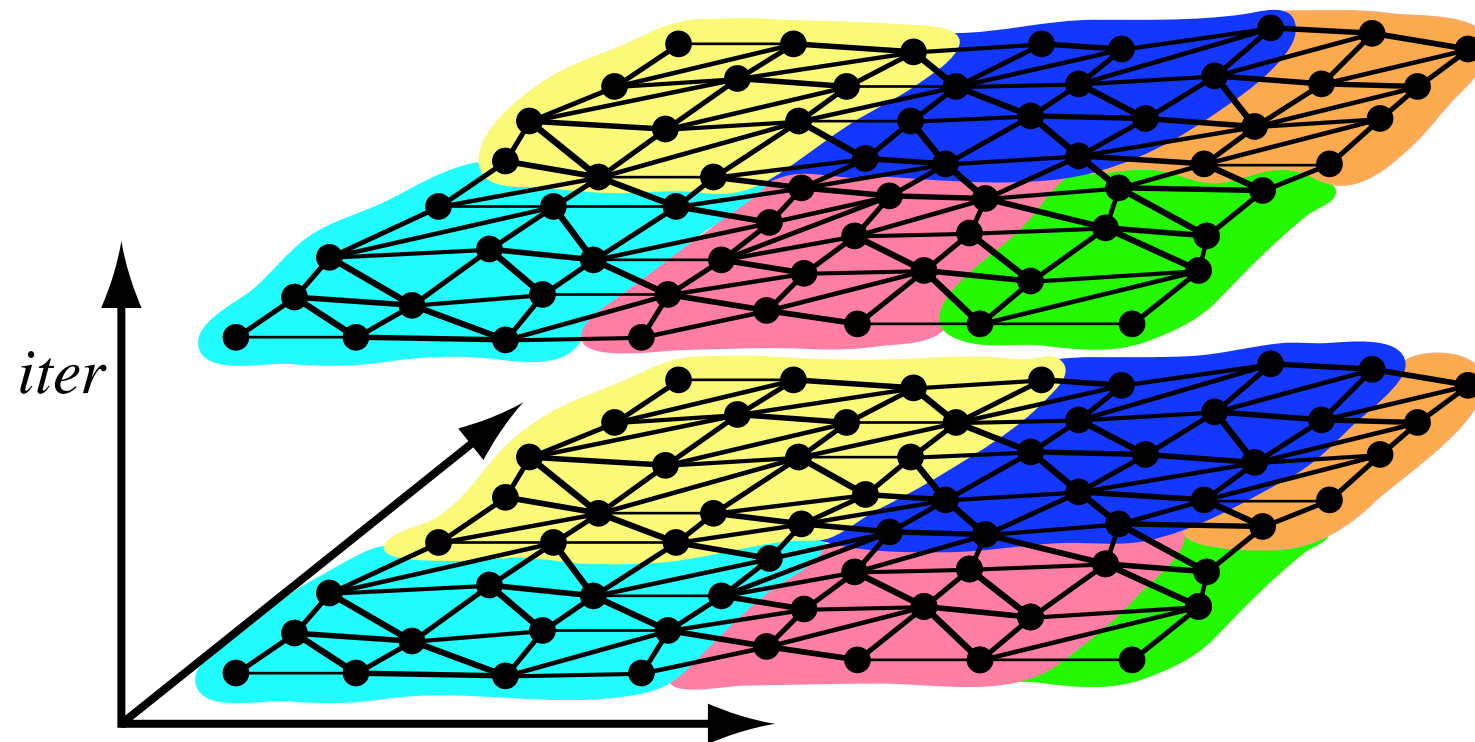
- Create seed partitioning at *iter* = 2
- Grow tiles to *iter* = 1 based on ordering of partitions

# Parallelize with Full Sparse Tiling



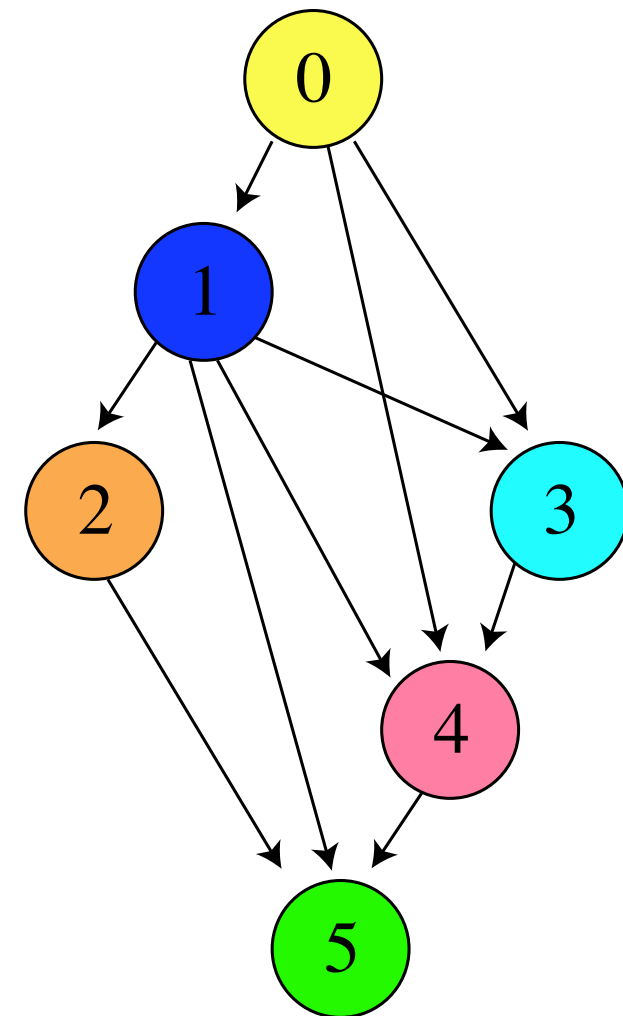
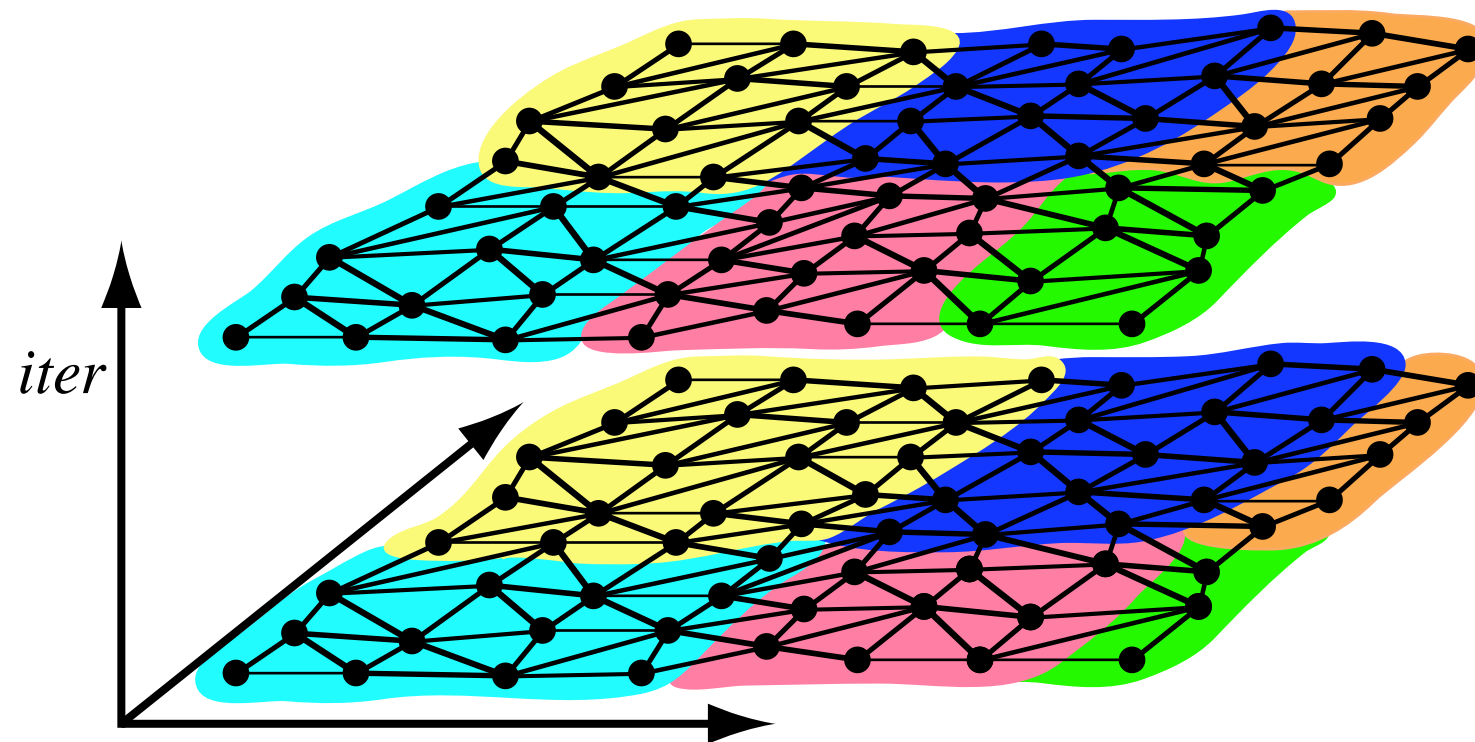
- Create seed partitioning at *iter* = 2
- Grow tiles to *iter* = 1 based on ordering of partitions

# Locality in Full Sparse Tiled Jacobi



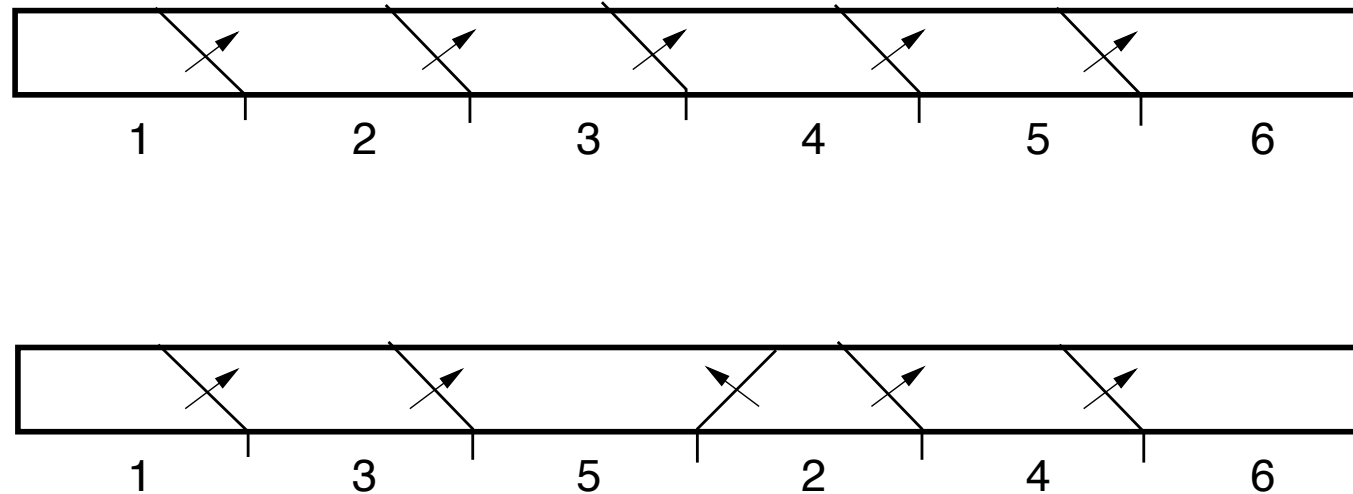
- **Intra-iteration**  
locality achieved by consecutively ordering matrix graph nodes within a seed partition
- **Inter-iteration**  
locality achieved with tile-by-tile execution

# Parallelism in Full Sparse Tiled Jacobi



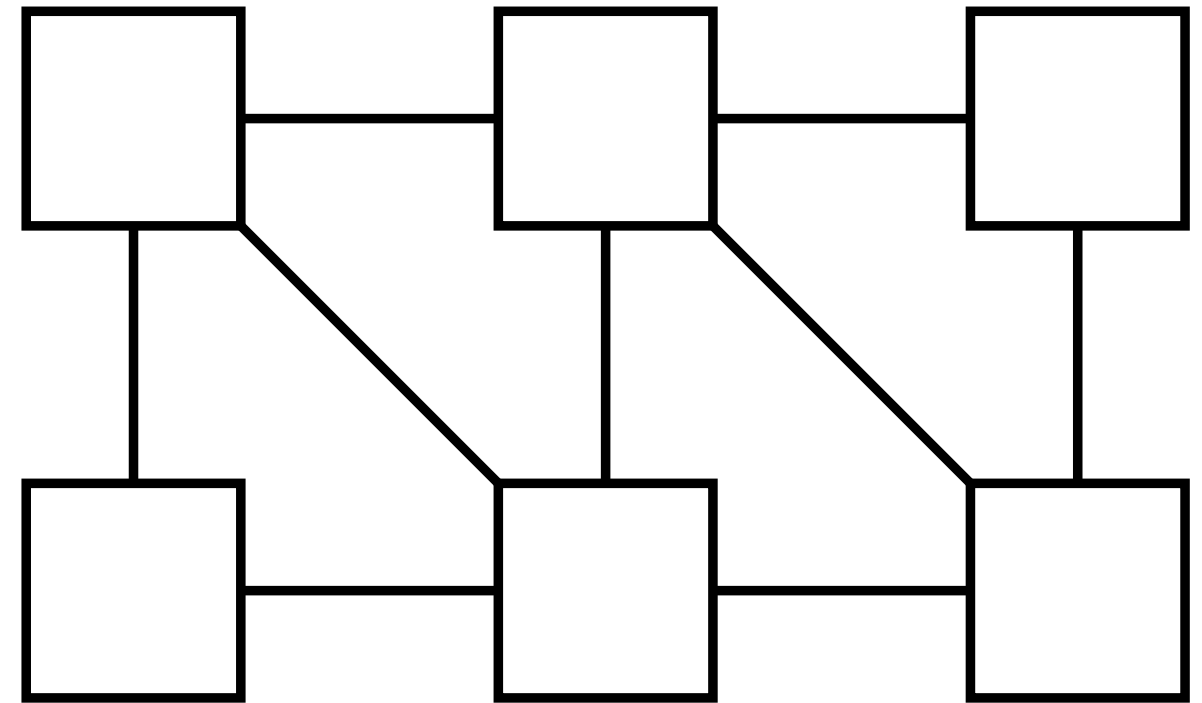
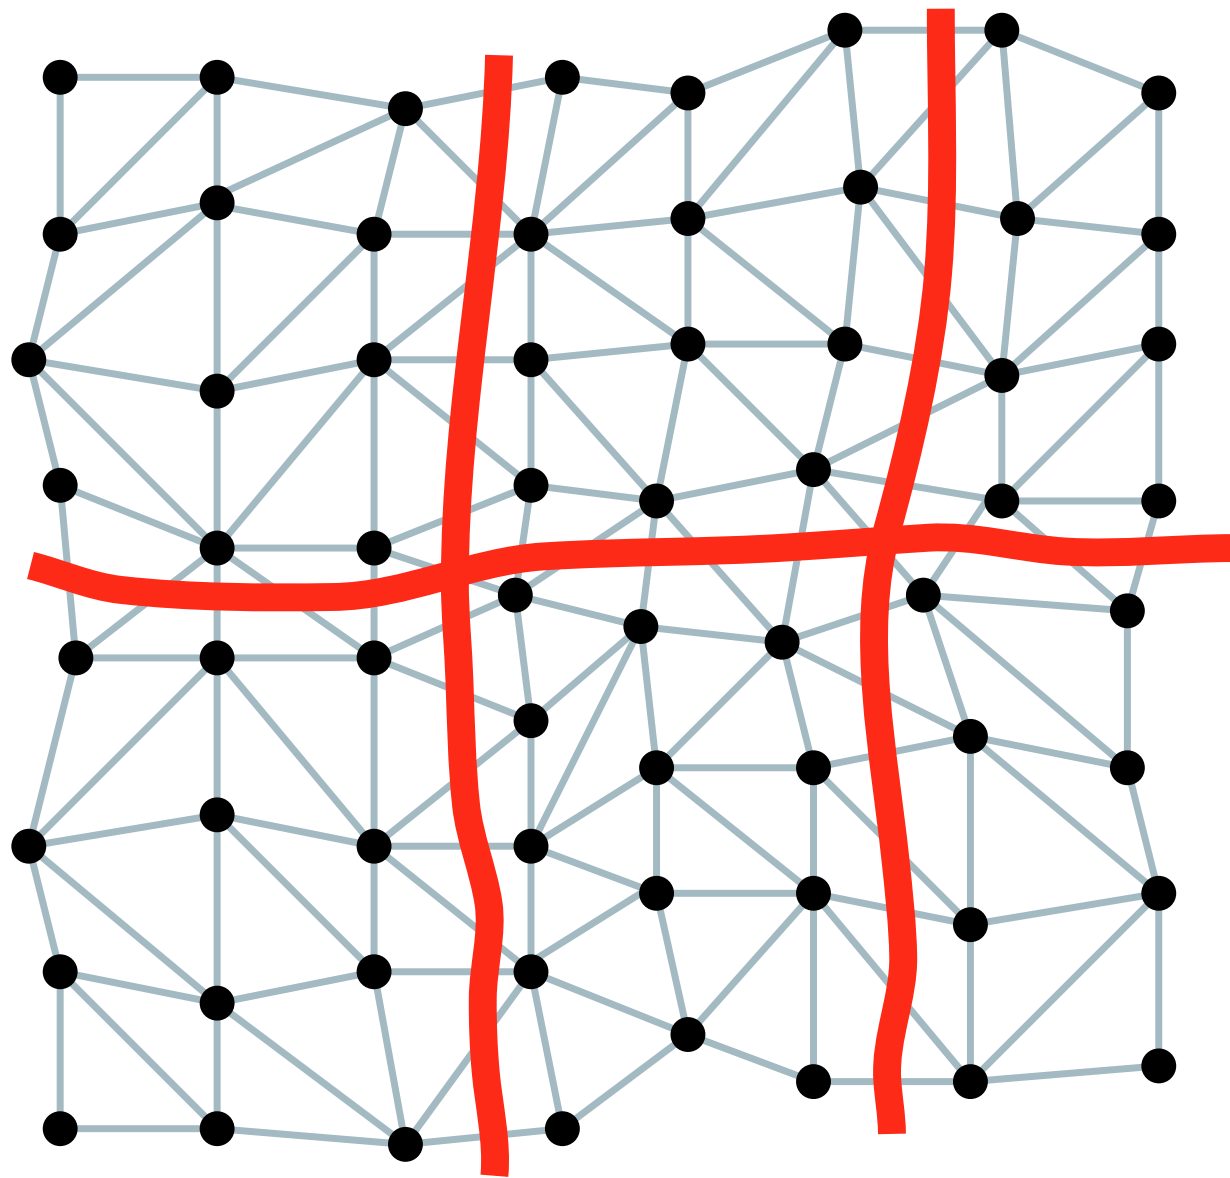
$$\begin{aligned}\text{Average parallelism} &= (\# \text{ tiles}) / (\# \text{ tiles in critical path}) \\ &= 6 / 5 = 1.2\end{aligned}$$

# How can we increase parallelism between tiles?



- Order that tile growth is performed matters
- Best is to first grow tiles whose seed partitions are not adjacent

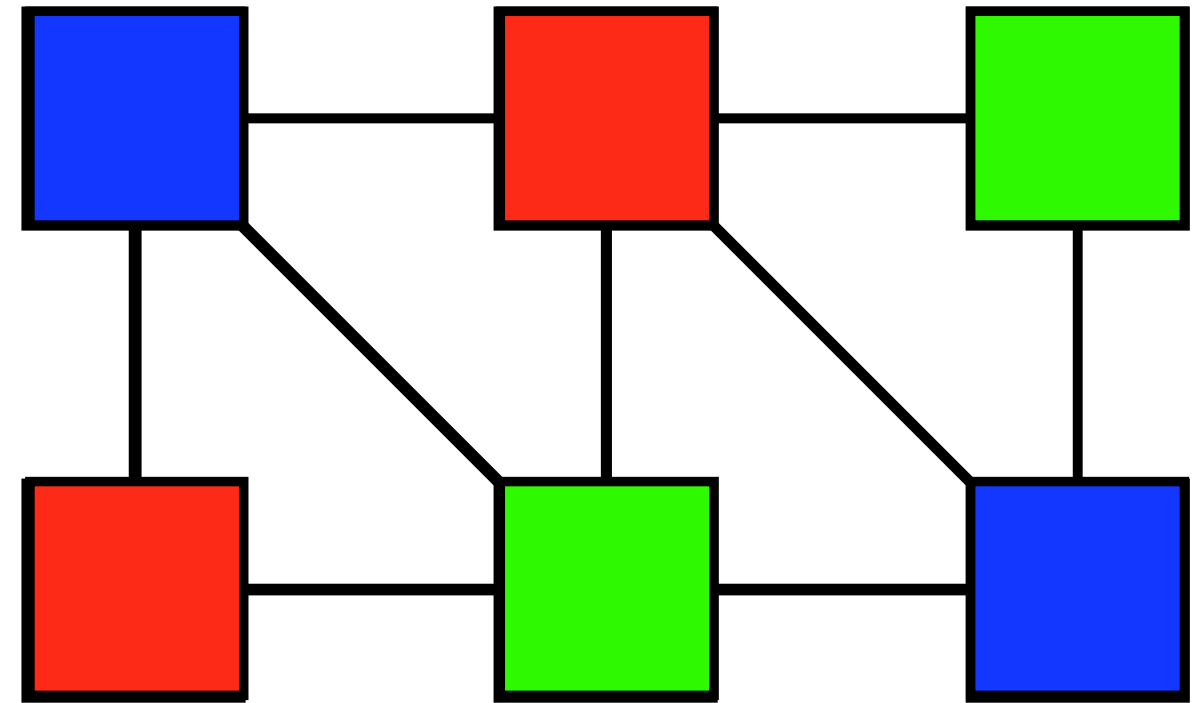
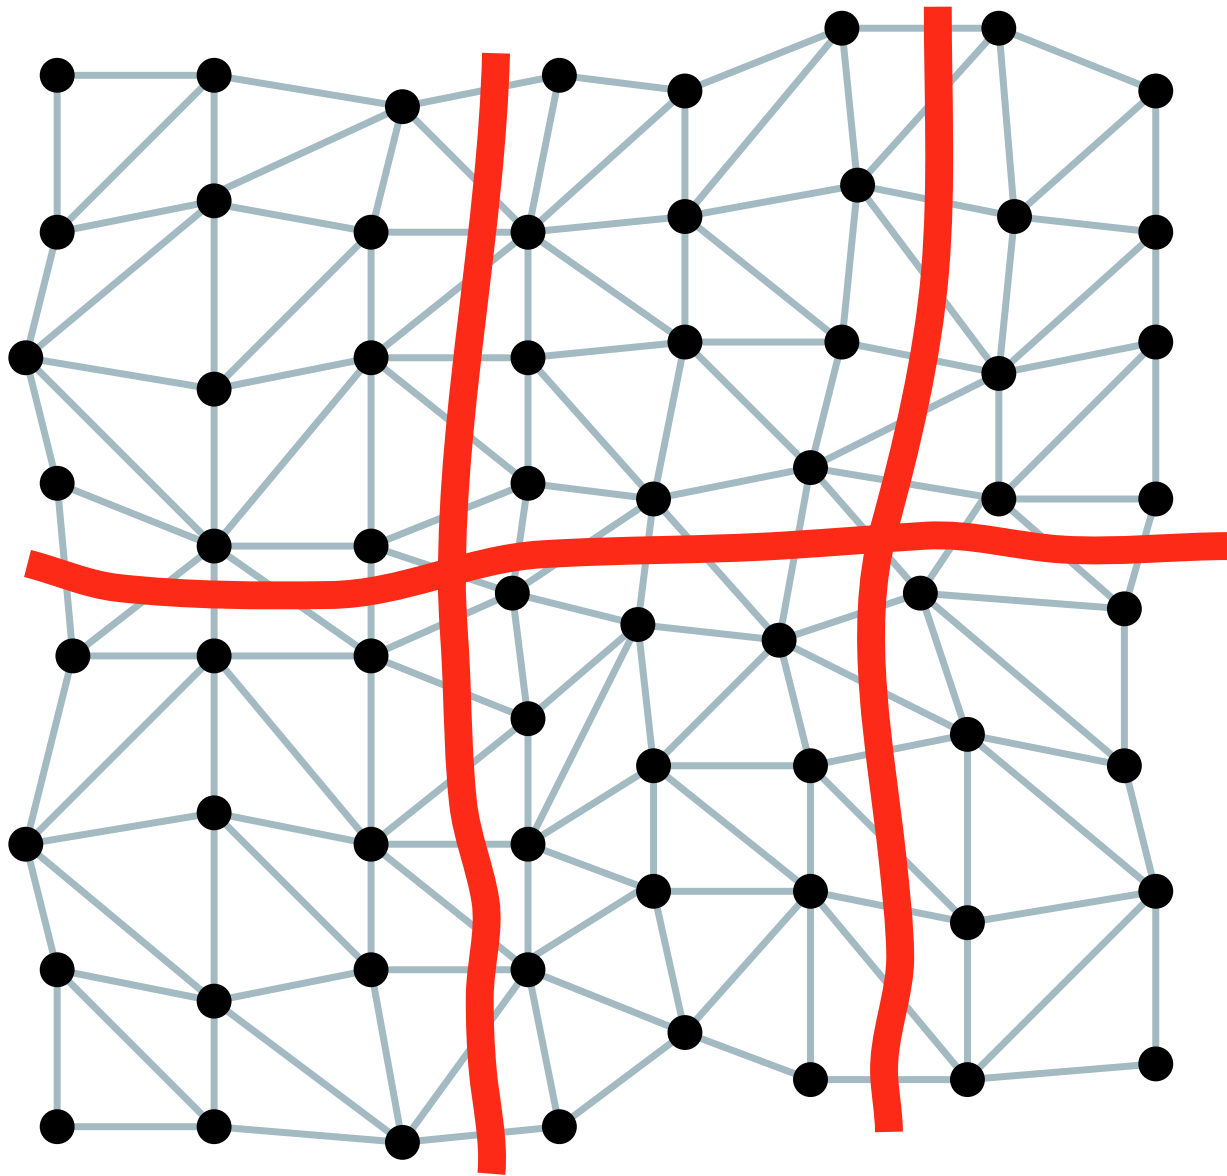
# Improving Average Parallelism Using Coloring



- Create a partition graph

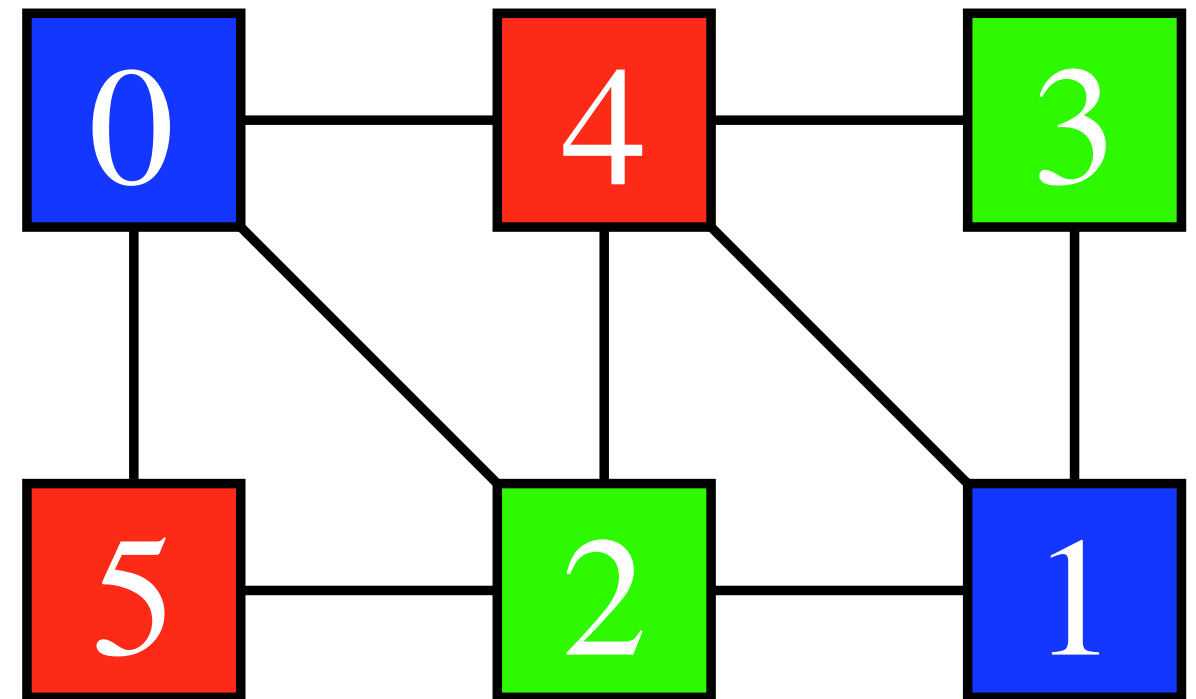
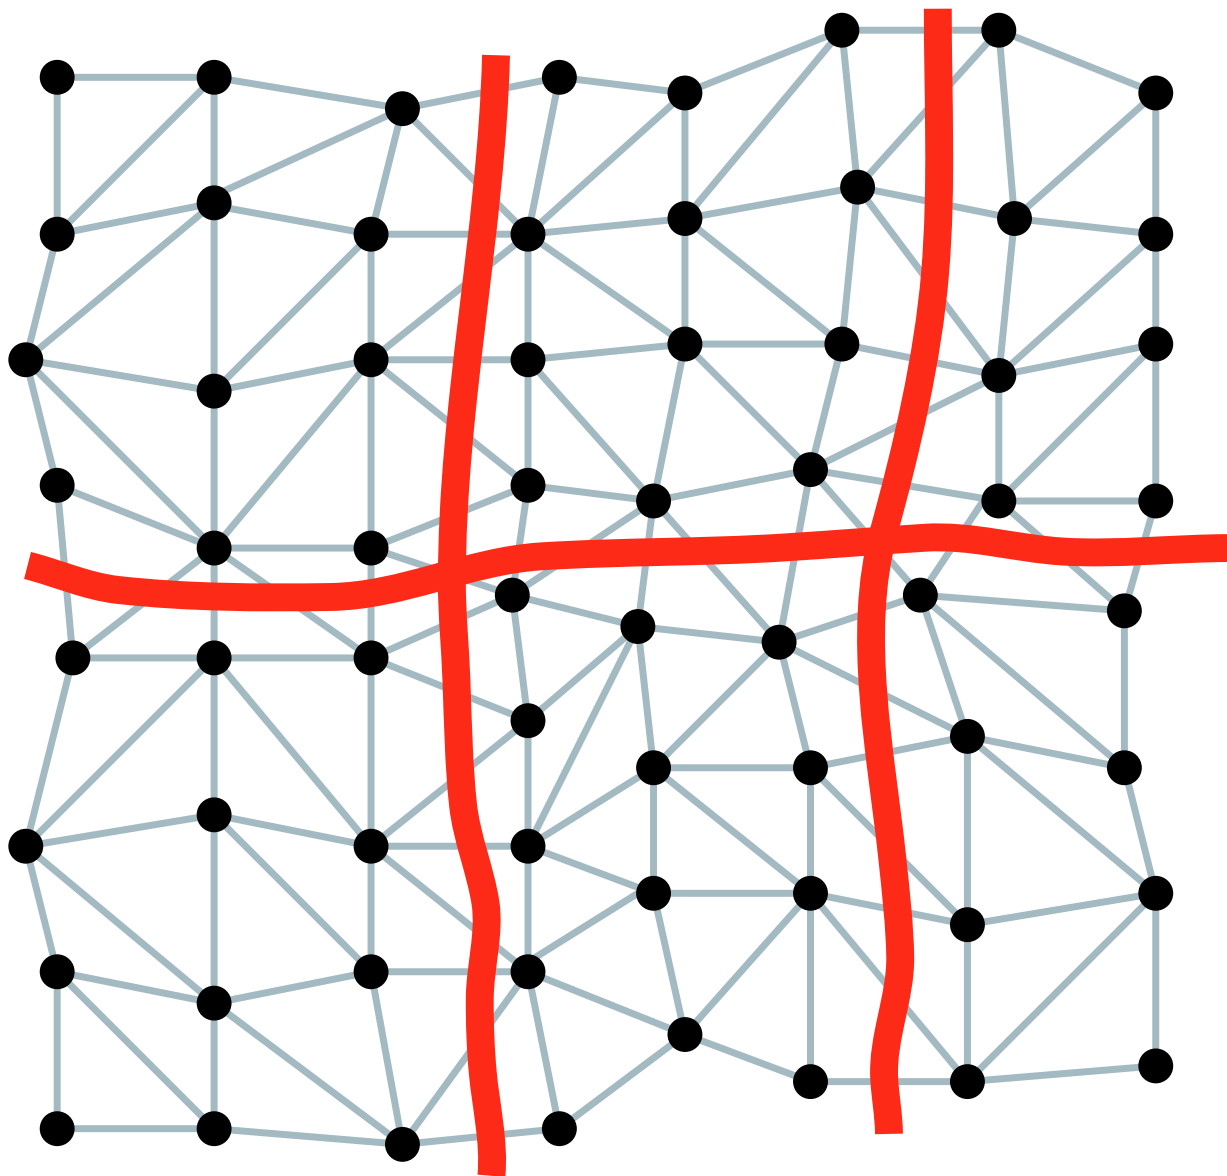


# Improving Average Parallelism Using Coloring



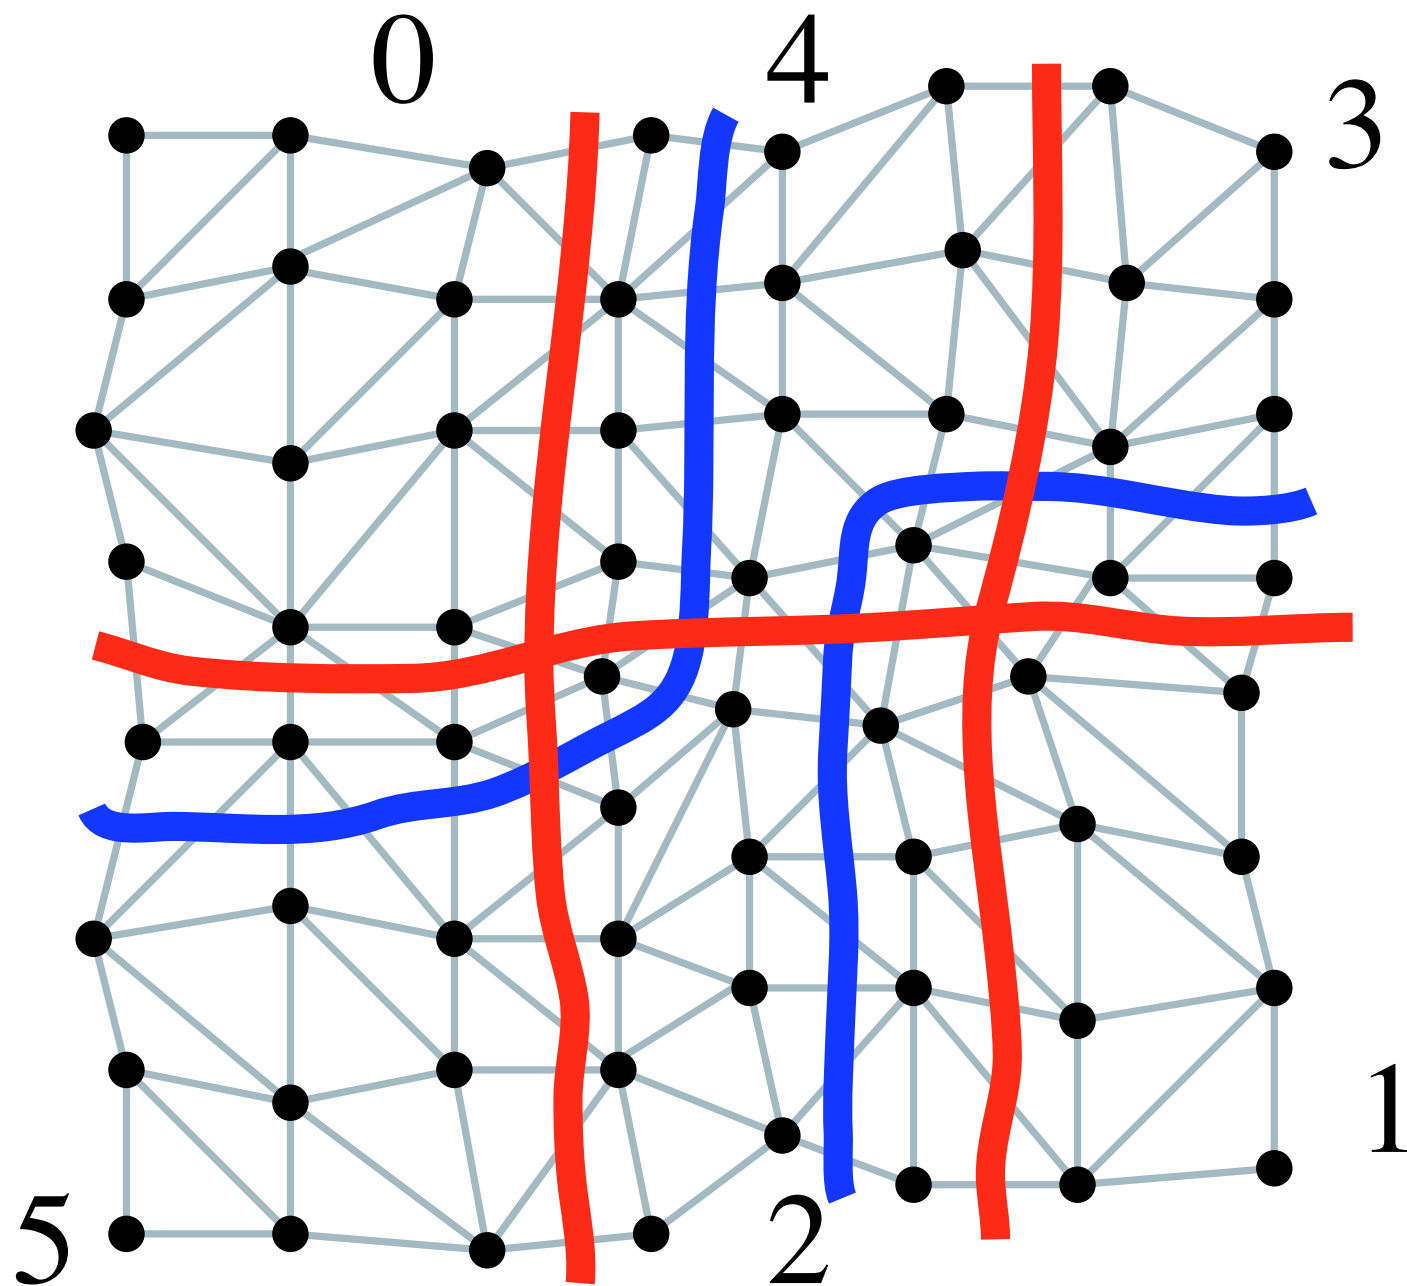
- Create a partition graph
- Color the partition graph

# Improving Average Parallelism Using Coloring



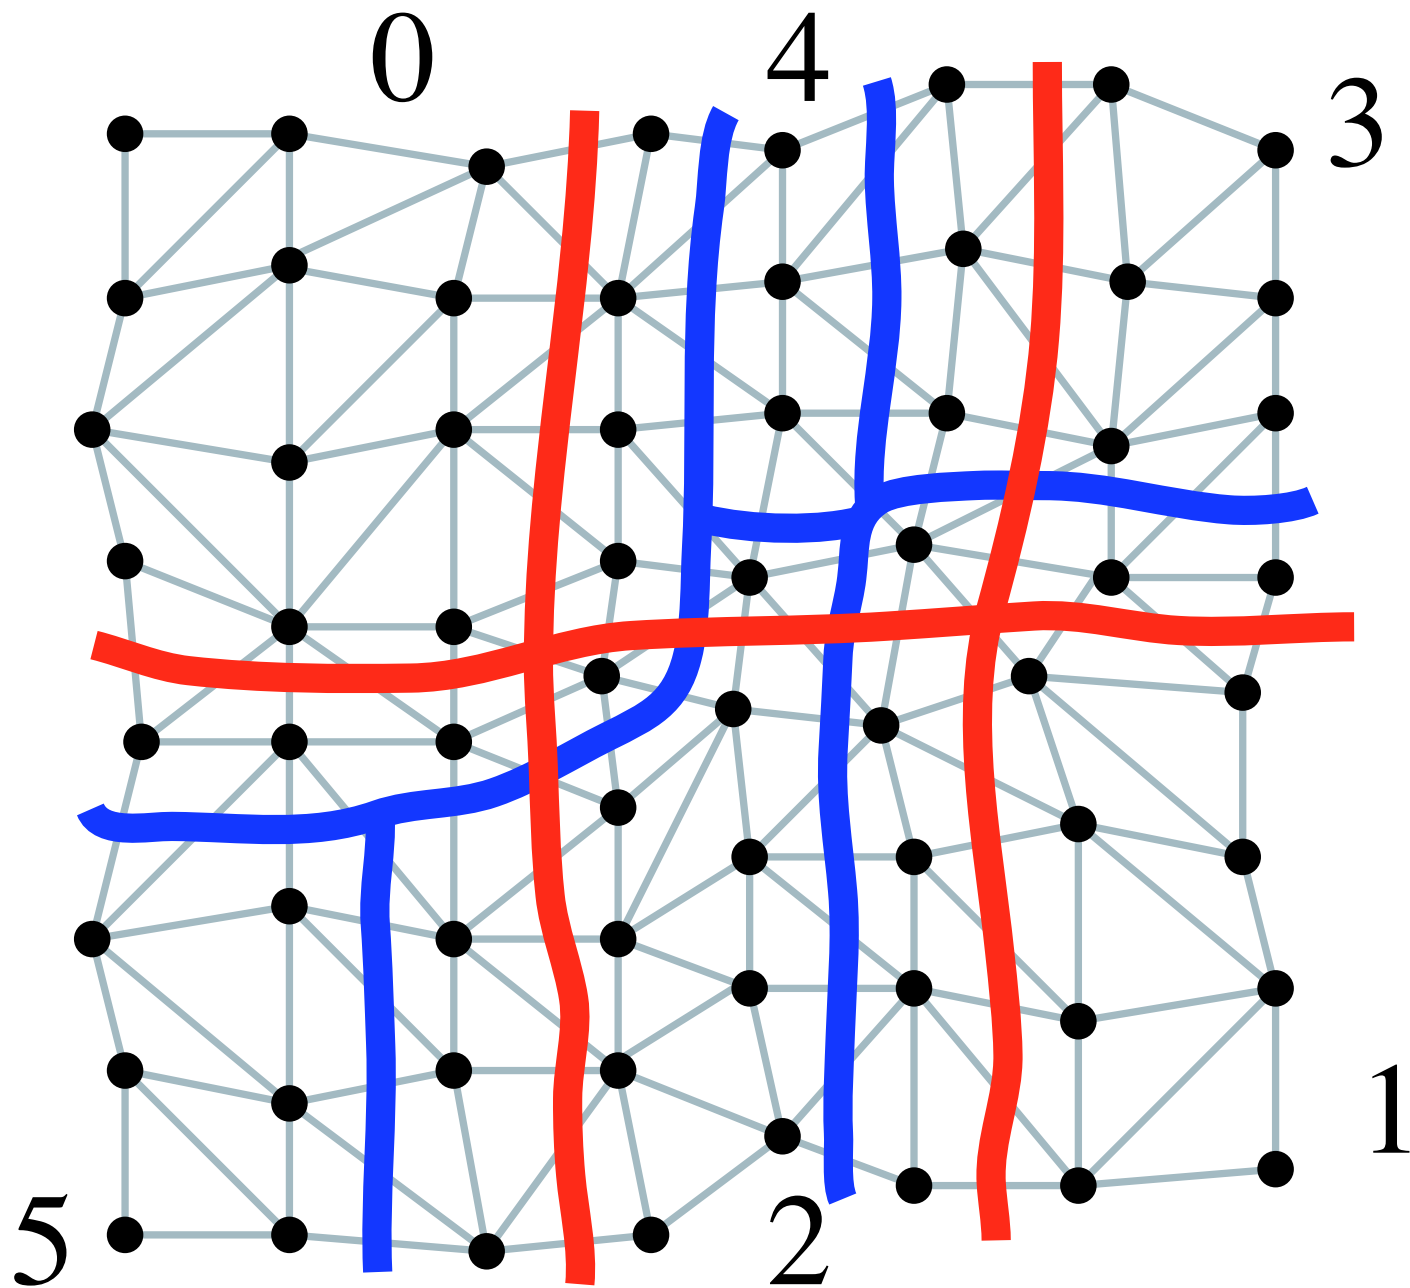
- Create a partition graph
- Color the partition graph
- Renumber partitions consecutively by color

# Grow Using New Partition Order



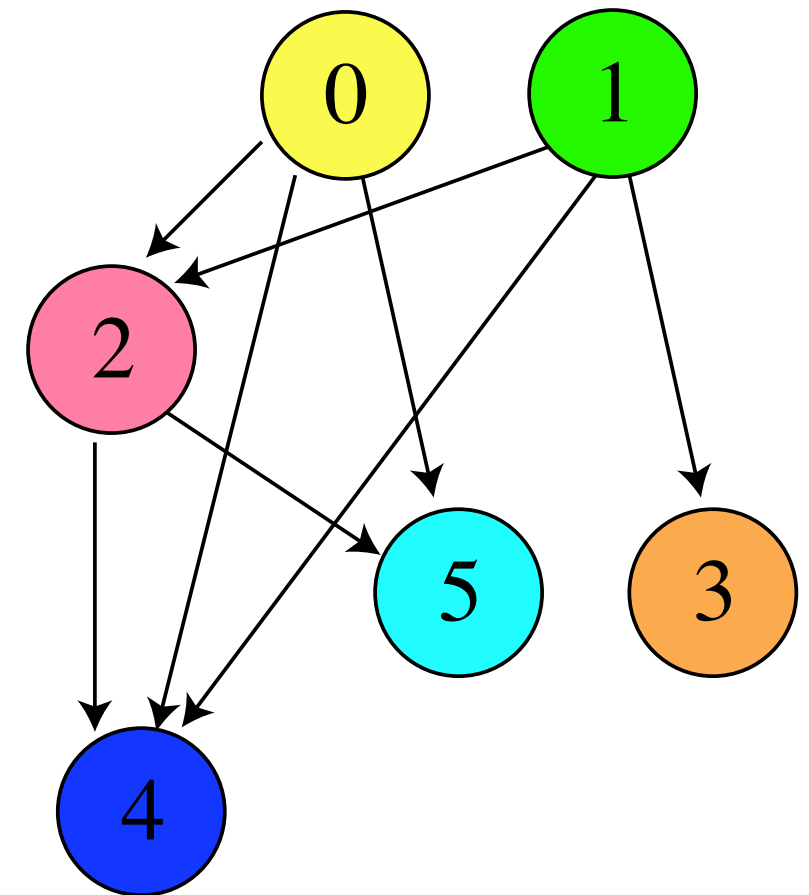
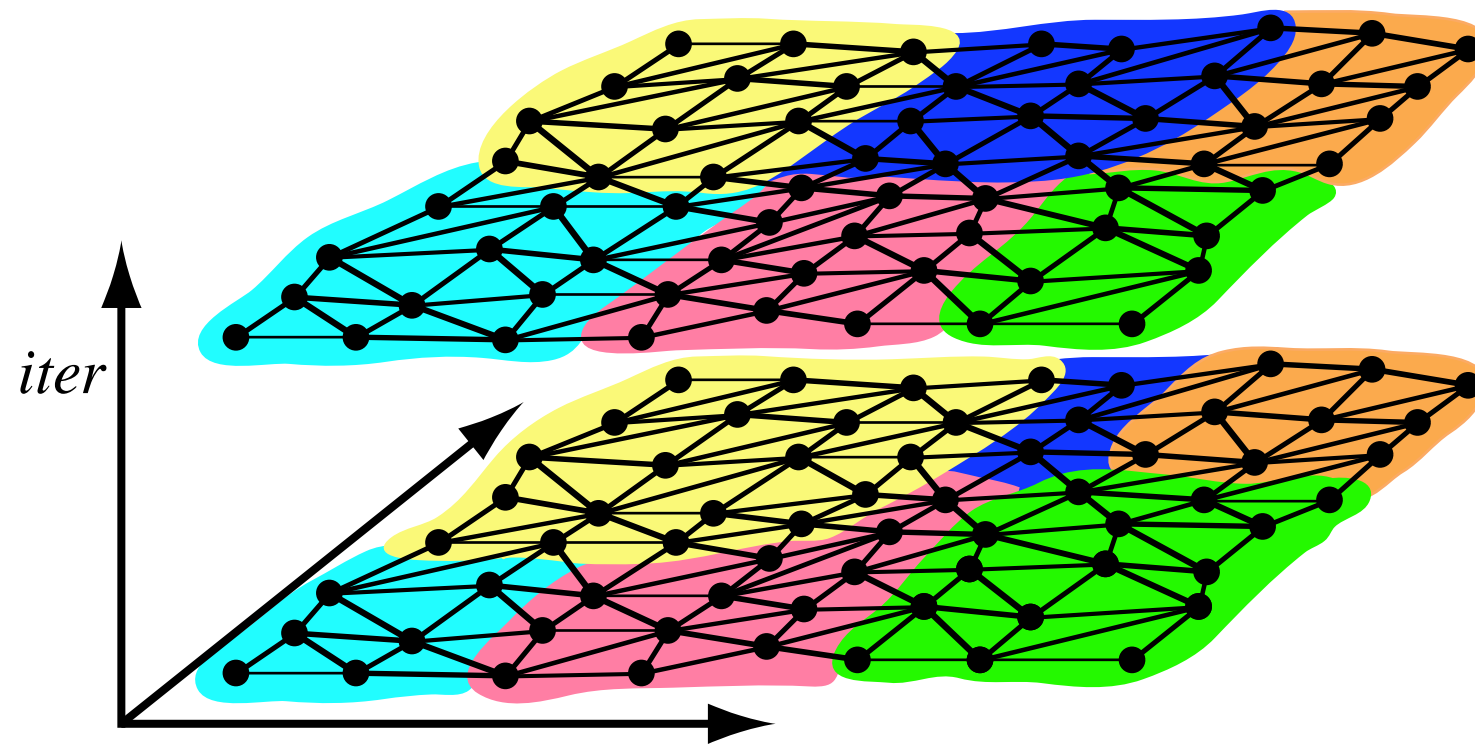
- Renumber the seed partition cells based on coloring
- Grow tiles using new ordering
- Notice that tiles 0 and 1 may be executed in parallel

# Re-grow Using New Partition Order



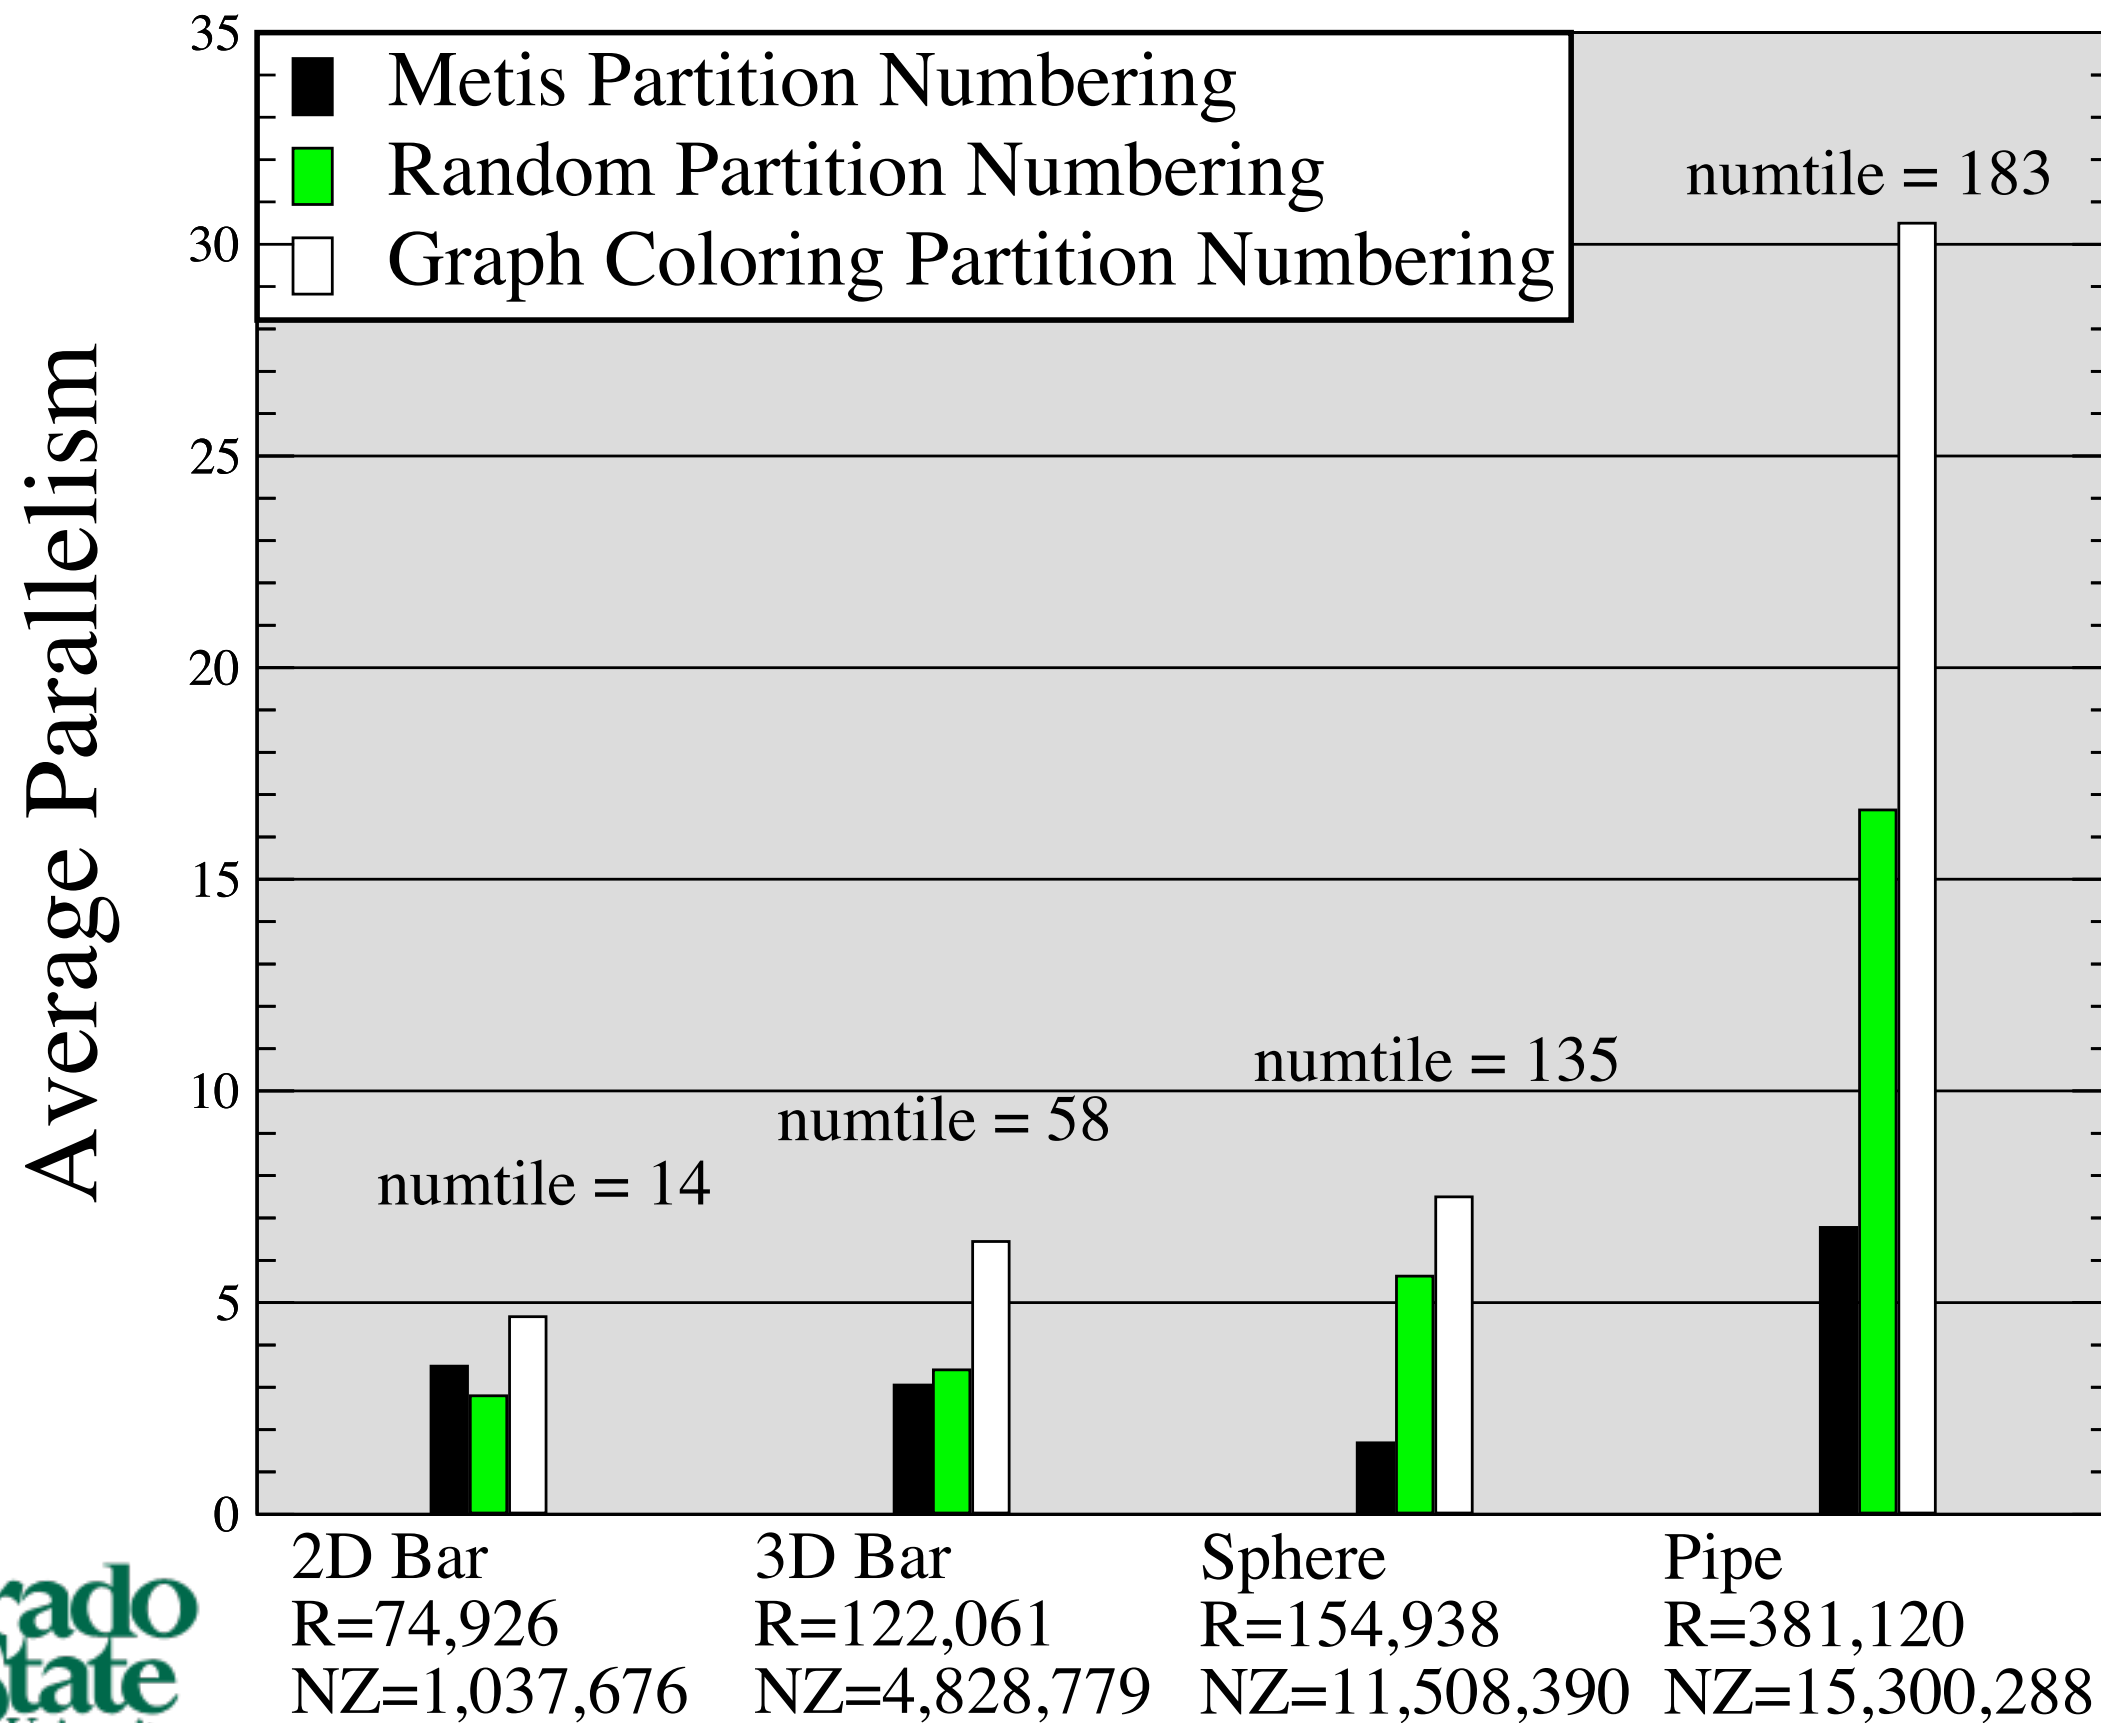
- Renumber the seed partition cells based on coloring
- Grow tiles using new ordering
- Notice that tiles 0 and 1 may be executed in parallel
- Tiles 4 and 5 may also be executed in parallel

# Average Parallelism is Improved

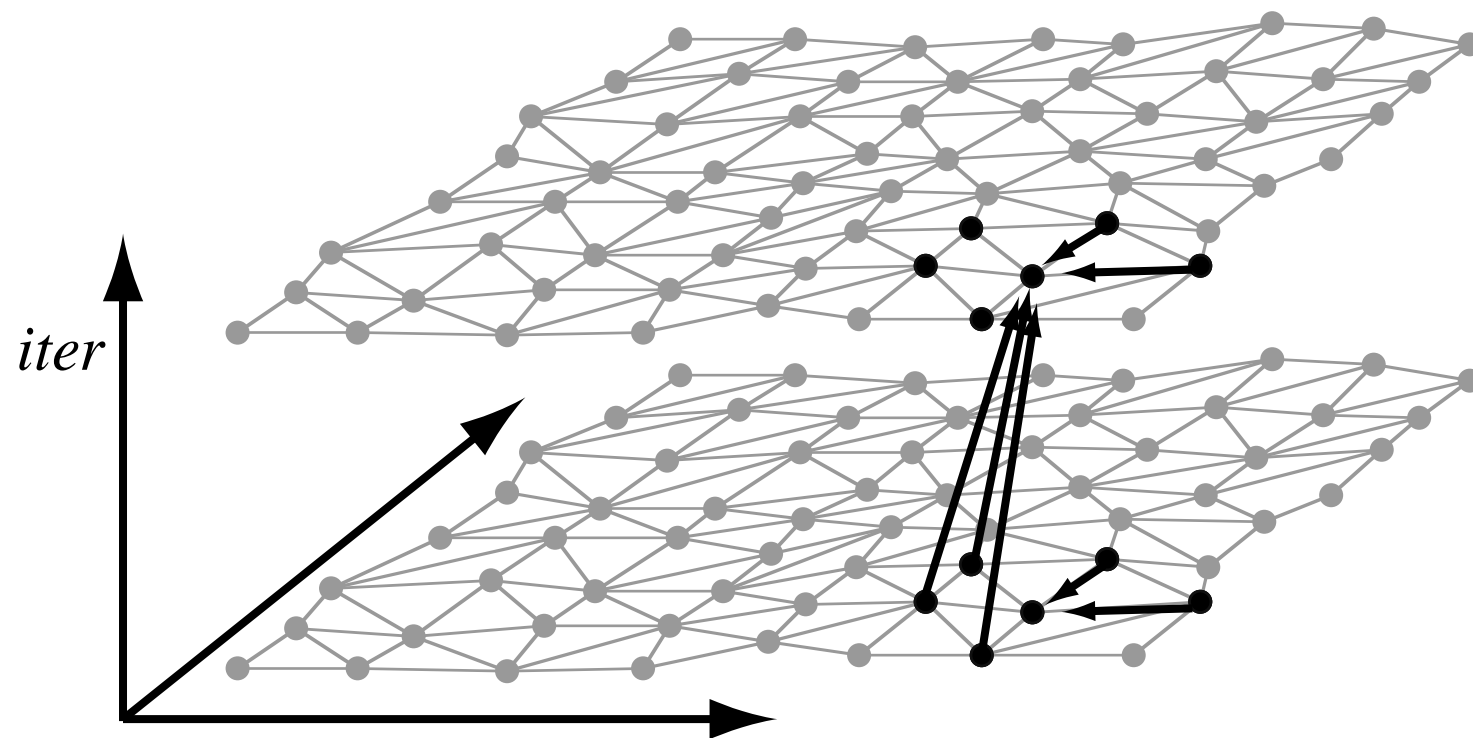


$$\begin{aligned}\text{Average parallelism} &= (\# \text{ tiles}) / (\# \text{ tiles in critical path}) \\ &= 6 / 3 = 2\end{aligned}$$

# Improvement with Real Matrix Graphs



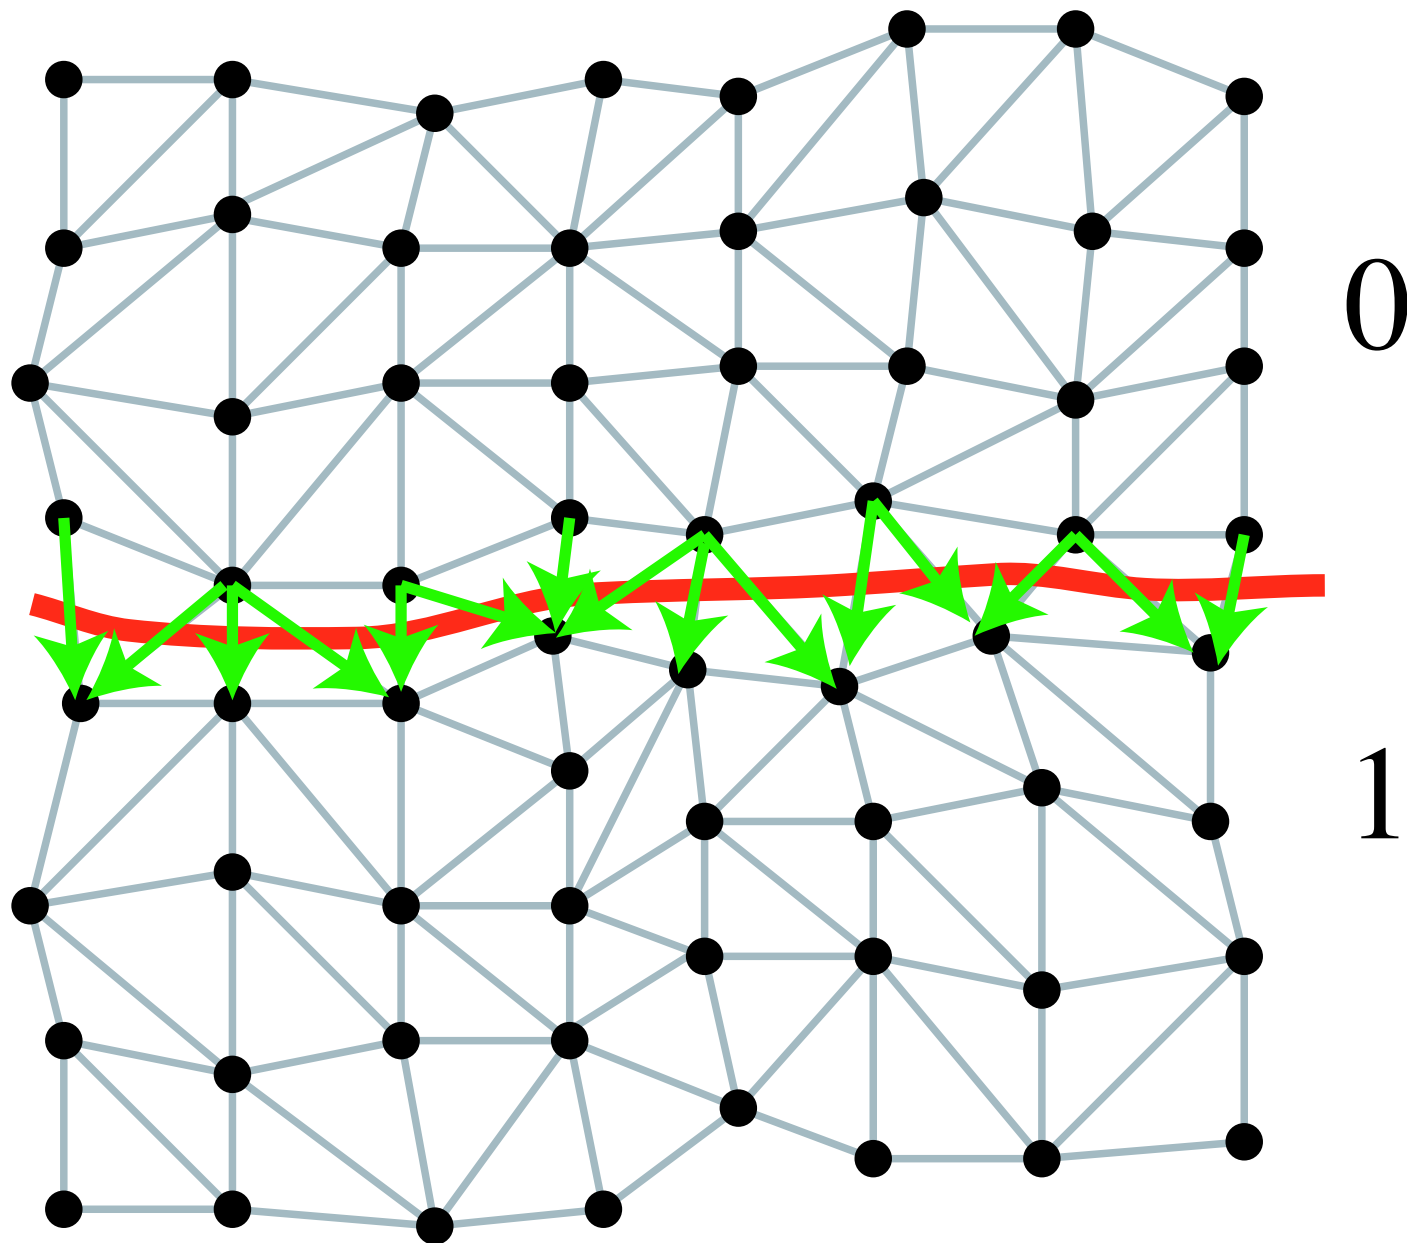
# Gauss-Seidel



- Loop carried dependences within convergence iteration as well as between them
- Dependences depend on the ordering of the nodes
- Nodes can be reordered apriori

# Owner Computes Method

## Nodal Gauss-Seidel [Adams 2001]

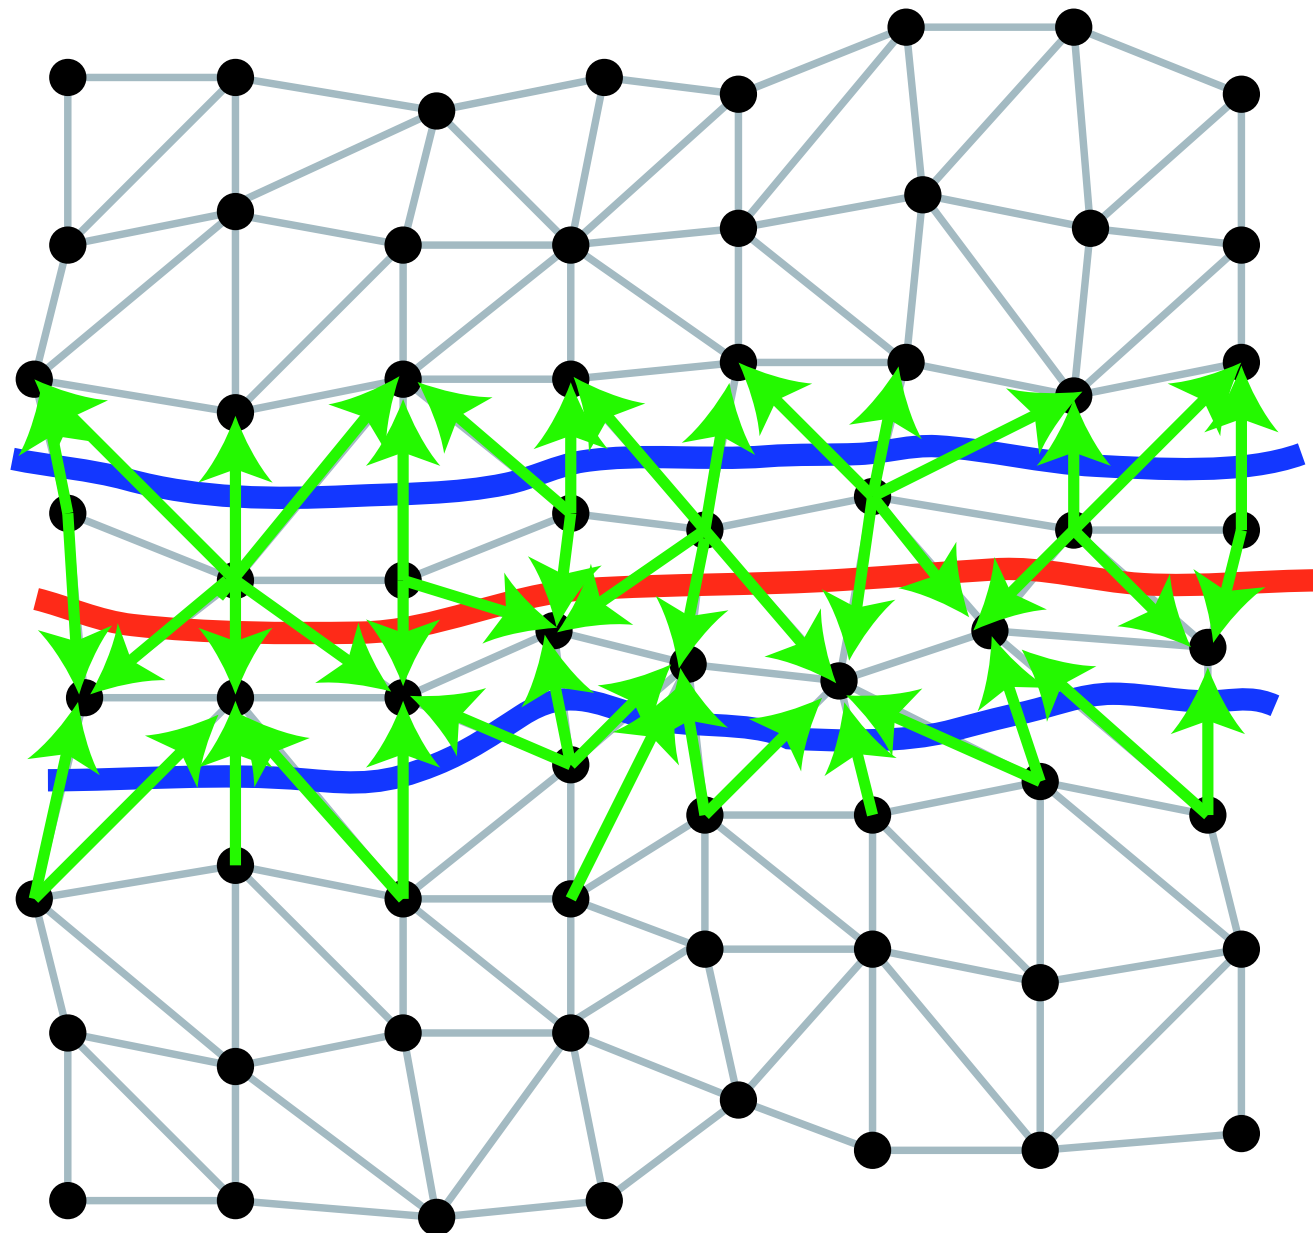


- Renumbers nodes (rows/columns in sparse matrix)
- Make data dependences between main partitions consistent



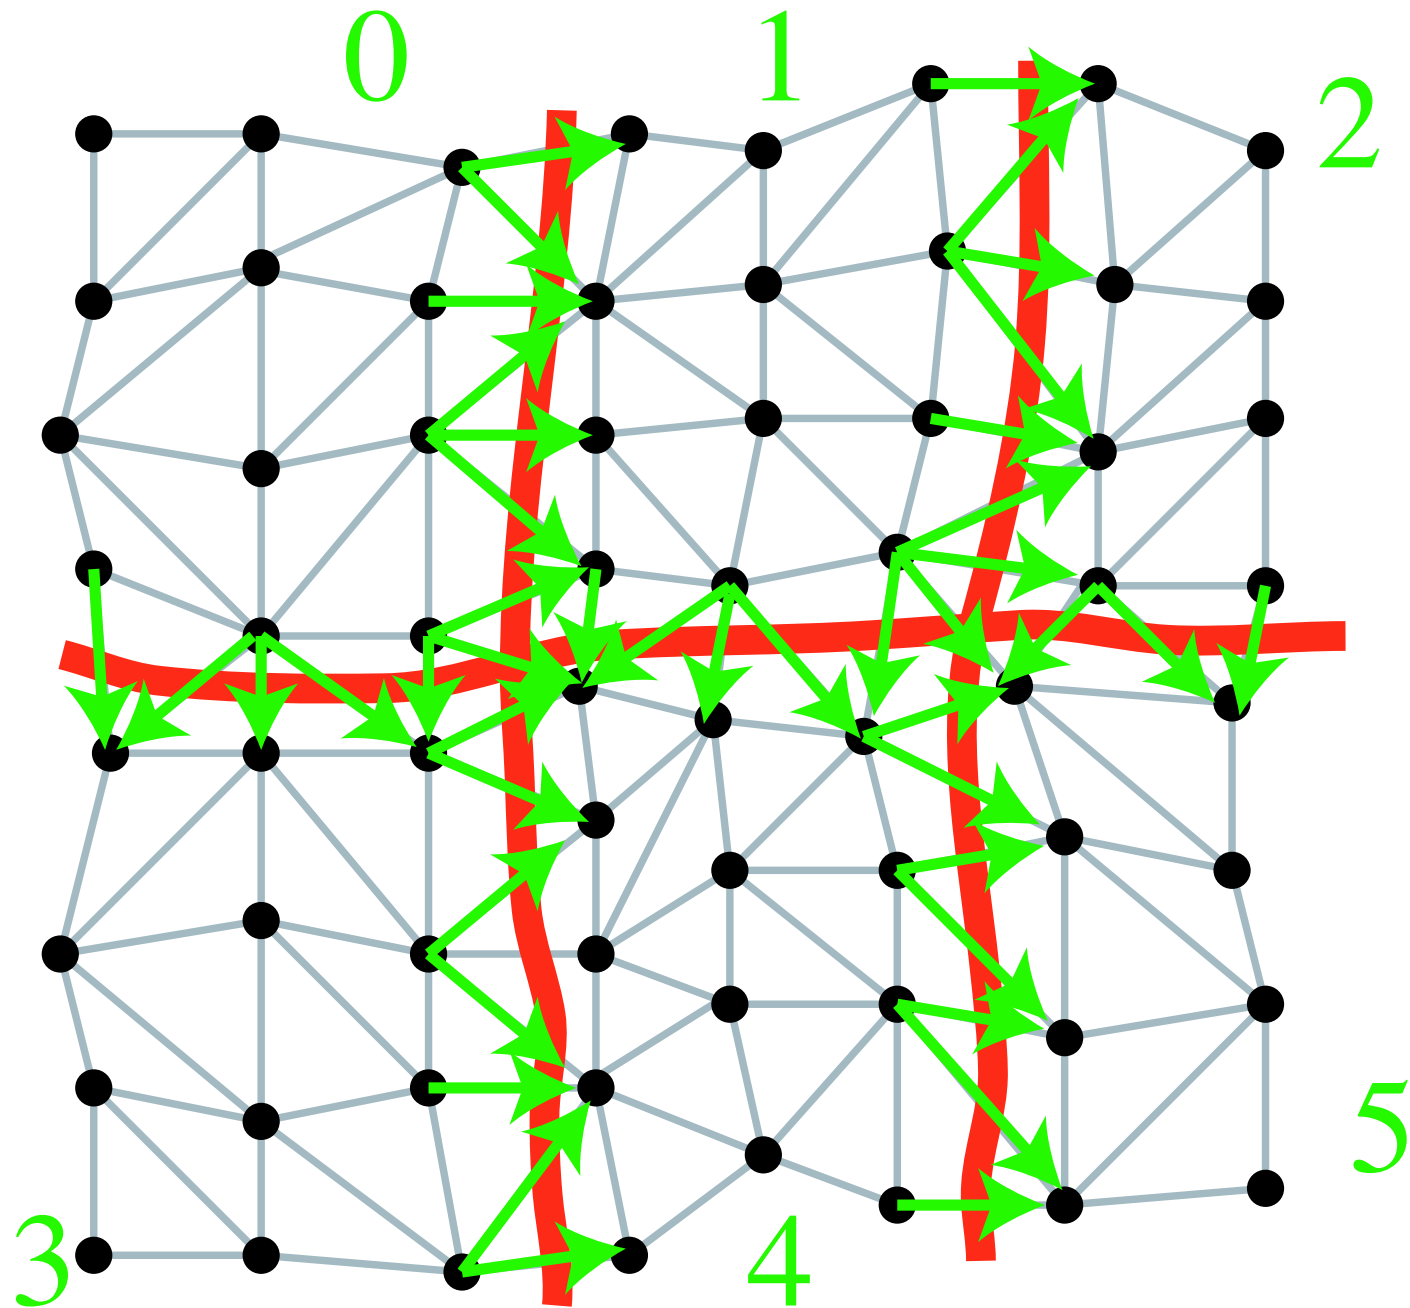
# Owner Computes Method

## Nodal Gauss-Seidel [Adams 2001]



- Renumbers nodes (rows/columns in sparse matrix)
- Make data dependences between main partitions consistent
- Subpartition for better parallelism

# Full Sparse Tiled Gauss-Seidel



- Tile growth creates and maintains a partial ordering between nodes in matrix graph
- Reorder nodes so that partial ordering is maintained

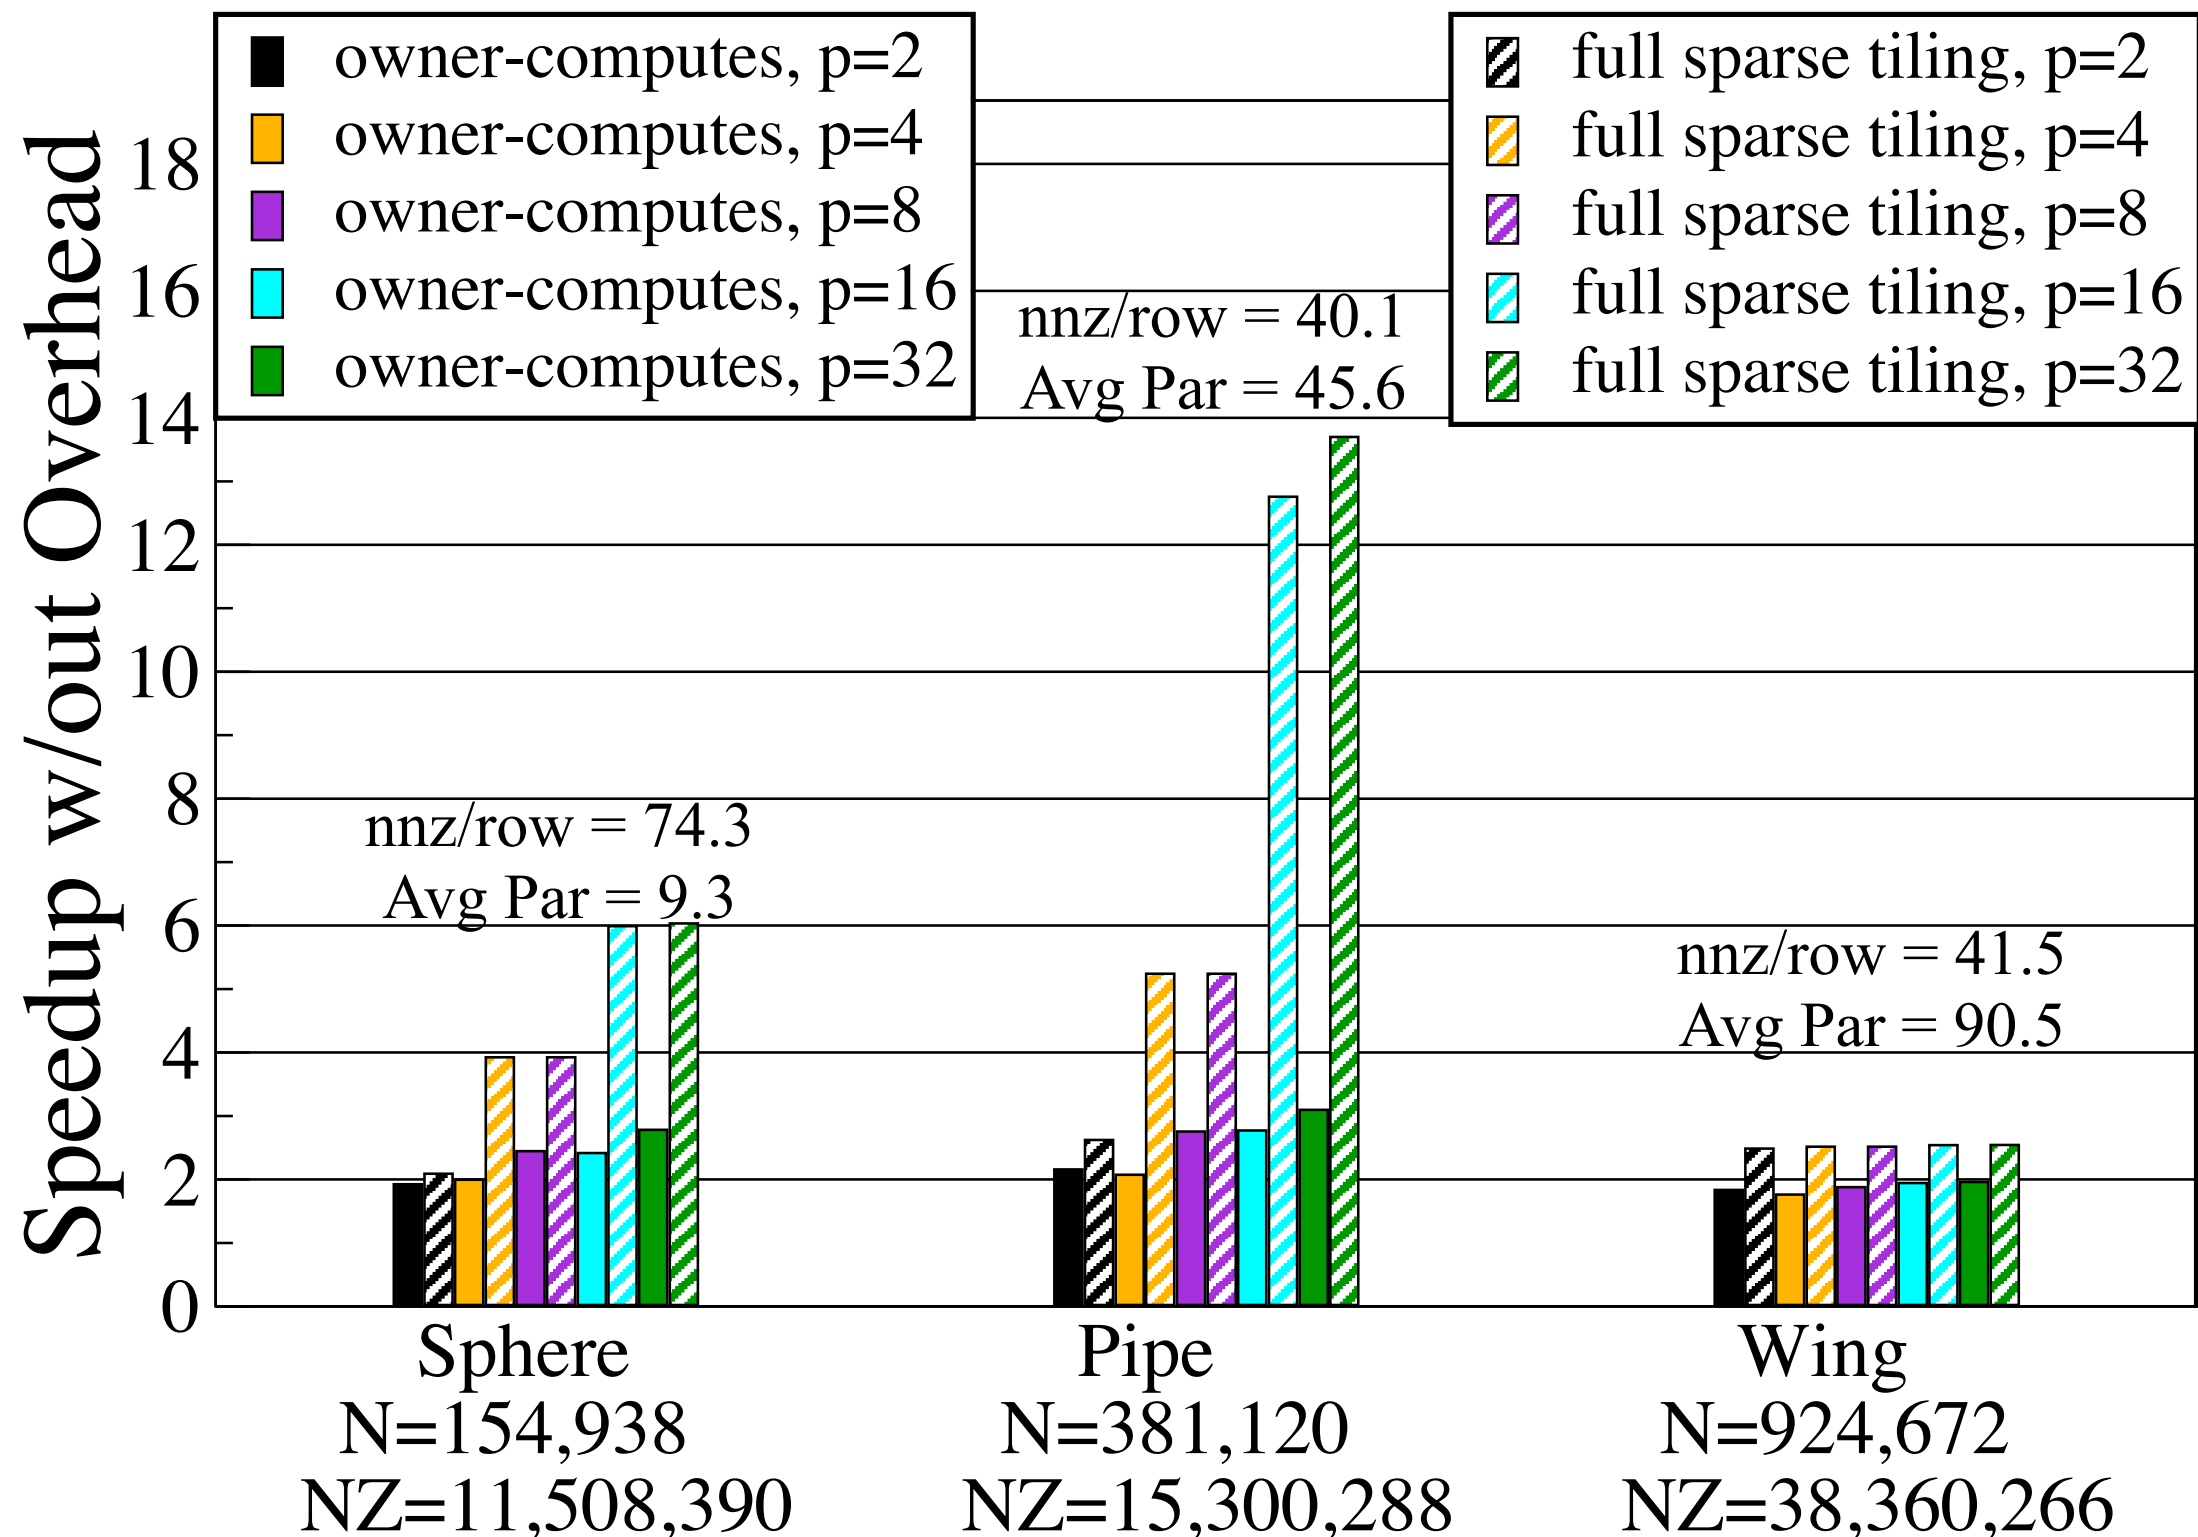
# Experimental Methodology

## Part 1

- **Baseline is Gauss-Seidel implemented CSR and uses provided ordering**
- **Owner computes implementation**
  - Partition matrix graph into equal-sized cells for each processor
  - Sub-partition and reorder on each processor for intra-iteration locality
  - Violate intra-iteration dependences for an idealistic parallel efficiency
- **Full sparse tiling implementation**

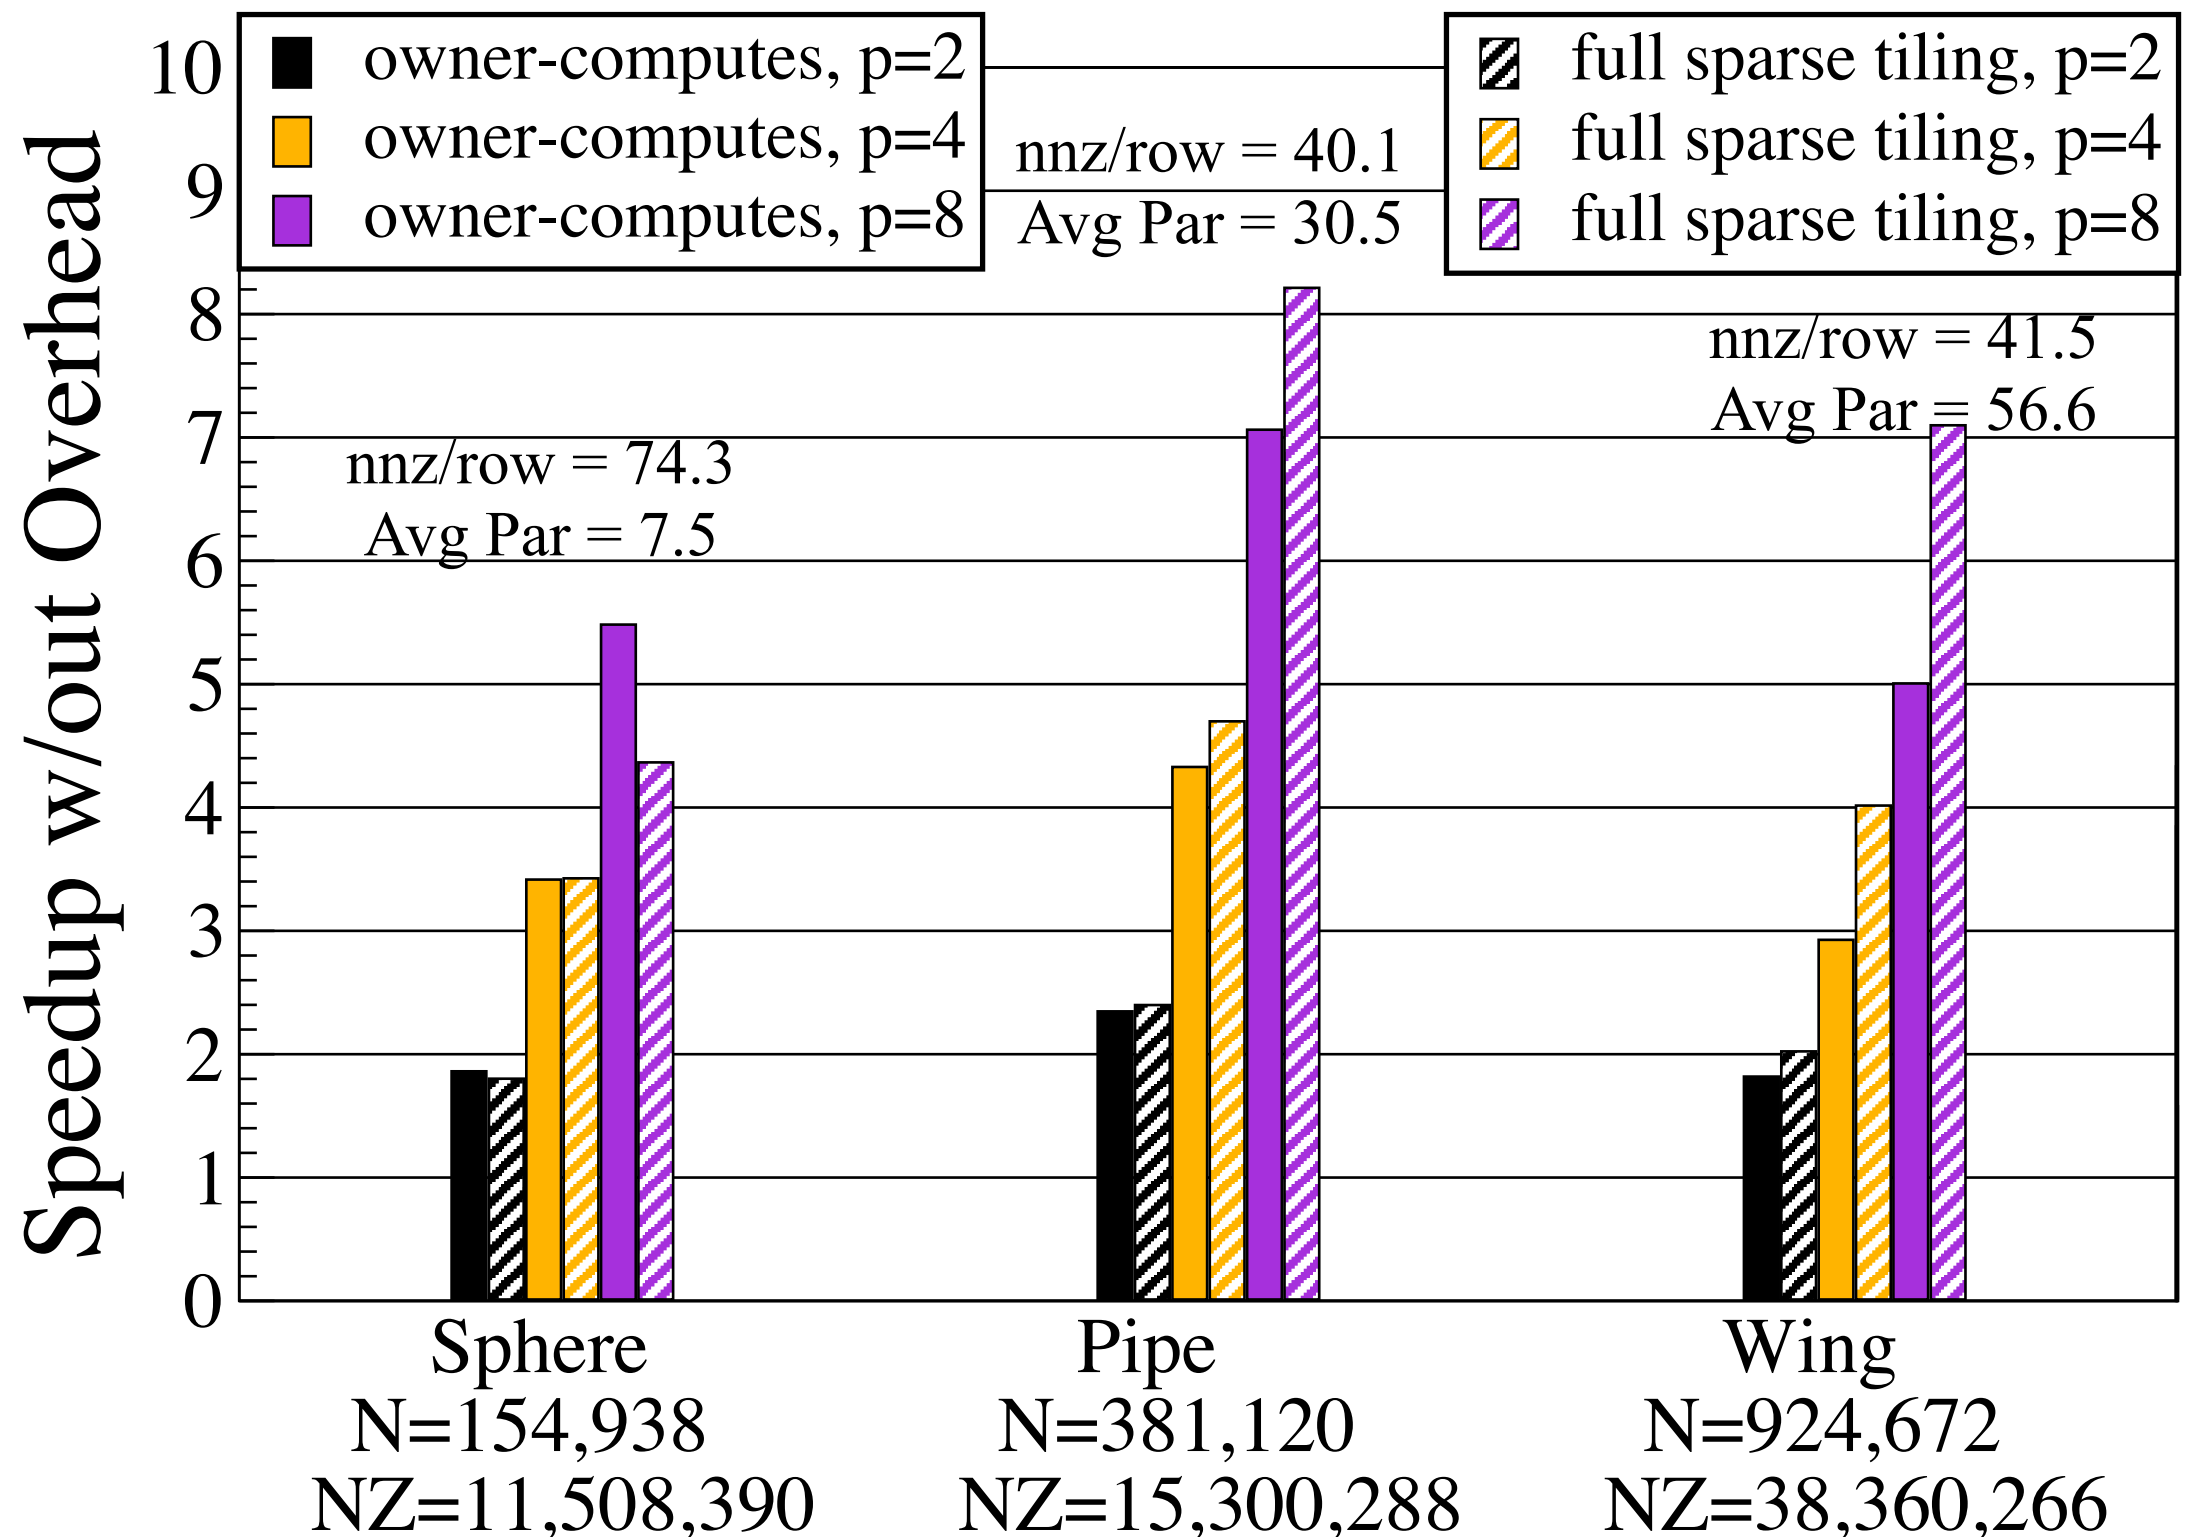
# Experimental Results

SUN HPC10000, 36 UltraSPARC II, 400MHz, 4MB L2



# Experimental Results

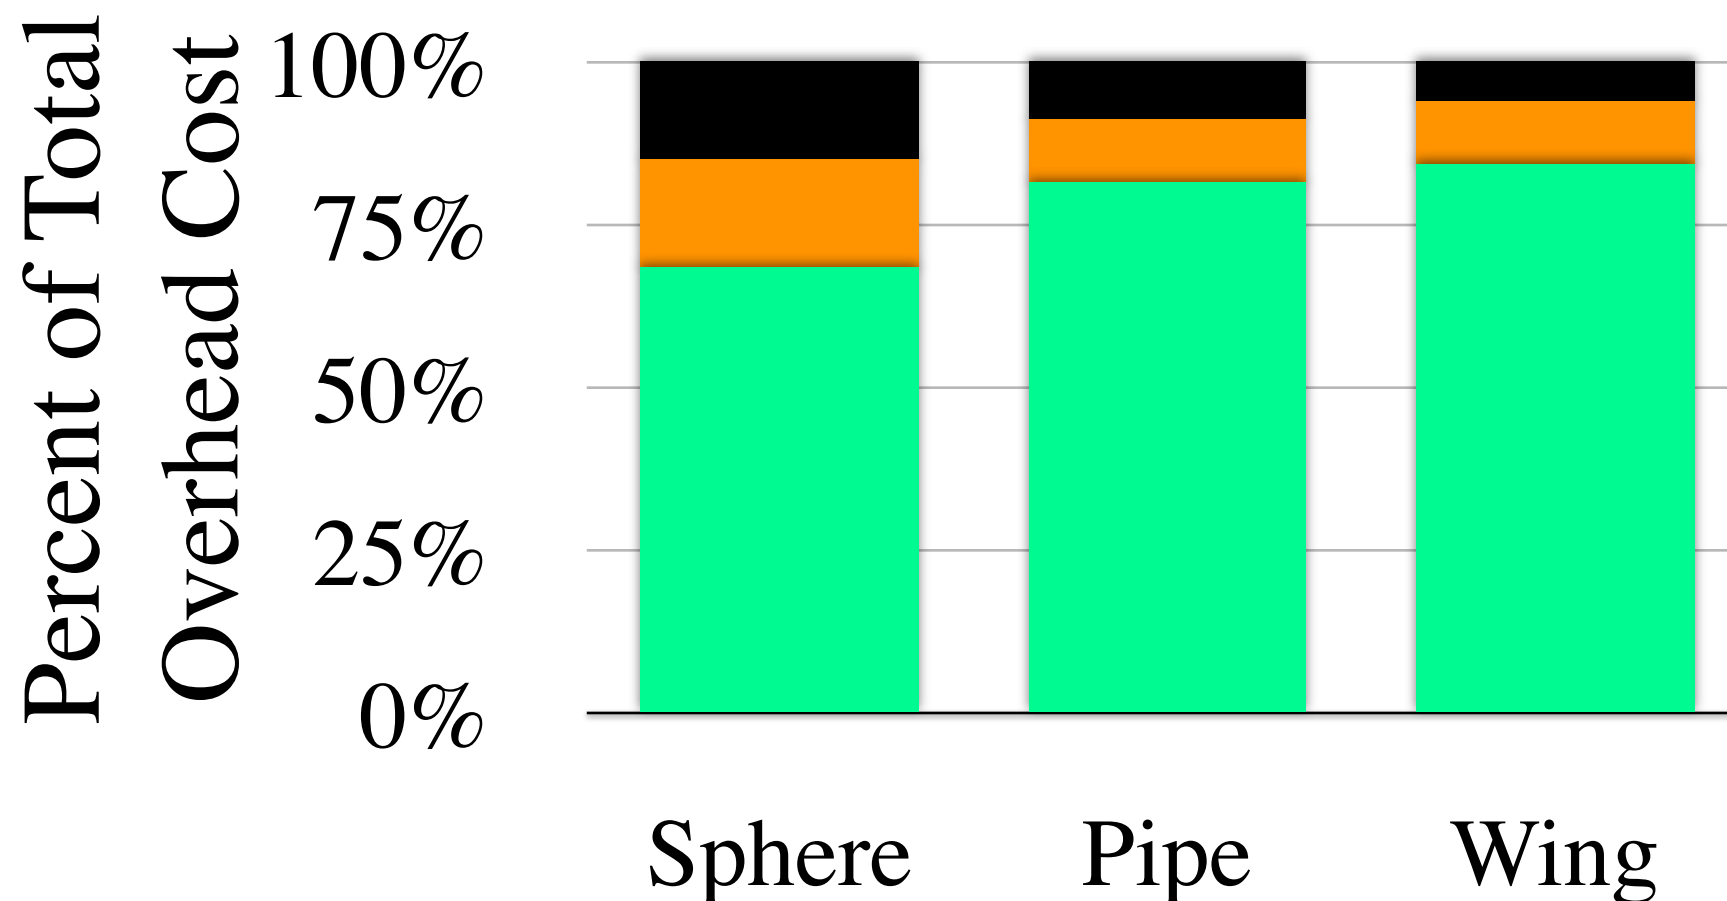
IBM Power 3, 375MHz, 8MB L2 cache



# Overhead of Full Sparse Tiling on IBM Power 3

- Break even at **114 thru 400** convergence iterations
- Partitions created with Metis package [Karypis98]

■ Partitioning    ■ Data Remapping    ■ Tile Growth



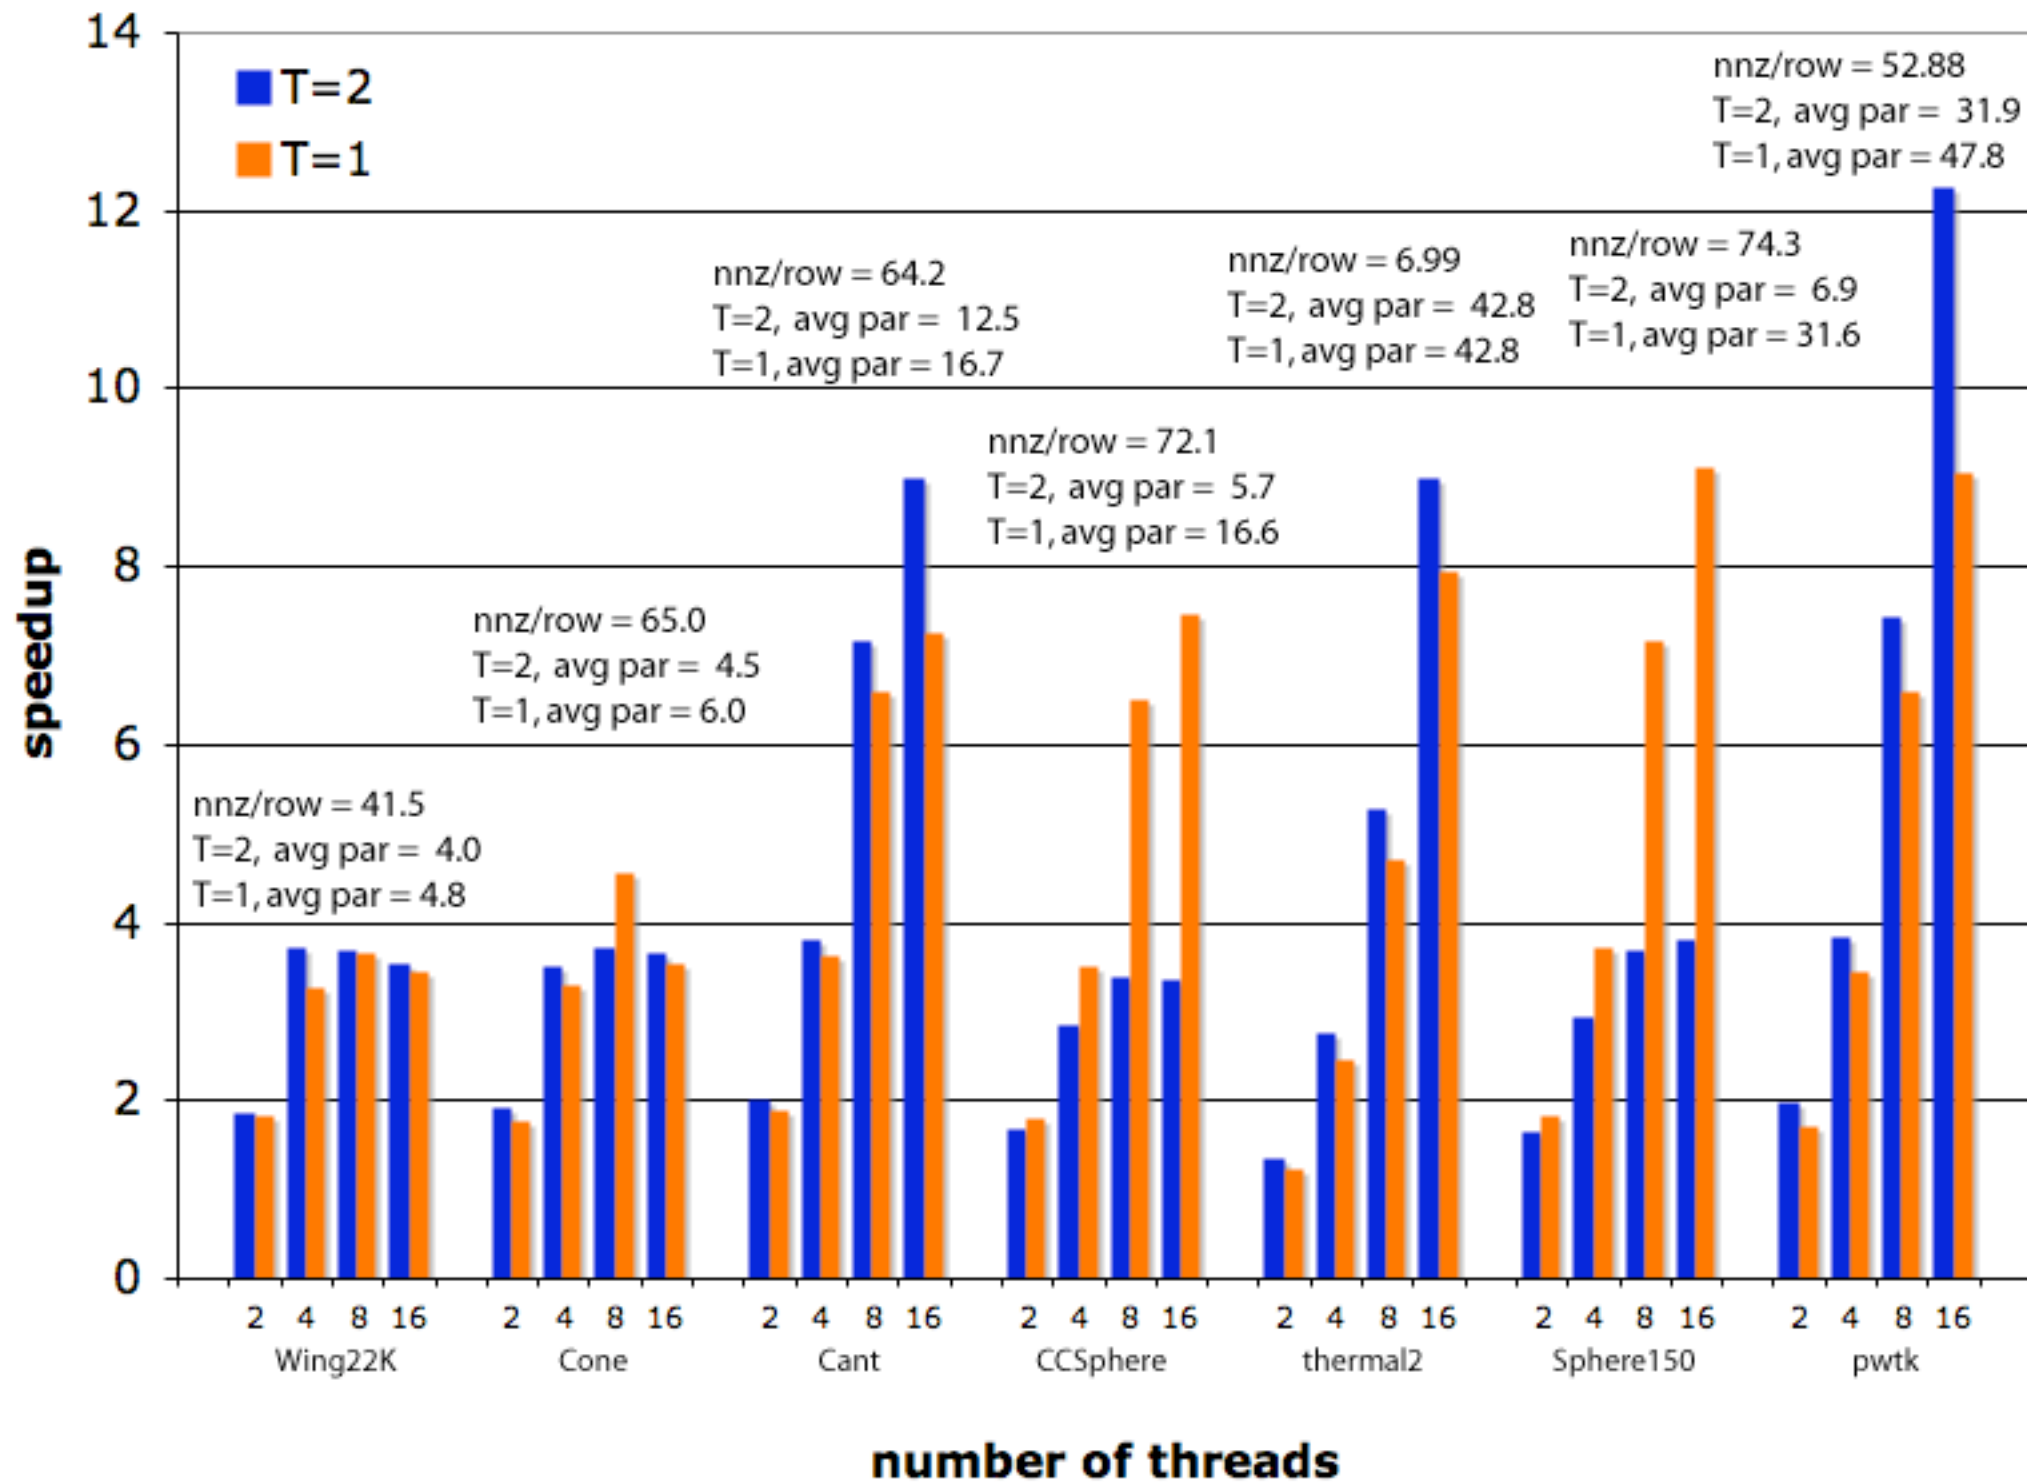
# Experimental Methodology

## Part 2

- **Baseline is one processor version of full sparse tiled Gauss-Seidel**
- **Owner computes implementation**
  - one convergence iteration is partitioned and a task graph is constructed between partitions,  $T = 1$
- **Full sparse tiling implementation**
  - one convergence iteration is partitioned and a task graph is constructed between partitions,  $T = 2$

# Experimental Results for Blue Ice

IBM Power 5+, 1.9GHz, 1.9MB L2 cache, 36MB L3 cache





# Parameters Identified for Possible Auto-tuning

- Given mesh characteristics such as the average non-zeros per row, size, etc., model and tune for the following:
  - seed partition size
  - number of iterations to tile across
  - static or dynamic tile scheduling for parallelism
  - parameters for either scheduling strategy
- Would overlapping tiles remove the parallelism versus data locality tradeoff?

# Conclusions

- Full sparse tiling exploits parallelism, intra-iteration data reuse, and inter-iteration data reuse
- Coloring the partition graph to number seed partitions improves the average parallelism in the tile dependence graph
- On shared memory processors, full sparse tiling can outperform owner-computes methods when enough parallelism is present
- More study needed of temporal data locality versus parallelism tradeoff