

Tools for HPC

APOLENT
Corporation

COREquations Division

Emerging Platforms

- Moore's Law
 - Power Wall
 - Memory Wall
 - ...
- Complex Parallel Systems are here
 - Core-to-Core comm./sync. only through L3 cache in AMD Barcelona and INTEL Nehalem proc. (L1/L2 private)
 - Processor-to-Processor comm./sync. even more costly in multi-socket configurations
 - Heterogenous multi-cores
Cell BE, Host + GPUs (NVIDIA, AMD), Host + Larrabee₂

Programming Requirements

- Need to carefully examine
 - Data Partitioning
 - Communication
 - Synchronization

(A standard data-parallel implementation may be insufficient)

Relevant Applications/Domains

- Image/Video Processing, Vision, etc.
- Scientific/Engg./Statistical Computation
 - Computer Simulation of Physics, Signal Proc., SCADA/HMI
- Game Physics Engines
 - (SONY) Bullet Physics, (NVIDIA) Physx, (INTEL) Havoc
- CG Animation
 - Autodesk 3DS Max/Maya/XSI, Houdini, PIXAR PRMan, Mental Images. Also (open-source) Blender, Pixie & Aqsis
- Bioinformatics, RNA Structure Prediction
- Geospatial Information Systems (GIS)
 - ESRI ArcGIS, PCI Geomatica. Also (open-source) GRASS

Application Hotspots

- Focus of Attention (Vision, Image Processing)
 - Image filtering, edge detection, feature extraction
- Motion Estimation: H.264 (Video Processing)
- Numerical Solutions to PDEs, Filters, Regression, Control & Optimization, Parameter Estimation (Sc./Engg./Stat.)
- Dynamics (Physics Simulation)
 - Fluid/Cloth/Wire/Body Simulation in Games/Movies
- Rendering (CG Animation)
 - Global Illumination/Ray Tracing
- Dynamic Programming Kernels (Bioinformatics)
- Orthorectification (GIS)

Hotspot Characteristics

- Highly Compute-Intensive
- High Data Volume
- Very well structured kernel fragments

Possible Solutions

- Automatically generate libraries
- Develop highly tuned libraries

Automatic Library Generation (The COREquations Engine)

Automatic Library Generation

- Domain Specific Compiler Engine
- Class of Problems: The Polyhedral Model
 - Incorporates most parts of four dwarfs in the " Berkeley View* "
(each was expected to require an independent technology)
 - Structured Grids
 - Dense Matrix
 - Most of Dynamic Programming
 - Some of Graphical Models


* The Berkeley View is a recent study by application/compiler/hardware specialists that outlines the most important problem domains for parallel platforms

The Polyhedral Model

- Most dynamic instr. from iterative kernels
- Manual code optimization is error prone and time consuming
- We wish to develop equivalent high level compiler optimizations
 - Analyze program behavior
 - Design custom transformations
 - Guarantee validity

Parallel Code Generation

- Host owns data (partitions, communicates & synchronizes)
- "Compute Engines" process data
- Host reports results back to user
- Heterogeneous threads



Specifications

- Single Assignment Loops (actually equations)
- Matrix Multiplication (The HPC "Hello World")

```
affine matrix_product {P, Q, R | P>0 && Q>0 && R>0}
  given float A {i,k | 0<=i<P && 0<=k<Q };
  given float B {k,j | 0<=k<Q && 0<=j<R };
  returns float C {i,j,k | 0<=i<P && 0<=j<R && k==Q-1 };
  using float temp_C {i,j,k | 0<=i<P && 0<=j<R && -1<=k<Q };
  through
    temp_C[i,j,k] = case
      { | k>=0 } : temp_C[i,j,k-1] + (A[i,k] * B[k,j]);
      { | k==-1 } : 0;
    esac;
  C[i,j,k] = temp_C[i,j,k];
```

Program Parameters

Input Variables

Output Variable

Local Variable

Equation for Accumulator

Equation for Result

Script-Driven Compilation

```
source("Setup.bsh");
```

Read the program

```
ConnectServer("localhost");
```

```
ReadProgram("matrix_product.alphabets");
```

```
SetMemoryMap("matrix_product","C",
```

Desired memory layout of output

```
"(P,Q,R,i,j,k->P,Q,R,i,j)","0,0");
```

```
MPICodeGen(8,"32,32,32");
```


Generate code for 8 processors,
tiling for parallelism with sizes
32x32x32 along i,j and k

Memory Maps are only needed for output variables

Tool finds optimal memory layout for local variables

Performance I: Scaling

- Near-Perfect Scaling



Performance II: Raw Comparison


- Compare to Hand-Optimized Sequential Code
- Before, our MPI Code Generator, we wanted to generate Good Sequential Code

Sequential Code Generation

- Case Study: PDE Solver from Atmospheric Sc.
 - 2D-3D stencil. Five-point update
 - Over time, elements on an area update their value based on previous value and North, South, East and West neighbors
 - Periodic boundary conditions
 - West border depends on East border, etc.
 - SW corner depends on SE and NW corners
- Hand-Optimized by Nathan Burnett, a Masters Student at CSU
 - Performed optimizations such as loop interchange, fusion, fission, unrolling, removal of lookup-tables ...

Performance (Hand-Optimization)

- Intel CORE2 Duo (2.3Ghz)



EI PDE Solver

through

```
h[i,j,k] = case
    { | k==0 } : h_init[j,i];
    { | 0<k } : h[i,j,k-1] - delt * H0 * delta[i,j,k-1];
esac;
```

Main Variable

lap_h[i,j,k] =

case

```
{ | i==0 && j==0 } : (h[i+1,j,k] + h[2im_h-1,j,k] + h[i,j+1,k] + h[i,2jm_h-1,k]
SW Corner          - 4*h[i,j,k]) * inv_d_sq;
```

...

```
{ | i==0 && 0<=j<2jm_h } : (h[i+1,j,k] + h[2im_h-1,j,k] + h[i,j+1,k] + h[i,j-1,k]
West Border            - 4*h[i,j,k]) * inv_d_sq;
```

...

```
{ | 0<i<2im_h-1 && 0<j<2jm_h-1 } : (h[i+1,j,k] + h[i-1,j,k] + h[i,j+1,k] + h[i,j-1,k]
Main Body              - 4*h[i,j,k]) * inv_d_sq;
```

esac;

Key Computation

```
delta[i,j,k] = case
```

```
{ | k==0 } : 0;
```

```
{ | 0<k } : delta[i,j,k-1] - delt * g * lap_h[i,j,k];
```

esac;


Auxillary Variable

```
results[j,i] = h[i,j,TMAX];
```

Output

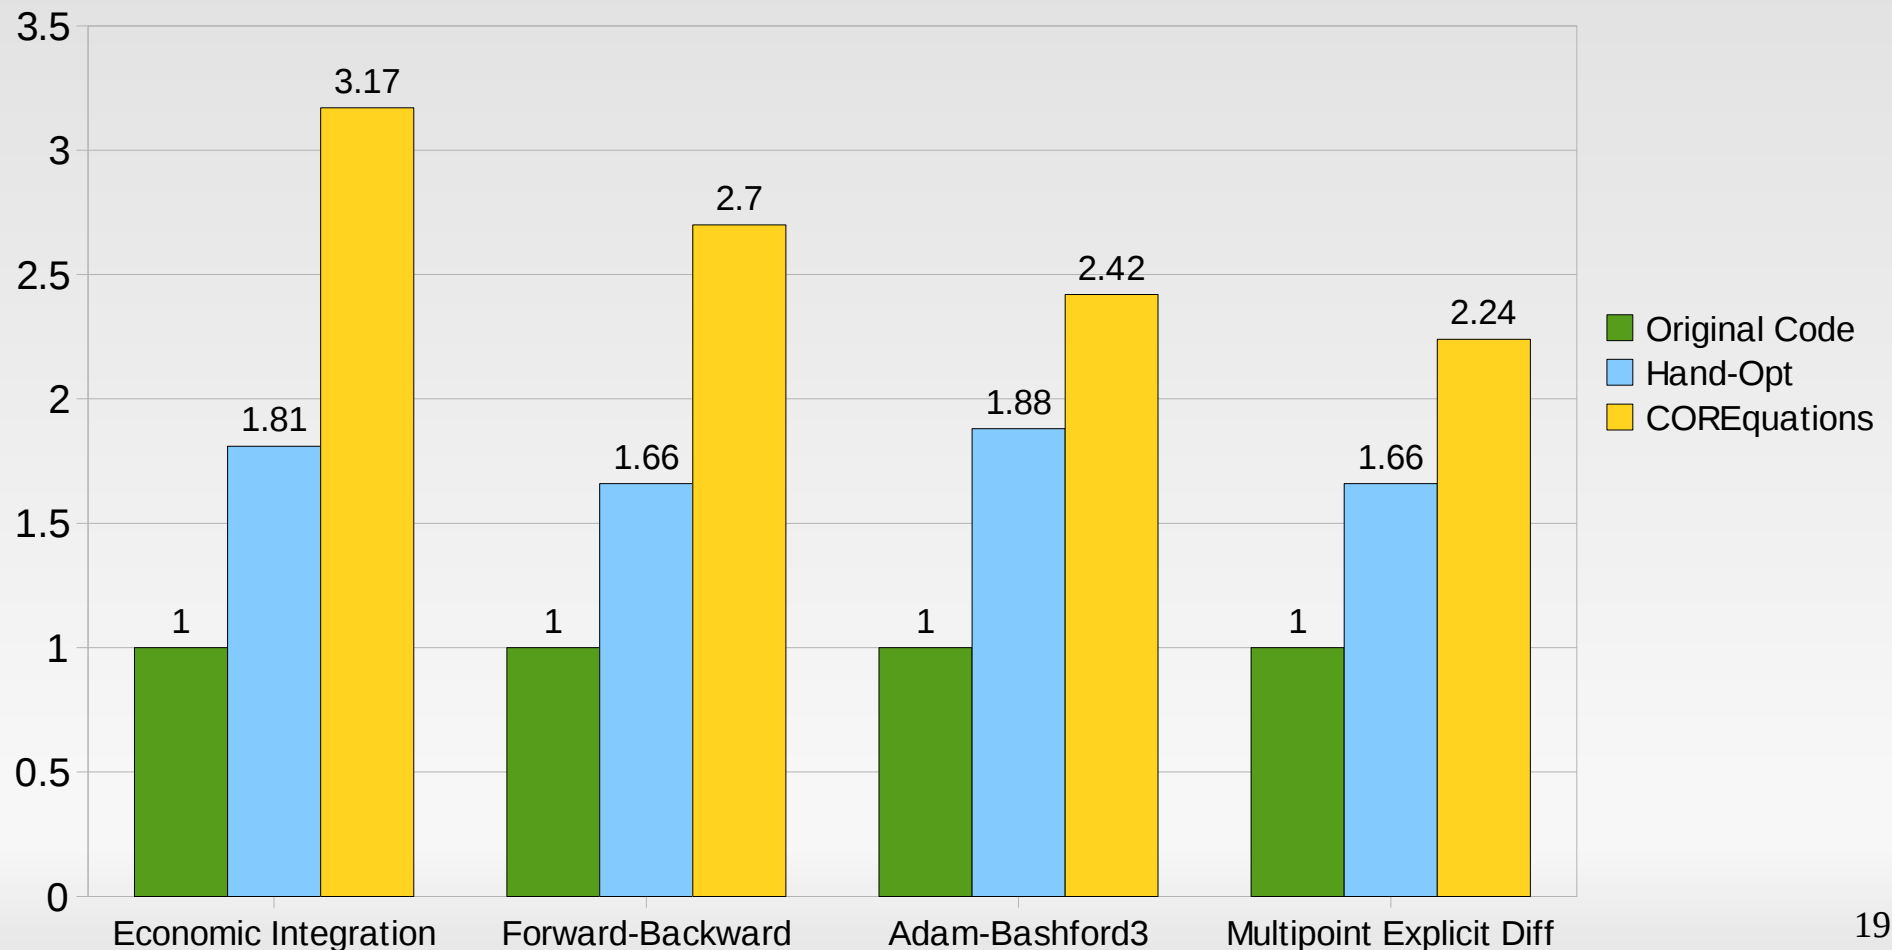
Performance (Hand-Optimization)

- Intel CORE2 Duo (2.3Ghz)



Then, I provided code for his equations

- Intel CORE2 Duo (2.3Ghz)



Why is generated code so good?

- Memory Optimization
 - (Prefetch-friendly) Alignment
 - Storage Minimization
- Loop Optimizations (just much more of them): Automatic loop-generator specializes code for all boundary cases
 - Aggressive Fusion
 - Aggressive Fission

Triangular Matrix Multiplication

- Beyond "Hello World"
- $P=Q=R$ ($=N$ say). Add triangular constraints

```
affine triangular_matrix_product {N|N>0}
  given float A {i,k | 0<=i<N && 0<=k<=i };
      float B {k,j | 0<=k<N && k<=j<N };
returns float C {i,j,k | 0<=i<=j && 0<=j<N && k==i } ||
      {i,j,k | 0<=i<N && 0<=j<i && k==j } ;
using float temp_C {i,j,k | 0<=i<N && 0<=j<N && -1<=k<=(i,j) };
through
temp_C[i,j,k] = case
    { | k>=0 } : temp_C[i,j,k-1]+(A[i,k]*B[k,j]);
    { | k==-1 } : 0;
    esac;
C[i,j,k] = temp_C[i,j,k];
```

Triangular Matrix Multiplication


- Sequential Code

```
for i = 0 to n-1 {  
  for j = 0 to n-1 {  
    C[i,j] = 0;  
    for k = 0 to min(i,j) {  
      C[i,j] += A[i,k]*B[k,j]  
    }  
  }  
}
```

Why write equations when you can write code?


- Iteration Space is a polyhedron

- Now, let us generate MPI code
(What is each Compute Engine iterating over?)



Triangular Matrix Multiplication

- Targeting memory-constrained architectures
 - Need to tile all dimensions (not necessarily cubic)
 - Many Partial Tiles
 - All Results are from Partial Tiles
 - Result distributed in block-cyclic along i




CE 1 requires a partial block of inputs from B and a partial block from A
CE 2 requires a partial block of inputs from B but a full block of inputs from A

In Short ...

- Writing MPI code even for a straightforward specialization of Matrix Multiplication is extremely time consuming and error-prone
- In COREquations, just run the same script for the new specification

Performance

- Again, Near-Perfect Scaling



Design Space Exploration

- Sequential Code for Matrix Multiplication

```
for i = 0 to n-1 {  
  for j = 0 to n-1 {  
    C[i,j] = 0;  
    for k = 0 to N-1 {  
      C[i,j] += A[i,k] * B[k,j]  
    }  
  }  
}
```

Subsequent reads of A are along the inner index.
Memory accesses to A are prefetch friendly.

Subsequent reads of B are along the outer index.
Memory accesses to B are not prefetch friendly.

- What about writes to C?

Writes to C are to the same location
throughout the inner loop

Design Space Exploration

- A number of choices in code-generation

- Tiling: (shape and size)

- ```
FixedTiledCodeGen("48,32,63");
```

- Different Memory Layouts of Variables

- ```
SetMemoryMap("matrix_product", "B",  
              "(P,Q,R,k,j->P,Q,R,j,k)", "0,0");
```

- Permutation of Loops

- With j as the inner loop, A and C are accessed only once along the inner loop. B[k,j] becomes Prefetch-friendly.

- ```
CoB("matrix_product", "C", "(P,Q,R,i,j,k->P,Q,R,i,k,j)");
```

- ```
CoB("matrix_product", "temp_C", "(P,Q,R,i,j,k->P,Q,R,i,k,j)");
```

Future Work

- We're going to make it easier to program
 - Higher Level Language
 - Matrix Product should simply be specified as
$$C[i,j] = \text{SUM}([k], A[i,k]*B[k,j]);$$
 - Domain Inferencing
- More Analysis (Scheduling, etc.)
- Get the Human-Out-of-the-Loop
- Porting to Newer Architectures (GPUs, Cell, Larrabee, ...)

Summary

- For specifications in Alphabets
 - MPI is easy
 - Design Space Exploration is easy
 - Scalable Performance
- Auto-verification
 - Compiler-assisted programming to avoid double-or incomplete-definitions, out-of-bound references ...
 - Major source of errors is eliminated

Generality of Solution

- A new language for generating HPC code for compute-intensive kernels, but also ...
An Intermediate Representation for optimization of loops in conventional languages

Appendix

Inexpensive Design Space Exploration

Simple Matrix Product

```
Bsh Workspace: 0
File Font
BeanShell
2.0b4 - by Pat Niemeyer (pat@pat.net)
bsh % source("Setup.bsh");
bsh % ConnectServer("localhost");
bsh % ReadProgram("../examples/embedded/matrix_product.alphabets");
bsh % SetMemoryMap("matrix_product", "C", "(P,Q,R,i,j,k->P,Q,R,i,j)", "0,0");
bsh % CodeGen();
bsh % GenerateWrappers();
bsh %
```

Refer to script on slide 7

```
Bsh Works
int i_, j_, k_;
#define S1(i_,k_,j_) temp_C((i_), (k_), (j_)) = (temp_C((i_), (k_), (j_)-1)) + ((A((i_), (j_))) * (B((j_), (k_))))
#define S2(i_,k_,j_) temp_C((i_), (k_), (j_)) = 0
#define S3(i_,k_,j_) C((i_), (k_), (j_)) = temp_C((i_), (k_), (j_))
/* Generated from by CLoog v0.14.1 gmp bits in 0.06s. */
/* CLoog asked for 17260 KBytes. */
if (Q_ >= 2) {
  for (i_ = 0; i_ <= P_-1; i_++) {
    for (j_ = 0; j_ <= R_-1; j_++) {
      S2(i_, j_, -1);
      for (k_ = 0; k_ <= Q_-2; k_++) {
        S1(i_, j_, k_);
      }
      k_ = Q_-1;
      k_ = Q_-1;
      S1(i_, j_, Q_-1);
      S3(i_, j_, Q_-1);
    }
  }
}
if (Q_ == 1) {
  for (i_ = 0; i_ <= P_-1; i_++) {
    for (j_ = 0; j_ <= R_-1; j_++) {
      S2(i_, j_, -1);
      S1(i_, j_, 0);
      S3(i_, j_, 0);
    }
  }
}
```

```
File Edit View Terminal Tabs Help
[ggupta@kenshin 01:31:11 temp_dir] gcc -O3 -o simple_mm matrix_product-seqC.c matrix_product-main.c
[ggupta@kenshin 01:35:05 temp_dir] ./simple_mm 1024 1024 1024
Execution time : 33.002811 sec.
[ggupta@kenshin 01:35:40 temp_dir] █
```

Part of Generated Code

Execution Time

Tiled Matrix Product

```
Bsh Workspace: 2
File Font
BeanShell
2.0b4 - by Pat Niemeyer (pat@pat.net)
bsh % source("Setup.bsh");
bsh % ConnectServer("localhost");
bsh % ReadProgram("../examples/embedded/matrix_product.alphabets");
bsh % AShow();
affine matrix_product {P,Q,R | P>=1 && Q>=1 && R>=1}
given
  float A {i,k | k>=0 && i>=0 && Q-k>=1 && P-i>=1};
  float B {k,j | j>=0 && k>=0 && R-j>=1 && Q-k>=1};
returns
  float C {i,j,k | Q-k==1 && j>=0 && i>=0 && R-j>=1 && P-i>=1};
using
  float temp_C {i,j,k | k>=-1 && j>=0 && i>=0 && R-j>=1 && Q-k>=1 &&
P-i>=1};
through

temp_C[i,j,k] = case
  { | k>=0} : (temp_C[i,j,k-1]+(A[i,k]*B[k,j]));
  { | k== -1} : 0;
esac;

C[i,j,k] = temp_C[i,j,k];
|
.

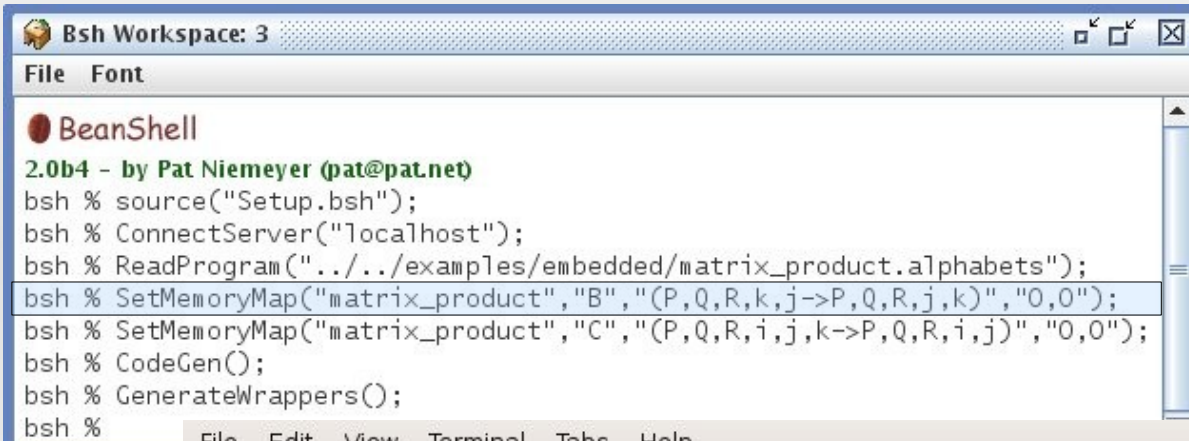
bsh % SetMemoryMap("matrix_product","C","(P,Q,R,i,j,k->P,Q,R,i,j)","0,0");
bsh % FixedTiledCodeGen("32,32,32");
bsh % GenerateWrappers();
bsh %
```

Loaded Program

Bsh Workspace: 2

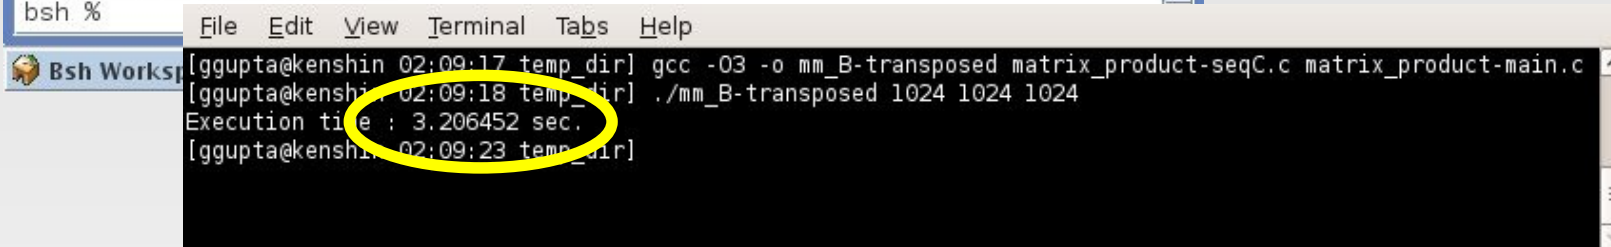
```
File Edit View Terminal Tabs Help
[ggupta@kenshin 01:53:05 temp_dir] gcc -O3 -o tiled_mm matrix_product-fixedtiled.c matrix_product-main.c -lm
[ggupta@kenshin 01:53:08 temp_dir] ./tiled_mm 1024 1024 1024
Execution time : 15.598294 sec.
[ggupta@kenshin 01:53:25 temp_dir]
```

Changing the memory layout of B



The screenshot shows the BeanShell IDE interface. The title bar reads "Bsh Workspace: 3". Below the menu bar (File, Font), the text "BeanShell" is displayed. The main editor area contains the following code:

```
2.0b4 - by Pat Niemeyer (pat@pat.net)
bsh % source("Setup.bsh");
bsh % ConnectServer("localhost");
bsh % ReadProgram("../examples/embedded/matrix_product.alphabets");
bsh % SetMemoryMap("matrix_product", "B", "(P,Q,R,k,j->P,Q,R,j,k)", "0,0");
bsh % SetMemoryMap("matrix_product", "C", "(P,Q,R,i,j,k->P,Q,R,i,j)", "0,0");
bsh % CodeGen();
bsh % GenerateWrappers();
bsh %
```

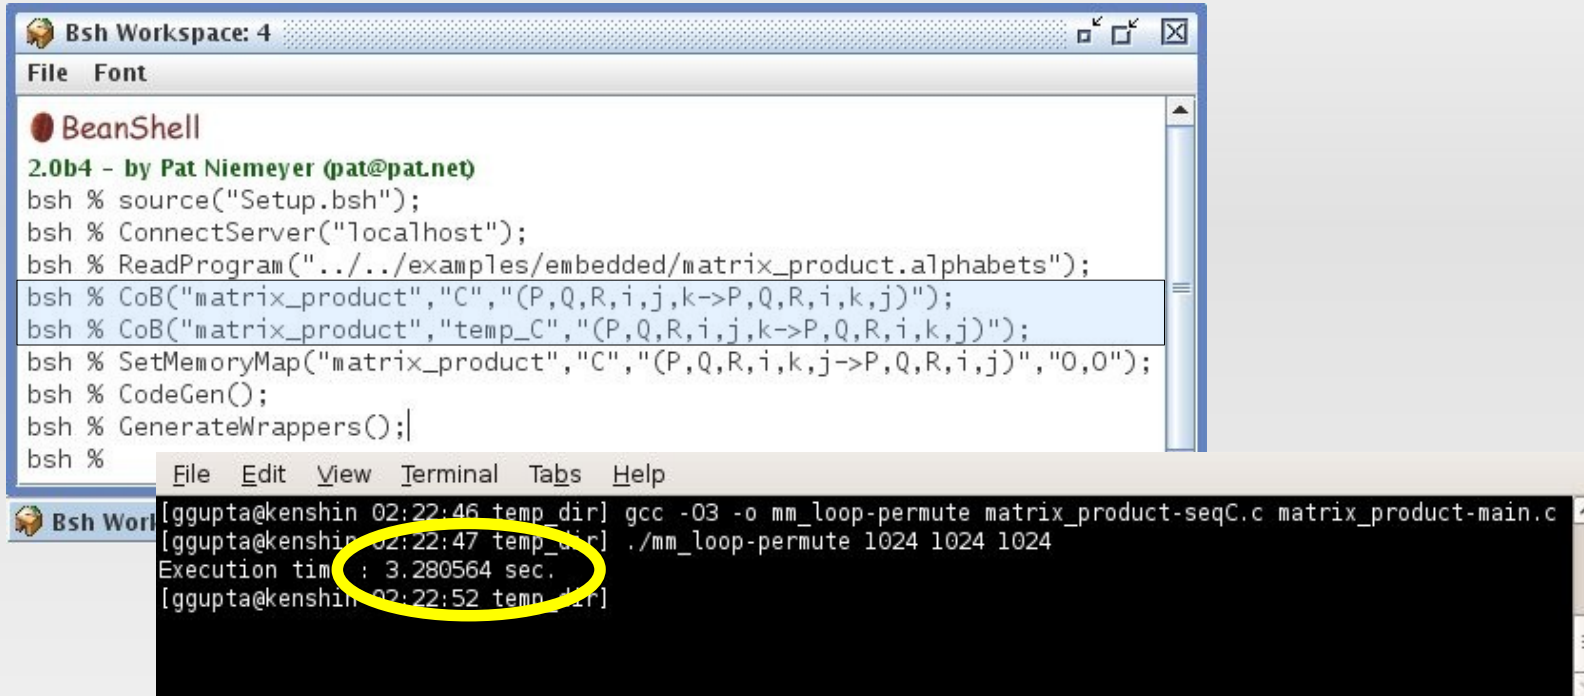


The screenshot shows a terminal window with the following output:

```
[ggupta@kenshin 02:09:17 temp_dir] gcc -O3 -o mm_B-transposed matrix_product-seqC.c matrix_product-main.c
[ggupta@kenshin 02:09:18 temp_dir] ./mm_B-transposed 1024 1024 1024
Execution time : 3.206452 sec.
[ggupta@kenshin 02:09:23 temp_dir]
```

- Significant performance upgrade by having a prefetch-friendly layout of the input array B
- Can we obtain similar performance with the standard memory layout for B

Permuting loop indices



The screenshot shows a BeanShell workspace window titled "Bsh Workspace: 4" with a menu bar (File, Font) and a toolbar. The BeanShell version is 2.0b4, by Pat Niemeyer (pat@pat.net). The code in the workspace is as follows:

```
bsh % source("Setup.bsh");  
bsh % ConnectServer("localhost");  
bsh % ReadProgram("../examples/embedded/matrix_product.alphabets");  
bsh % CoB("matrix_product", "C", "(P,Q,R,i,j,k->P,Q,R,i,k,j)");  
bsh % CoB("matrix_product", "temp_C", "(P,Q,R,i,j,k->P,Q,R,i,k,j)");  
bsh % SetMemoryMap("matrix_product", "C", "(P,Q,R,i,k,j->P,Q,R,i,j)", "0,0");  
bsh % CodeGen();  
bsh % GenerateWrappers();  
bsh %
```

Below the workspace is a terminal window titled "Bsh Worl". It shows the compilation and execution of the program:

```
[ggupta@kenshin 02:22:46 temp dir] gcc -O3 -o mm_loop-permute matrix_product-seqC.c matrix_product-main.c  
[ggupta@kenshin 02:22:47 temp dir] ./mm_loop-permute 1024 1024 1024  
Execution time : 3.280564 sec.  
[ggupta@kenshin 02:22:52 temp dir]
```

The execution time "3.280564 sec." is circled in yellow in the original image.

- By having j as the innermost loop index, accesses (to arrays A and B) preserve locality